

# Topic V2D

Instruction-Level Parallelism

Reading: (Section 4.10)

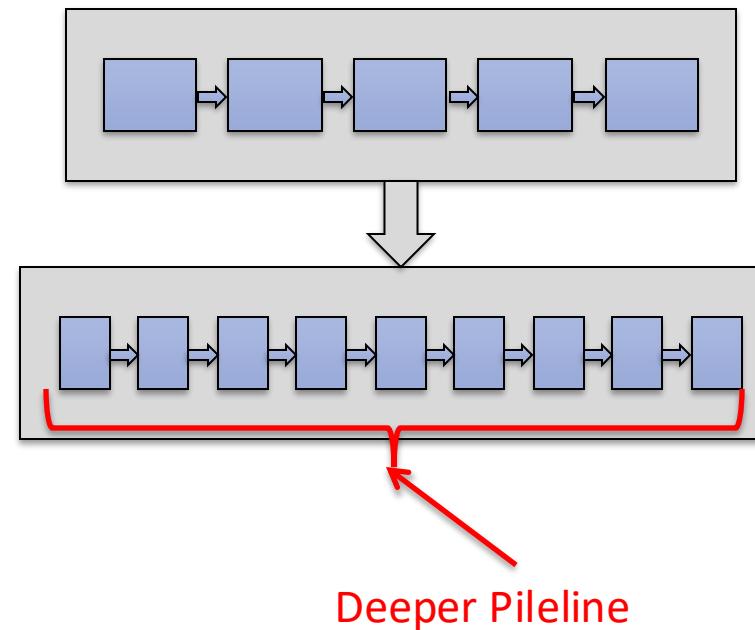
# Instruction-Level Parallelism (ILP)

Pipelining: executing multiple instructions in parallel

To increase ILP

Deeper pipeline

Less work per stage  $\Rightarrow$  shorter clock cycle



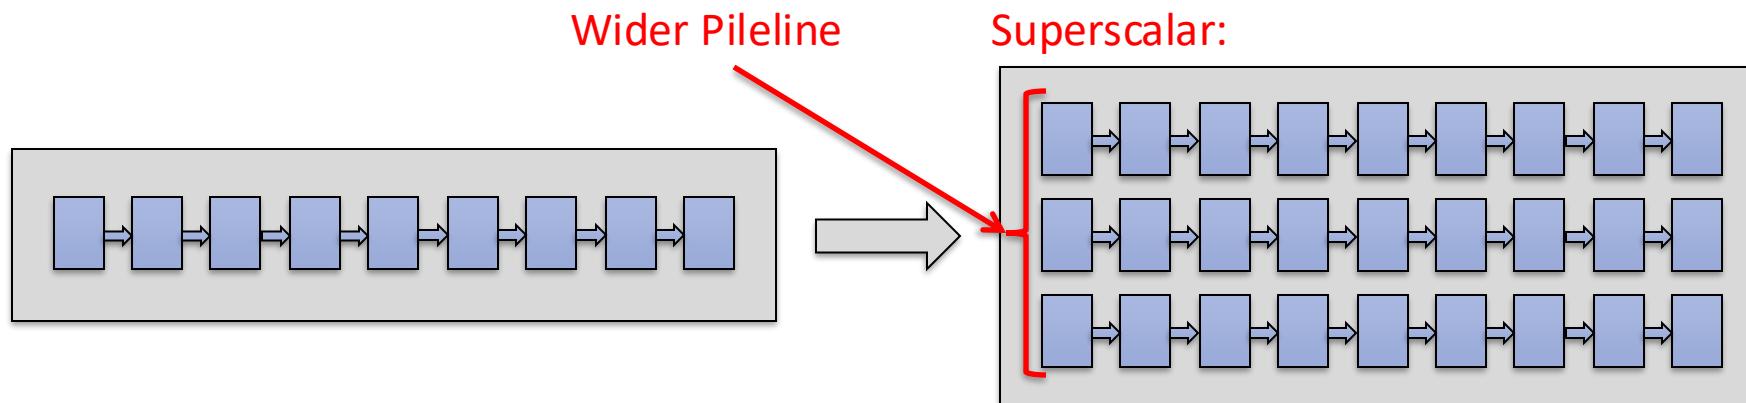
# Multiple Issue

Replicate pipeline stages  $\Rightarrow$  multiple pipelines

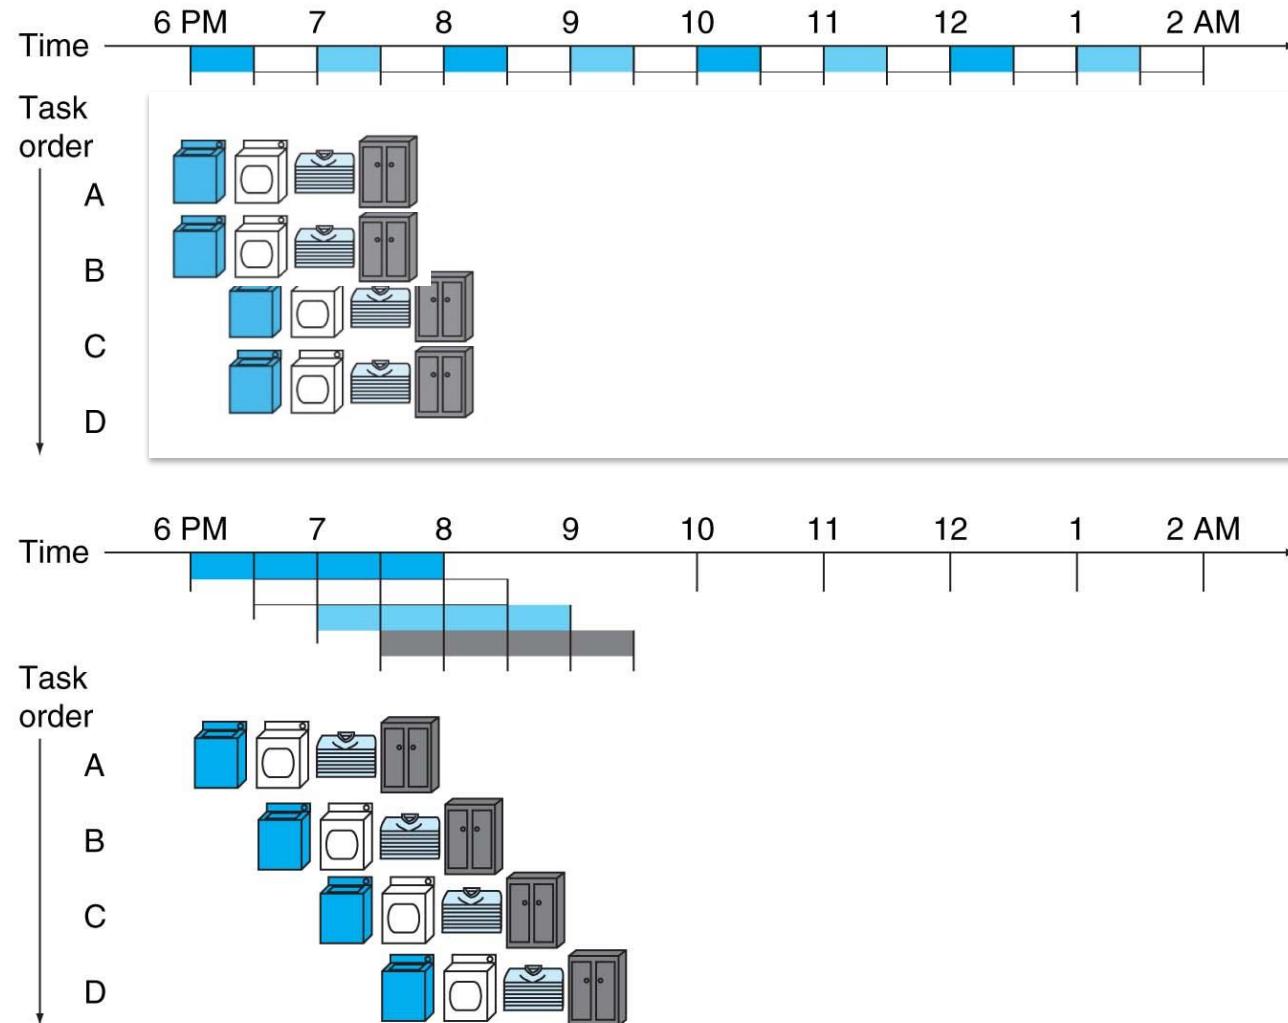
Start multiple instructions per clock cycle

CPI < 1, so use Instructions Per Cycle (IPC)

But dependencies limit ILP in practice



# Multiple Issue



# Multiple Issue

Static multiple issue - Very Long Instruction Word (VLIW)

Compiler groups instructions to be issued together

Packages them into “issue slots”

Compiler detects and avoids hazards

Dynamic multiple issue

CPU examines instruction stream and chooses instructions to issue each cycle

Compiler can help by reordering instructions

CPU resolves hazards using advanced techniques at runtime

# Static Multiple Issue

Compiler groups instructions into “issue packets”

- Group of instructions that can be issued on a single cycle

- Determined by pipeline resources required

Think of an issue packet as a very long instruction

- Specifies multiple concurrent operations

- ⇒ Very Long Instruction Word (VLIW)

# Dynamic Pipeline Scheduling

Allow the CPU to execute instructions out of order to avoid stalls

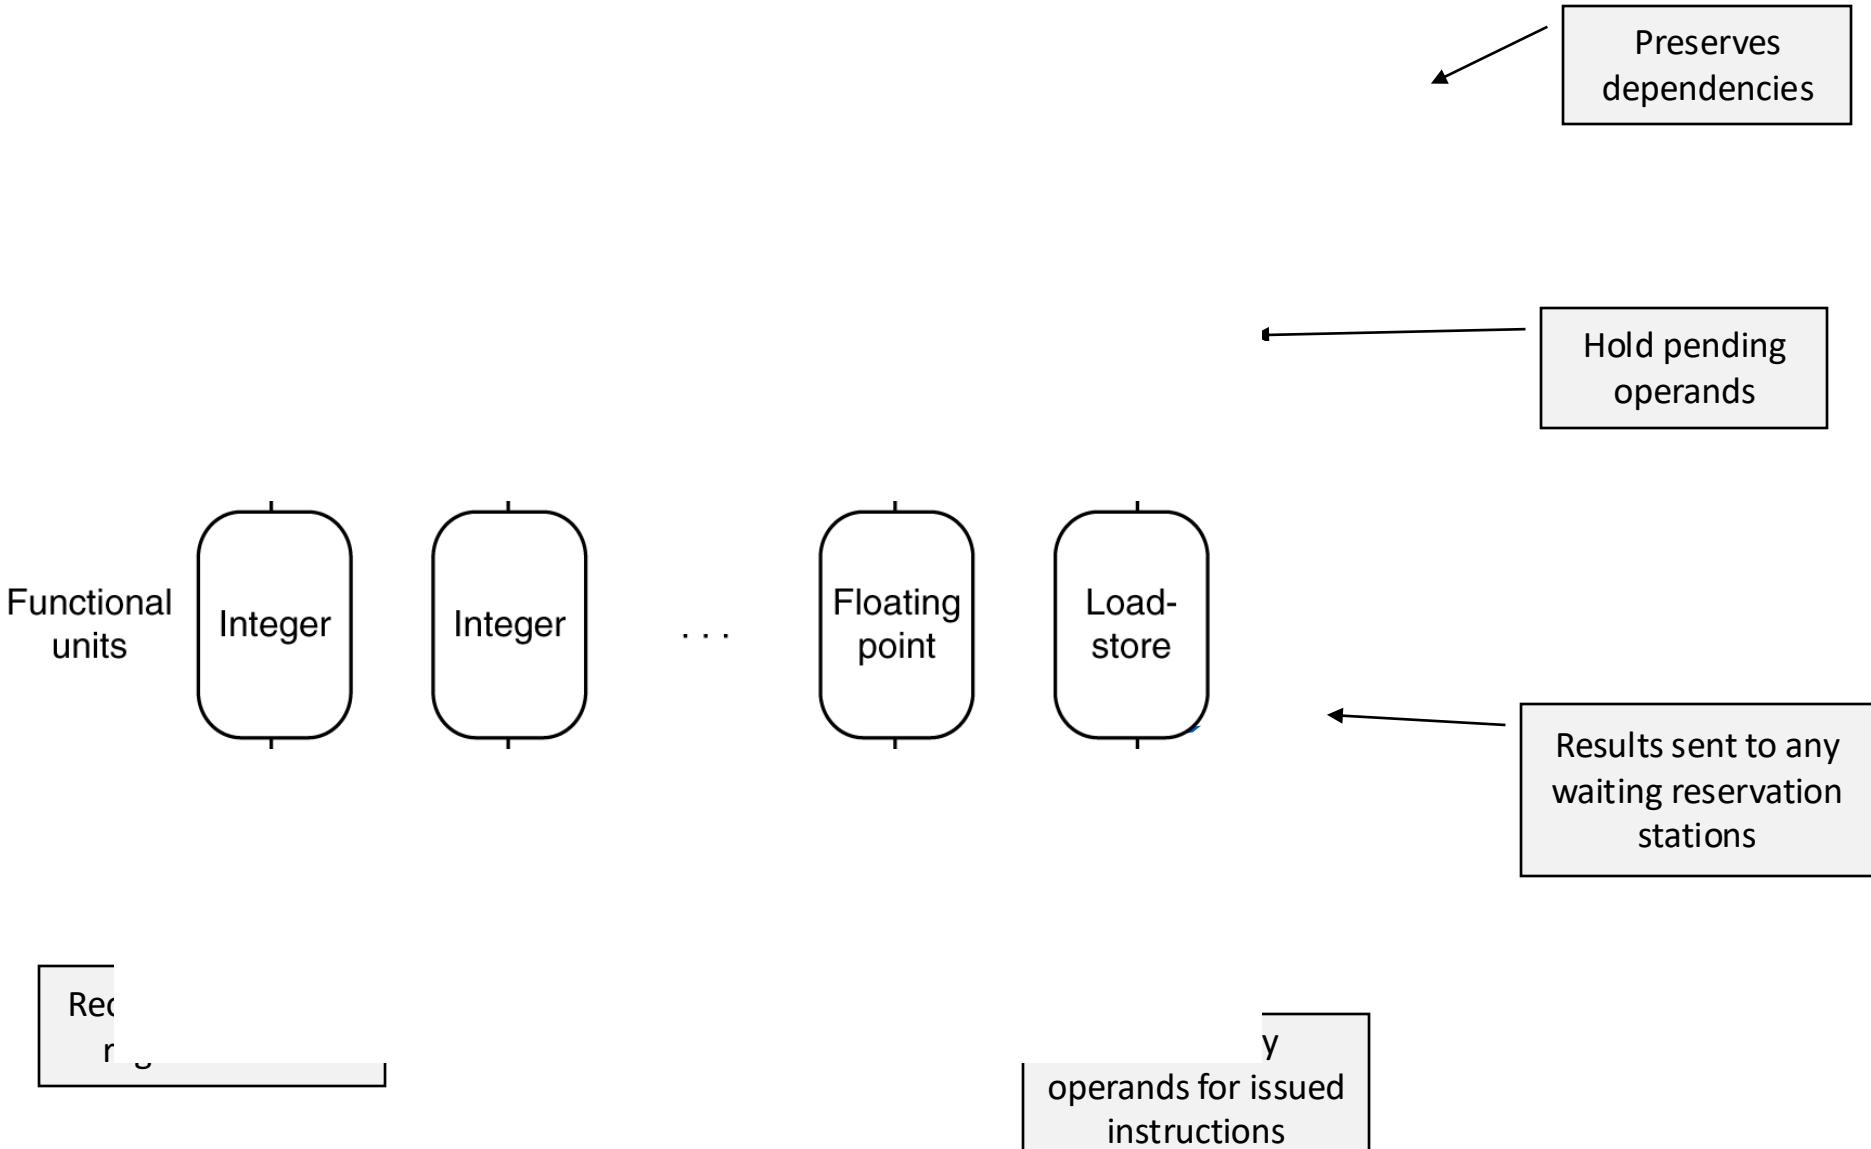
But commit result to registers in order

Example

```
ld    x31, 0(x21)
add  x1, x31, x2
sub  x23, x23, x3
andi x5, x23, 20
```

Can start **sub** while **add** is waiting for **ld**

# Dynamically Scheduled CPU



# Register Renaming

Reservation stations and reorder buffer provide register renaming

When issuing instruction to a reservation station:

- If operand is available in register file or reorder buffer

- Copy to reservation station

- If operand is not yet available

- It will be provided to the reservation station by a functional unit

# Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

# Does Multiple Issue Work?

## The **BIG** Picture

Yes, but not as much as we'd like

Programs have real dependencies that limit ILP

Some dependencies are hard to eliminate

- e.g., pointer aliasing

Some parallelism is hard to expose

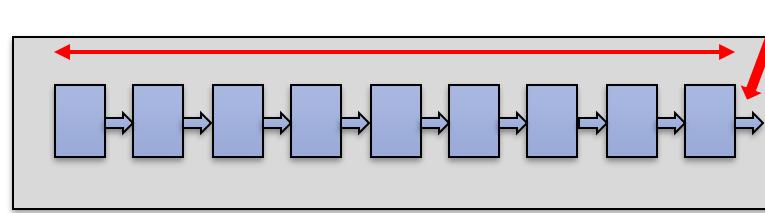
- Limited window size during instruction issue

Memory delays and limited bandwidth

- Hard to keep pipelines full

Speculation can help if done well

# Pipelining and ILP

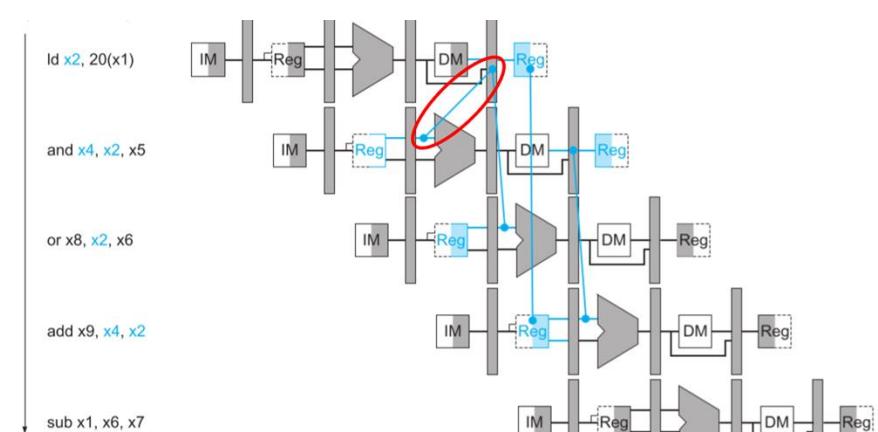
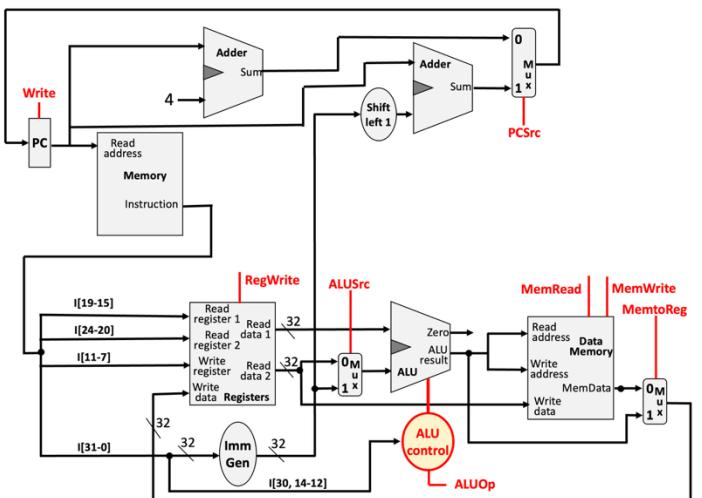
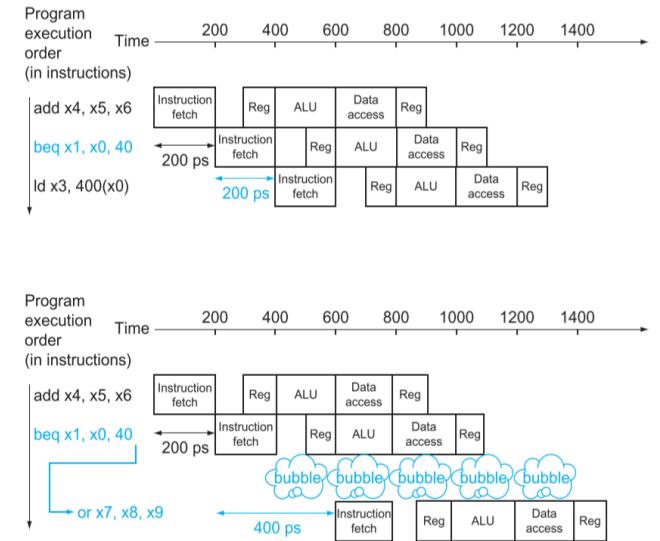


Pipelining improves instruction throughput using parallelism

More instructions completed per second

Latency for each instruction not reduced

# Hazards



control

structural

data

# Pipelining and ILP

Multiple issue and dynamic scheduling (ILP)

Dependencies limit achievable parallelism

Complexity leads to the power wall