

# Topic V26

Floating-Point  
Matrix Multiplication  
Reading: (Section 3.5)

# Floating Point (example)

Matrix multiplication of  $C = A * B$  (Double Precision Matrix Multiply)

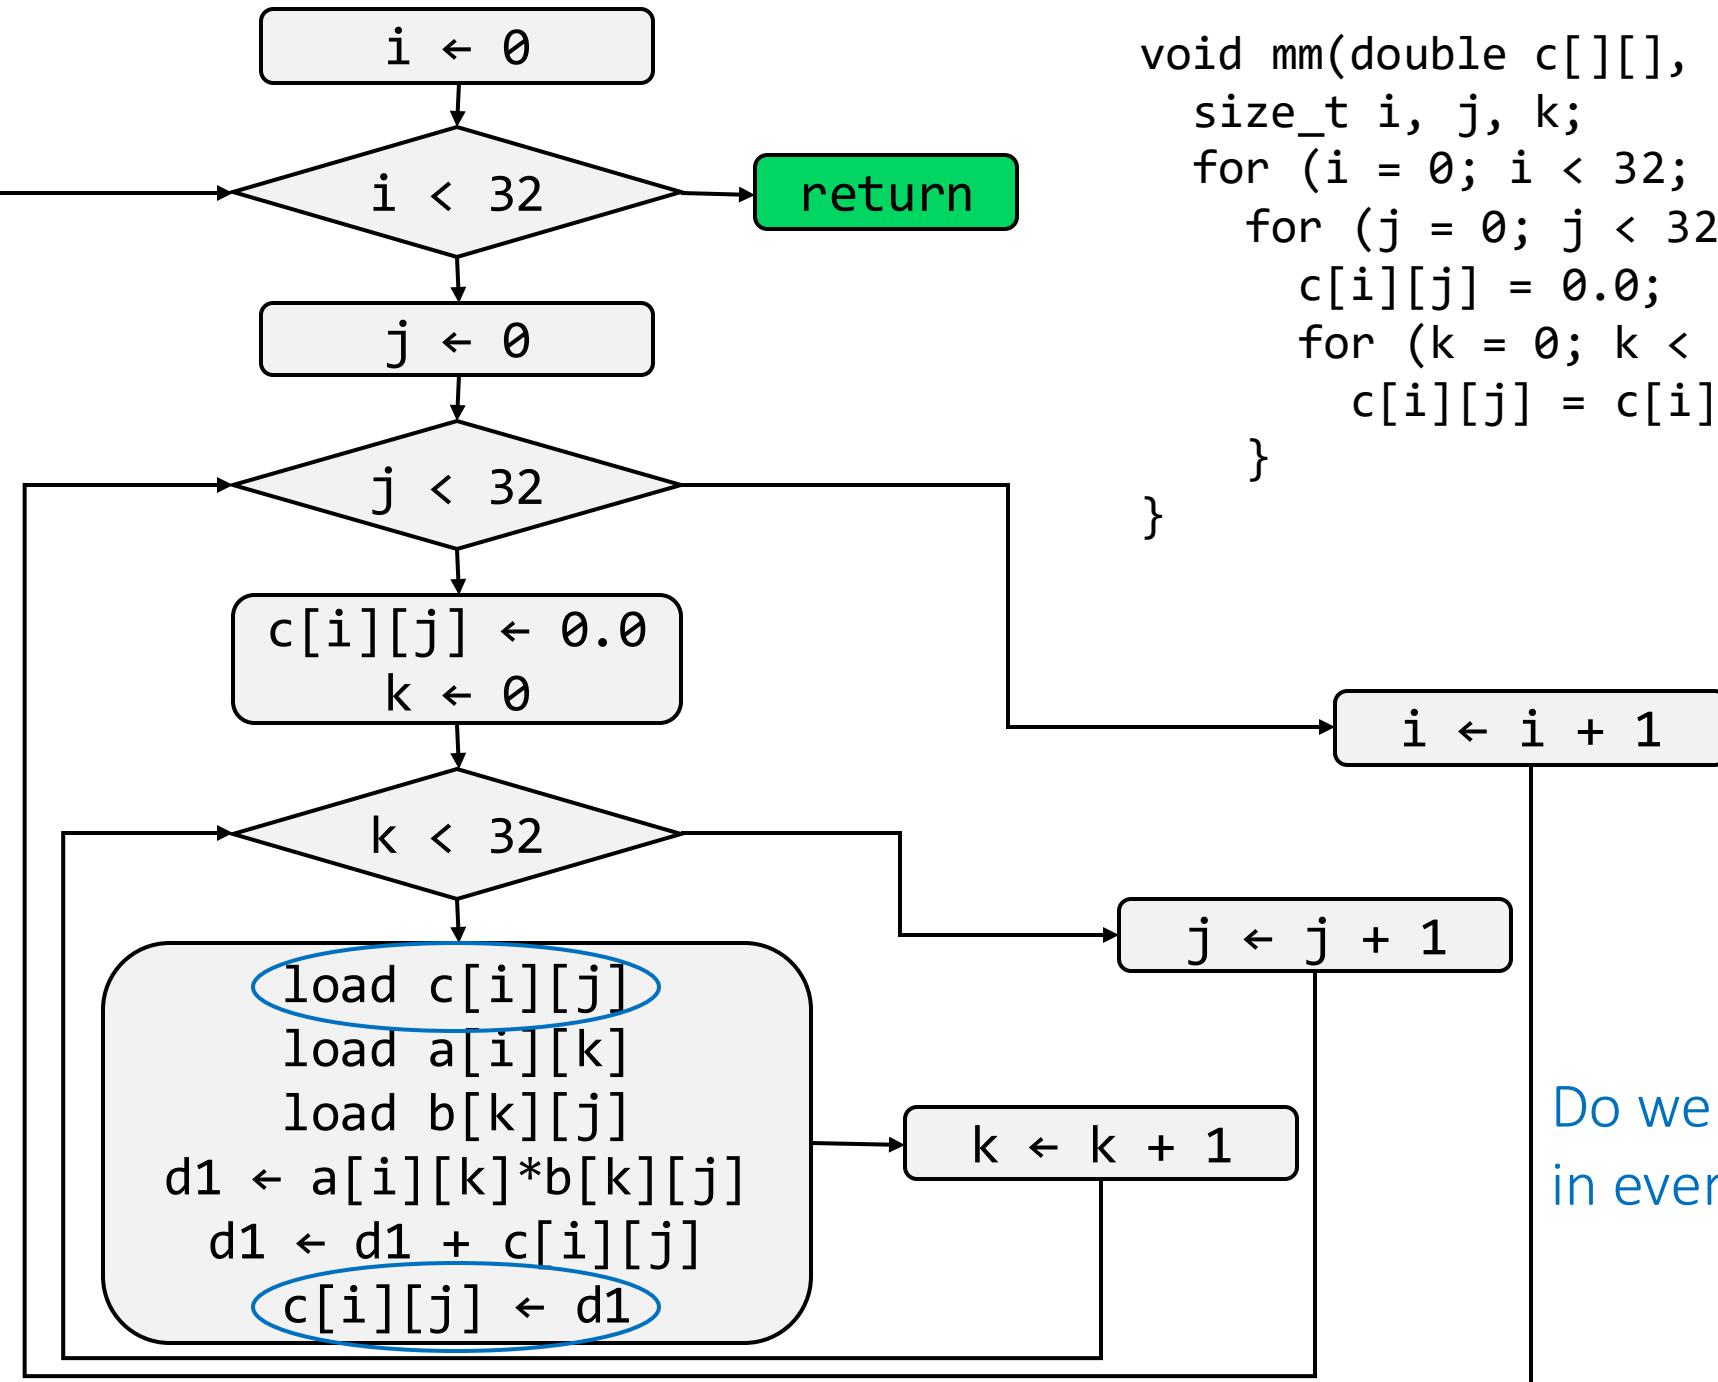
Assume C, A, and B are all square matrices with 32 elements in each dimension

This is an abuse of notation in the textbook; in C  
you cannot skip the size of previous dimensions

```
void mm(double c[][], double a[][], double b[][]){    Parameter Passing
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1){
            c[i][j] = 0.0;
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
}
```

base of  $c[] \leftrightarrow a_0$   
base of  $a[] \leftrightarrow a_1$   
base of  $b[] \leftrightarrow a_2$

Assumption  
 $i \leftrightarrow s_0$   
 $j \leftrightarrow s_1$   
 $k \leftrightarrow s_2$

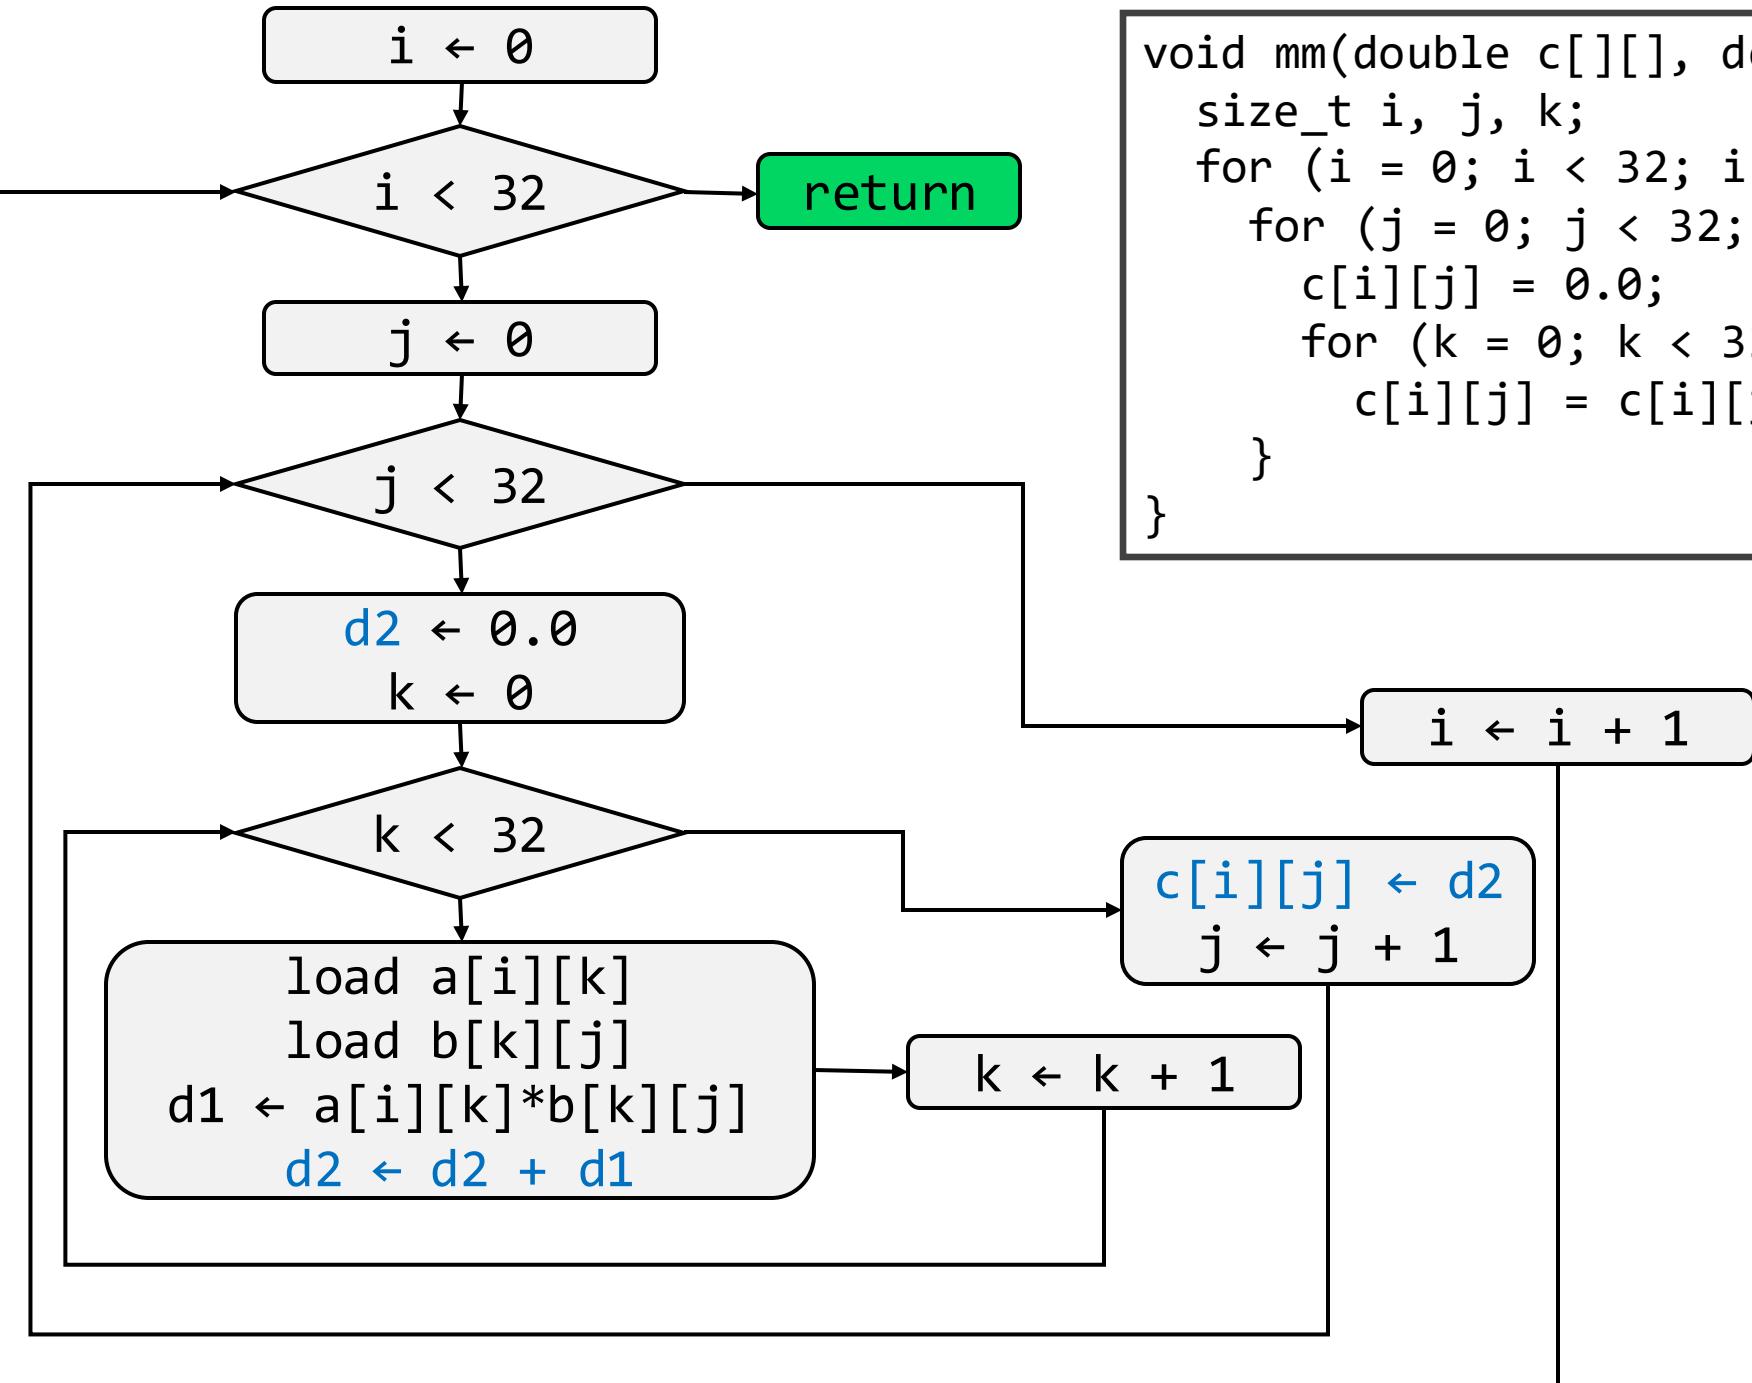


```

void mm(double c[][], double a[][], double b[][]){
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1){
            c[i][j] = 0.0;
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
    }
}

```

Do we need to load and store  $c[i][j]$  in every iteration of loop  $k$ ?



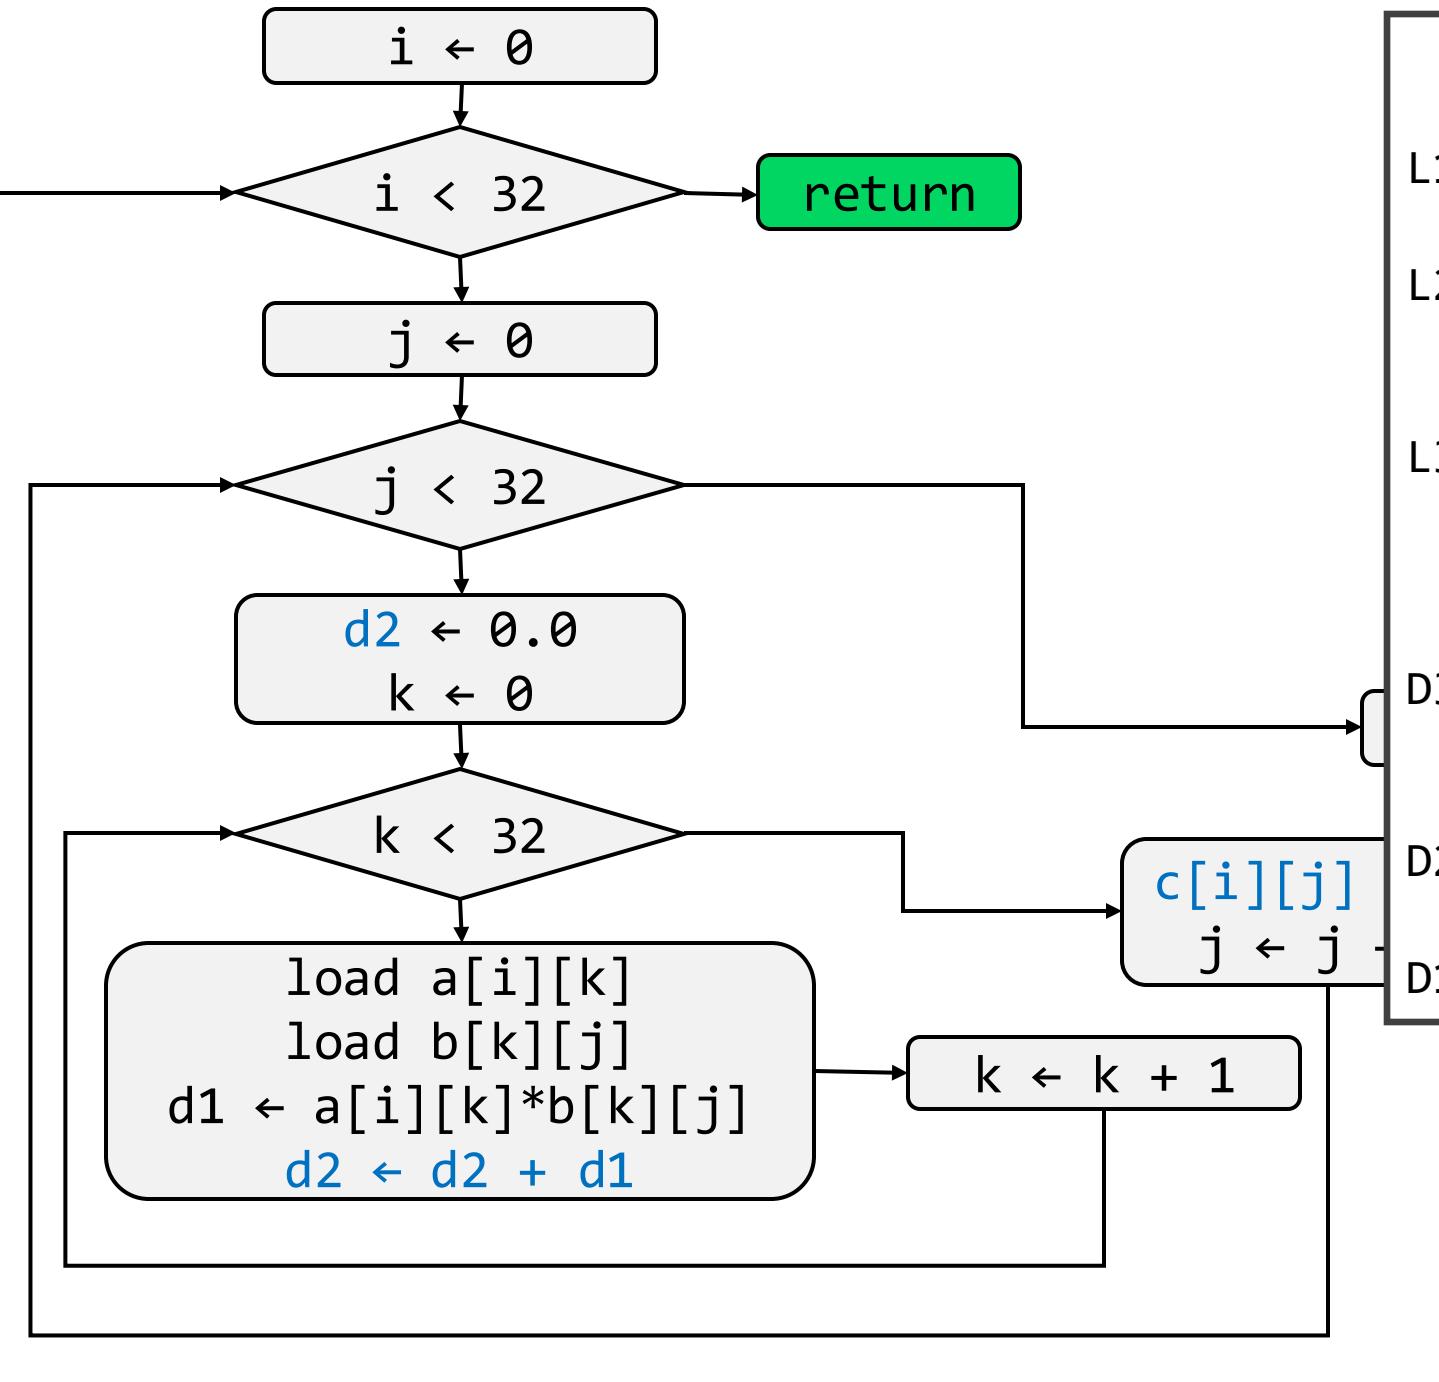
```

void mm(double c[][], double a[][], double b[][]){
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1){
            c[i][j] = 0.0;
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
}

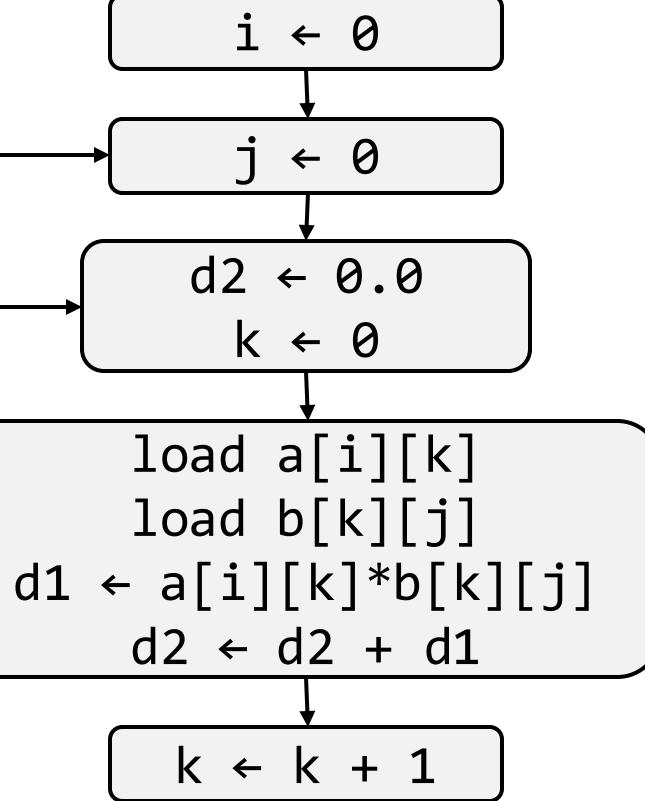
```

Parameter Passing  
base of  $c[]$   $\leftrightarrow a0$   
base of  $a[]$   $\leftrightarrow a1$   
base of  $b[]$   $\leftrightarrow a2$

Assumption  
 $i \leftrightarrow s0$   
 $j \leftrightarrow s1$   
 $k \leftrightarrow s2$



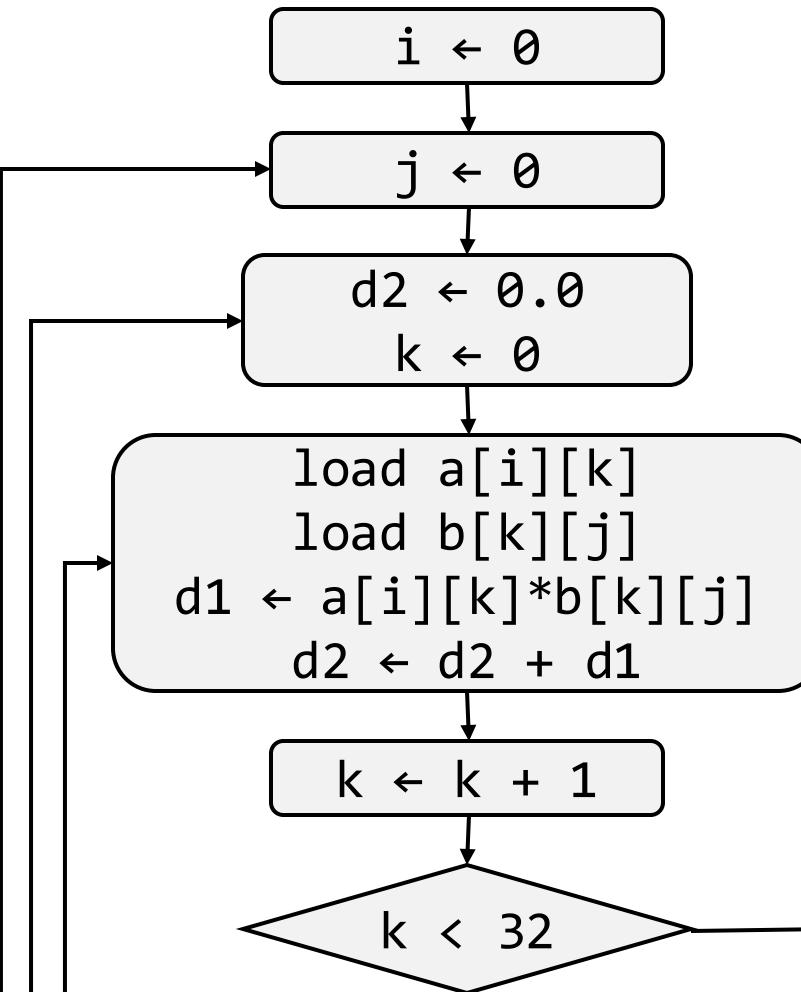
	li	t1, 32	# t1 ← 32
	li	s0, 0	# i ← 0
L1:	bgeu	s0, t1, D1	
	li	s1, 0	# j ← 0
L2:	bgeu	s1, t1, D2	
	<b>f4</b>	<b>← 0.0</b>	
	li	s2, 0	# k ← 0
L3:	bgeu	s2, t1, D3	
	<b>&lt;loop body&gt;</b>		
	addi	s2, s2, 1	# k ← k + 1
	jal	zero, L3	
D3:	<b>c[i][j]</b>	<b>← f4</b>	
	addi	s1, s1, 1	# j ← j + 1
	jal	zero, L2	
D2:	addi	s0, s0, 1	# i ← i + 1
	jal	zero, L1	
D1:	...		



```

void mm(double c[][], double a[][], double b[][]){
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1){
            c[i][j] = 0.0;
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
}

```



	li	t1, 32	# t1 ← 32
	li	s0, 0	# i ← 0
L1:	li	s1, 0	# j ← 0
L2:	<b>f4</b>	← 0.0	
	li	s2, 0	# k ← 0
L3:	<b>&lt;loop body&gt;</b>		
	addi	s2, s2, 1	# k ← k + 1
	bltu	s2, t1, L3	
	<b>c[i][j]</b>	← <b>f4</b>	
	addi	s1, s1, 1	# j ← j + 1
	bltu	s1, t1, L2	
	addi	s0, s0, 1	# i ← i + 1
	bltu	s0, t1, L1	
	...		

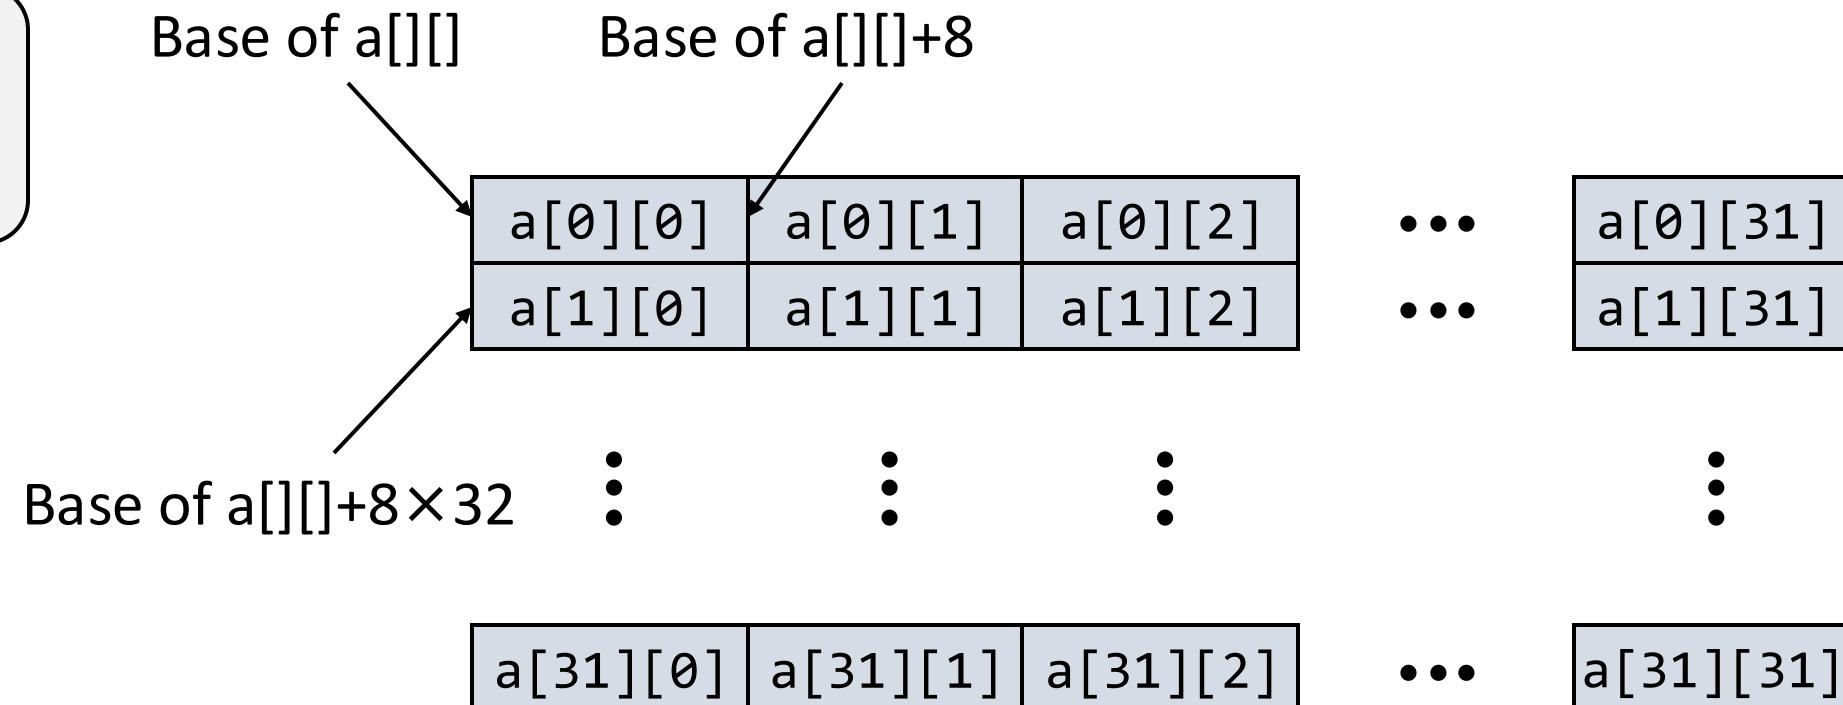
The assembly-like code corresponds to the control flow graph. It initializes  $t1$  to 32,  $i$  to 0, and  $j$  to 0. The loop body begins at label L3, where  $k$  is initialized to 0. Inside the loop,  $s2$  is updated to  $s2 + t1$ . The value of  $c[i][j]$  is set to  $f4$ . Then,  $j$  is incremented by 1, and the condition  $s1 < t1$  is checked against  $t1$  to determine if the inner loop should continue. If  $j < 32$ , the loop continues. After the inner loop exits,  $i$  is incremented by 1, and the condition  $s0 < t1$  is checked against  $t1$  to determine if the outer loop should continue. If  $i < 32$ , the loop continues. The code ends with a return statement.

# The Loop Body

How do we load  $a[i][k]$  into a floating point register?

First we have to consider how a 2-dimensional matrix of doubles is stored in memory

```
load a[i][k]
load b[k][j]
d1 <- a[i][k]*b[k][j]
d2 <- d2 + d1
```



In general, the address of  $a[i][k]$  is given by:

$$\text{address}(a[i][k]) = \text{base of } a[][] + (i \times 32 + k) \times 8$$

# The Loop Body (Cont.)

In general, the address of  $a[i][k]$  is given by:

$$\text{address}(a[i][k]) = \text{base of } a[][] + (i \times 32 + k) \times 8$$

```
load a[i][k]
load b[k][j]
d1 <- a[i][k]*b[k][j]
d2 <- d2 + d1
```

RISC-V assembly for load  $a[i][k]$

```
L3:    slli    t2, s0, 5      # t2 <- 32 * i
        add     t2, t2, s2      # t2 <- 32 * i + k
        slli    t2, t2, 3      # t2 <- (32 * i + k) * 8
        add     t2, a1, t2      # t2 <- Addr(a[i][k])
        fld     f1, 0(t2)       # f1 <- a[i][k]
```

Write the code to load  $b[k][j]$  in  $f2$

RISC-V assembly for load  $b[k][j]$

```
slli    t2, s2, 5      # t2 <- 32 * k
add     t2, t2, s1      # t2 <- 32 * k + j
slli    t2, t2, 3      # t2 <- (32 * k + j) * 8
add     t2, a2, t2      # t2 <- Addr(b[k][j])
fld     f2, 0(t2)       # f2 <- b[k][j]
```

# The Loop Body (Cont.)

Once we have loaded  $a[i][k]$  into  $f1$  and  $b[k][j]$  into  $f2$  we can proceed to perform the multiply and the add

```
load a[i][k]
load b[k][j]
d1 <- a[i][k]*b[k][j]
d2 <- d2 + d1
```

RISC-V assembly for FP multiply and add

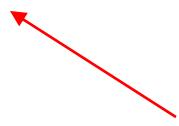
fmul.d	$f1, f1, f2$	# $f1 \leftarrow a[i][k] \times b[k][j]$
fadd.d	$f0, f0, f1$	# $c[i][j] + a[i][k] \times b[k][j]$

# Initializing and Storing f4

How can we initialize f4?

RISC-V assembly to initialize f4

fcvt.d.w f4, zero



Converts the integer zero into a floating-point zero

How can we store f4 in a[i][j]?

RISC-V assembly to store f4 in c[i][j]

```
slli t2, s0, 5      # t2 <= 32 * i  
add  t2, t2, s1     # t2 <= 32 * i + j  
slli t2, t2, 3      # <= (32 * i + j) * 8  
add  t2, a0, t2     # t2 <= Addr(c[i][j])  
fsd  f4, 0(t2)      # c[i][j] <= f4
```

```
li    t1, 32          # t1 <= 32  
li    s0, 0           # i <= 0  
L1: li    s1, 0          # j <= 0  
L2: f4 <= 0.0  
li    s2, 0           # k <= 0  
L3: <loop body>  
    addi s2, s2, 1      # k <= k + 1  
    bltu s2, t1, L3  
    c[i][j] <= f4  
    addi s1, s1, 1      # j <= j + 1  
    bltu s1, t1, L2  
    addi s0, s0, 1      # i <= i + 1  
    bltu s0, t1, L1  
...
```

```

mm:          ...
            li           t1, 32      # t1 < 32
            li           s0, 0       # i < 0
L1:          li           s1, 0       # j < 0
L2:          fcvt.d.w   f4, zero
            li           s2, 0       # k < 0
L3:          slli         t2, s0, 5  # t2 < 32 * i
            add          t2, t2, s2  # t2 < 32 * i + k
            slli         t2, t2, 3  # t2 < (32 * i + k) * 8
            add          t2, a1, t2  # t2 < Addr(a[i][k])
            fld          f1, 0(t2)  # f1 < a[i][k]
            slli         t2, s2, 5  # t2 < 32 * k
            add          t2, t2, s1  # t2 < 32 * k + j
            slli         t2, t2, 3  # t2 < (32 * k + j) * 8
            add          t2, a2, t2  # t2 < Addr(b[k][j])
            fld          f2, 0(t2)  # f2 < b[k][j]
            fmul.d      f1, f1, f2  # f1 < a[i][k] * b[k][j]
            fadd.d      f0, f0, f1  # c[i][j] + a[i][k] * b[k][j]
            addi         s2, s2, 1  # k < k + 1
            bltu        s2, t1, L3
            slli         t2, s0, 5  # t2 < 32 * i
            add          t2, t2, s1  # t2 < 32 * i + j
            slli         t2, t2, 3  # < (32 * i + j) * 8
            add          t2, a0, t2  # t2 < Addr(c[i][j])
            fsd          f4, 0(t2)  # c[i][j] < f4
            addi         s1, s1, 1  # j < j + 1
            bltu        s1, t1, L2
            addi         s0, s0, 1  # i < i + 1
            bltu        s0, t1, L1
            ...

```

load a[i][k] in f1

load b[k][j] in f2

store f4 in c[i][j]

```

mm:          ...
            li           t1, 32      # t1 < 32
            li           s0, 0       # i < 0
L1:          li           s1, 0       # j < 0
L2:          fcvt.d.w   f4, zero
            li           s2, 0       # k < 0
L3:          slli         t2, s0, 5    # t2 < 32 * i
            add          t2, t2, s2  # t2 < 32 * i + k
            slli         t2, t2, 3    # t2 < (32 * i + k) * 8
            add          t2, a1, t2  # t2 < Addr(a[i][k])
            fld          f1, 0(t2)  # f1 < a[i][k]
            slli         t2, s2, 5    # t2 < 32 * k
            add          t2, t2, s1  # t2 < 32 * k + j
            slli         t2, t2, 3    # t2 < (32 * k + j) * 8
            add          t2, a2, t2  # t2 < Addr(b[k][j])
            fld          f2, 0(t2)  # f2 < b[k][j]
            fmul.d      f1, f1, f2  # f1 < a[i][k] * b[k][j]
            fadd.d      f0, f0, f1  # c[i][j] + a[i][k] * b[k][j]
            addi         s2, s2, 1    # k < k + 1
            bltu        s2, t1, L3
            slli         t2, s0, 5    # t2 < 32 * i
            add          t2, t2, s1  # t2 < 32 * i + j
            slli         t2, t2, 3    # < (32 * i + j) * 8
            add          t2, a0, t2  # t2 < Addr(c[i][j])
            fsd          f4, 0(t2)  # c[i][j] < f4
            addi         s1, s1, 1    # j < j + 1
            bltu        s1, t1, L2
            addi         s0, s0, 1    # i < i + 1
            bltu        s0, t1, L1
            ...

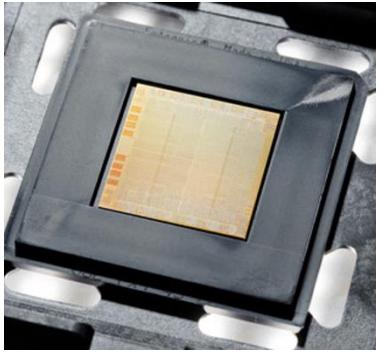
```

load a[i][k] in f1

load b[k][j] in f2

store f4 in c[i][j]

# Matrix Multiplication Accelerators



Hot Chips 32 | Server Processors

May 23, 2020 • David Schor • 0 Comments

THE MEMORY AREA NETWORK AT THE HEART OF  
IBM'S POWER10

September 3, 2020 Timothy Prickett Morgan

# IBM Releases Power ISA v3.1; To Present POWER10 At Hot Chips 32

## IBM BRINGS AN ARCHITECTURE GUN TO A CHIP KNIFE FIGHT

August 18, 2020 Timothy Prickett Morgan

BLOG - IBM Announces POWER10 Processor  
Technology

August 21st, 2020

IBM Reveals Next-Generation IBM POWER10 Processor  
New CPU co-optimized for Red Hat OpenShift for enterprise hybrid cloud

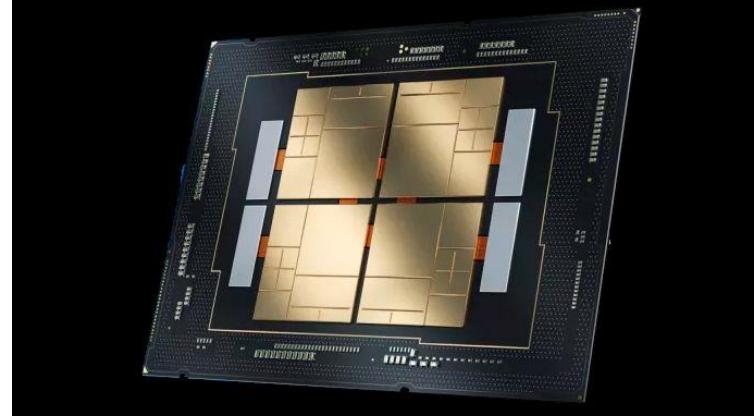
Architectures

# The x86 Advanced Matrix Extension (AMX) Brings Matrix Operations; To Debut with Sapphire Rapids

June 29, 2020

David Schor

Advanced Matrix Extension (AMX), AI, Intel, matrices, Sapphire Rapids, x86



*Intel Demos Sapphire Rapids Hardware Accelerator Blocks In Action At Innovation 2022*

by [Ryan Smith](#)

on September 28, 2022 7:20 PM EST



## Nvidia, Qualcomm Shine in MLPerf Inference; Intel's Sapphire Rapids Makes an Appearance.

By John Russell

# Introducing the Scalable Matrix Extension for the Armv9-A Architecture



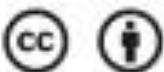
ARMv9 Scalable Matrix Extension Support Lands In Linux 5.19  
Written by Michael Larabel in Arm on 25 May 2022 at 04:40 AM EDT. 2 Comments

Journal of:

# Software: Practice and Experience

RESEARCH ARTICLE

 Open Access



## Fast matrix multiplication via compiler-only layered data reorganization and intrinsic lowering

Braedy Kuzma, Ivan Korostelev, João P. L. de Carvalho , José E. Moreira, Christopher Barton, Guido Araujo, José Nelson Amaral

First published: 14 May 2023 | <https://doi.org/10.1002/spe.3214>

# IEEE micro

The magazine for chip and silicon systems designers

VOLUME 42, NUMBER 5

SEPTEMBER/OCTOBER 2022



## Compiling for Accelerators

### THEME ARTICLE: COMPILING FOR ACCELERATORS

## Compiling for the IBM Matrix Engine for Enterprise Workloads

João P. L. de Carvalho , University of Alberta, Edmonton, AB, T6G 2S4, Canada

José E. Moreira , IBM Corporation, Yorktown Heights, NY, 10598, USA

José Nelson Amaral , University of Alberta, Edmonton, AB, T6G 2S4, Canada

ACM Transactions on  
**Architecture and  
Code Optimization**



## KERNELFARER: Replacing Native-Code Idioms with High-Performance Library Calls

JOÃO P. L. DE CARVALHO, University of Campinas (UNICAMP), Brazil  
BRAEDY KUZMA, IVAN KOROSTELEV, and JOSÉ NELSON AMARAL,  
University of Alberta, Canada  
CHRISTOPHER BARTON, IBM Corporation, Canada  
JOSÉ MOREIRA, IBM Corporation, USA  
GUIDO ARAUJO, University of Campinas (UNICAMP), Brazil

