

Topic V30

Cache Memory

Reading: (Section 5.3)

Cache Memory

A cache memory is a level of the memory hierarchy that is closer to the CPU than DRAM

Cache Organization (Data Storage)

A *cache block* contains multiple words

Example: a cache with four-word blocks



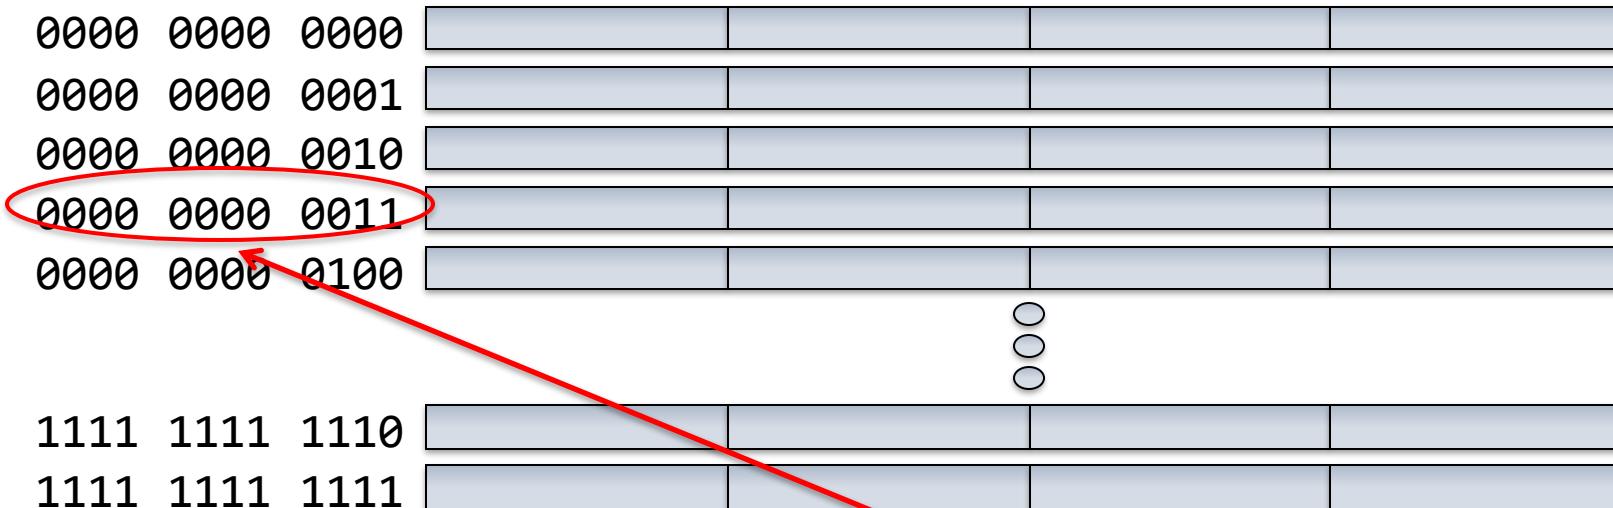
Cache Organization (Data Storage)

Each cache has multiple such blocks

Example: a cache with 4K ($4K = 4096$) four-word blocks:

We want to give a unique identifier to each block

How many bits do we need for each identifier?

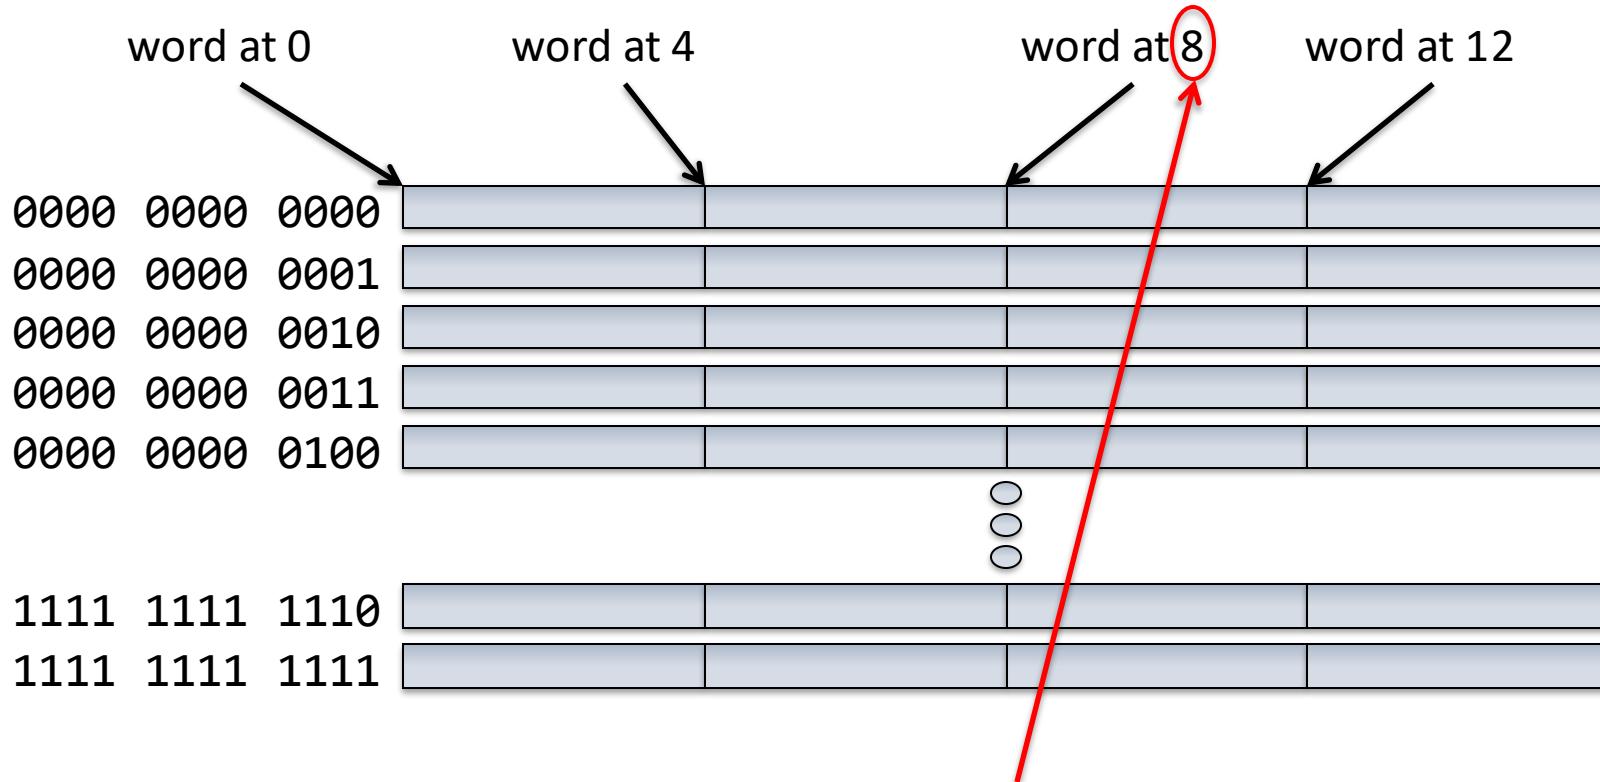


The unique ID of each cache entry is the *index* of the entry

Offset

How do we find a word inside a block?

We measure the distance, in bytes, from the beginning of the block

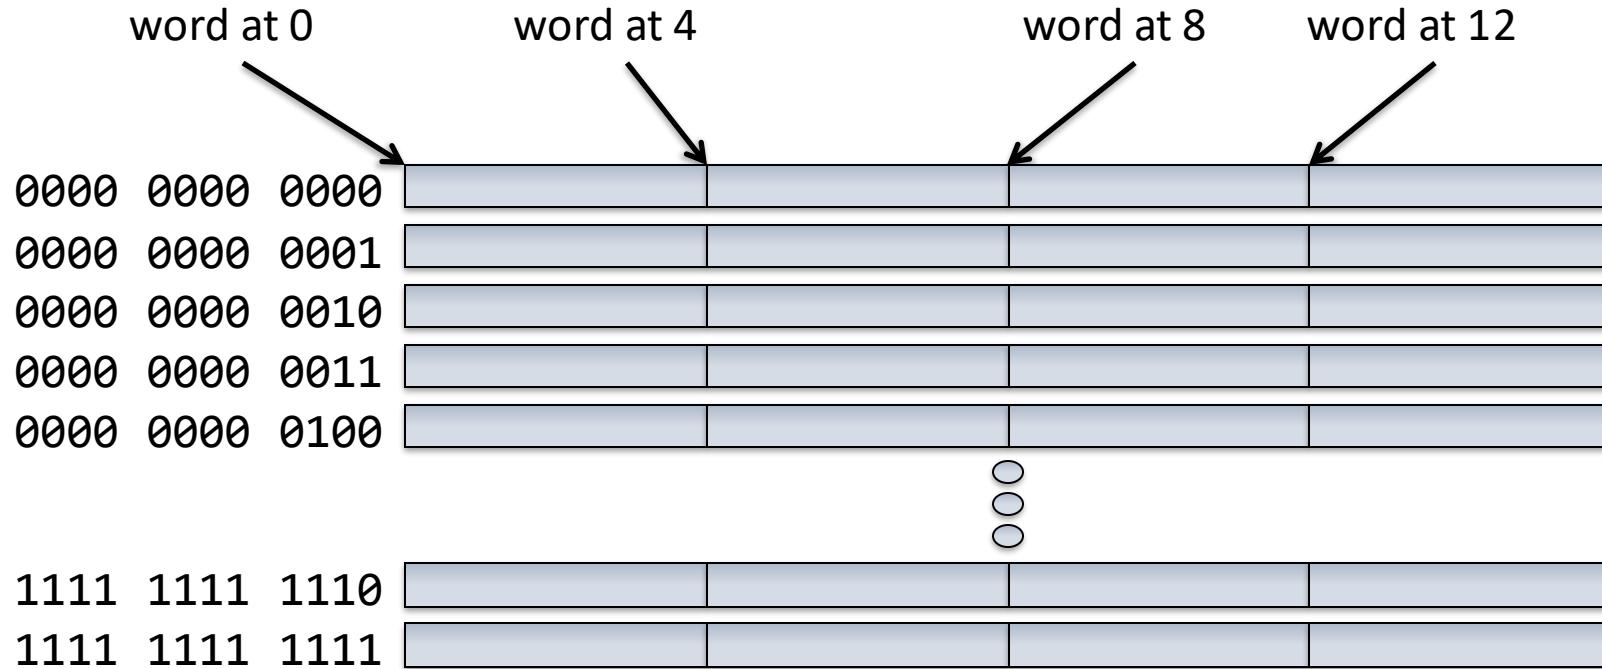


The position of a word within a block is the word *offset*

Offset

How do we find a word inside a block?

We measure the distance, in bytes, from the beginning of the block



In the example, how many bits do I need for the offset?

Each block has 16 bytes \Rightarrow need 4 bits.

Finding a word in the cache

Example:

Processor requests word at address **0xFF00 0028**

Address 0xFF00 0028

1111 1111 0000 0000 0000 0000 0010 1000

word at 0

word at 4

word at 8

word at 12

0000 0000 0000

0000 0000 0001

0000 0000 0010

0000 0000 0011

0000 0000 0100

1111 1111 1110

1111 1111 1111



...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

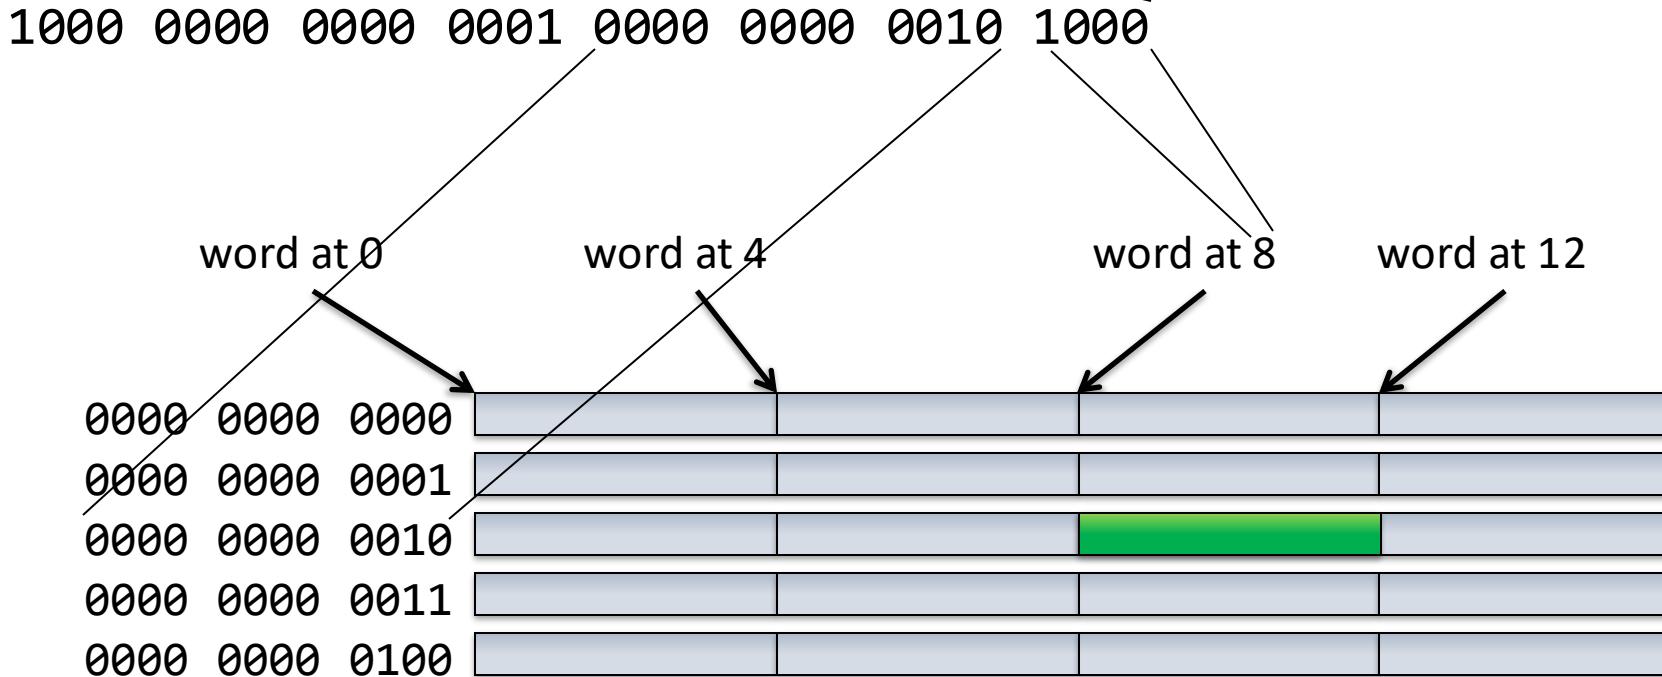
...

...

...

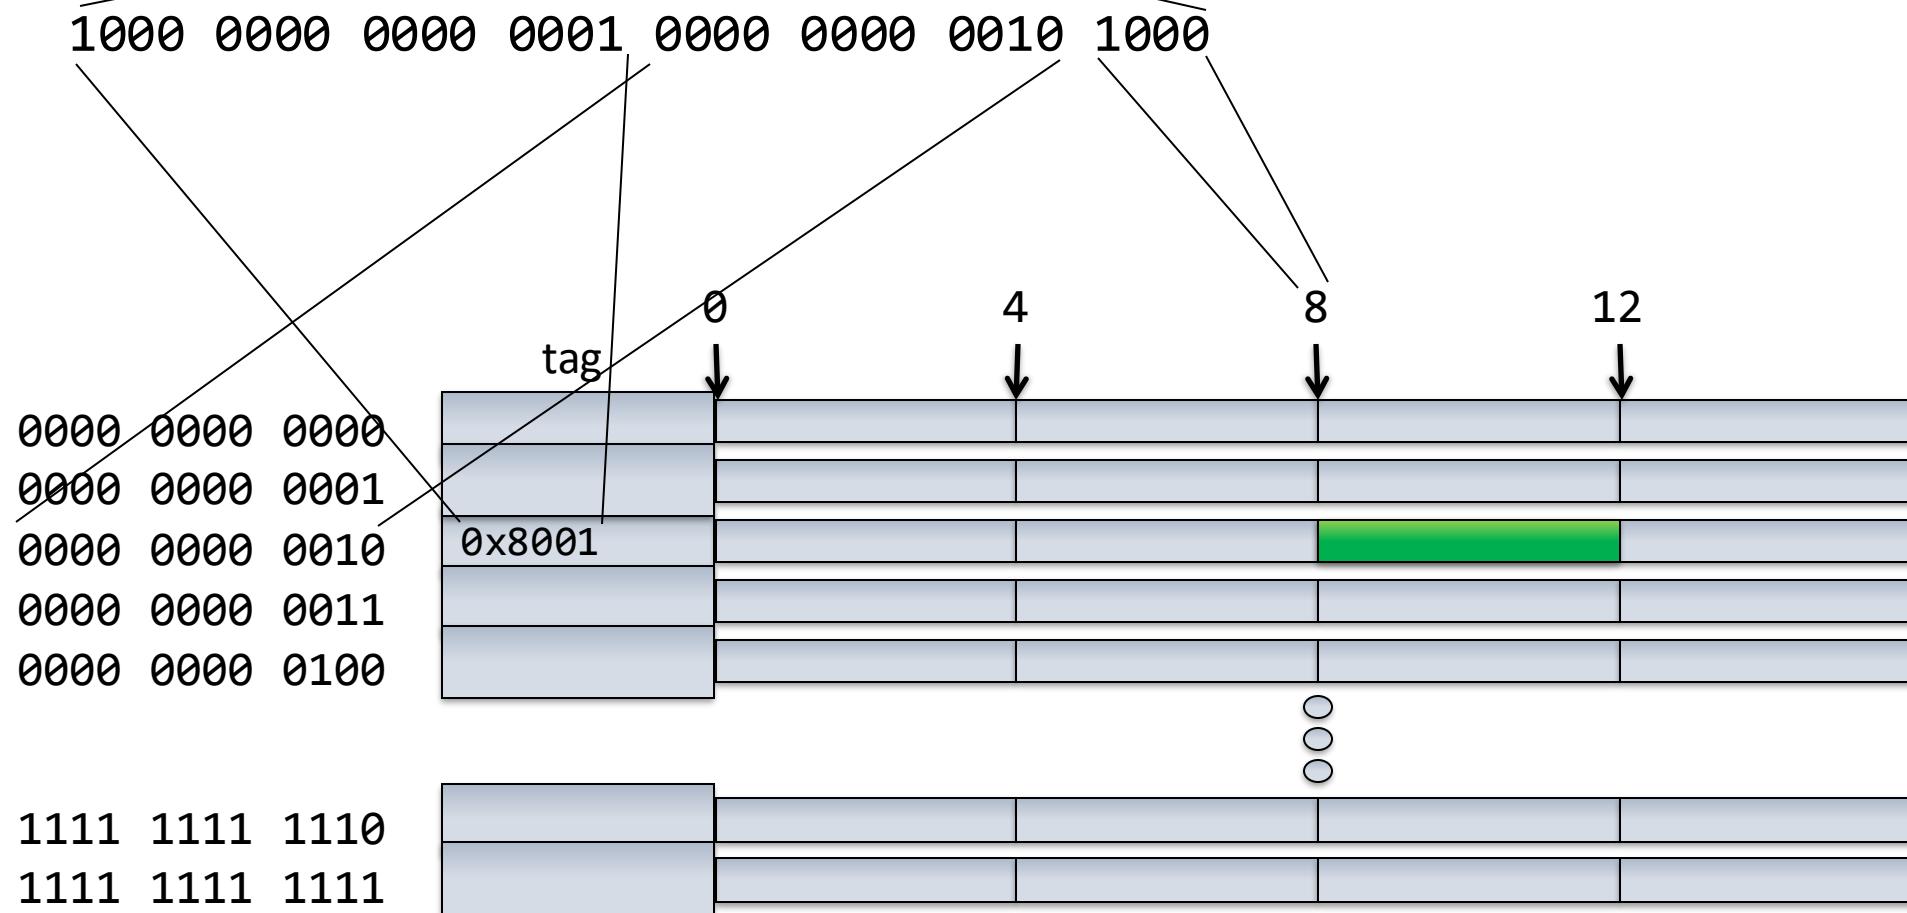
...

Now processor accesses
address **0x8001 0028**



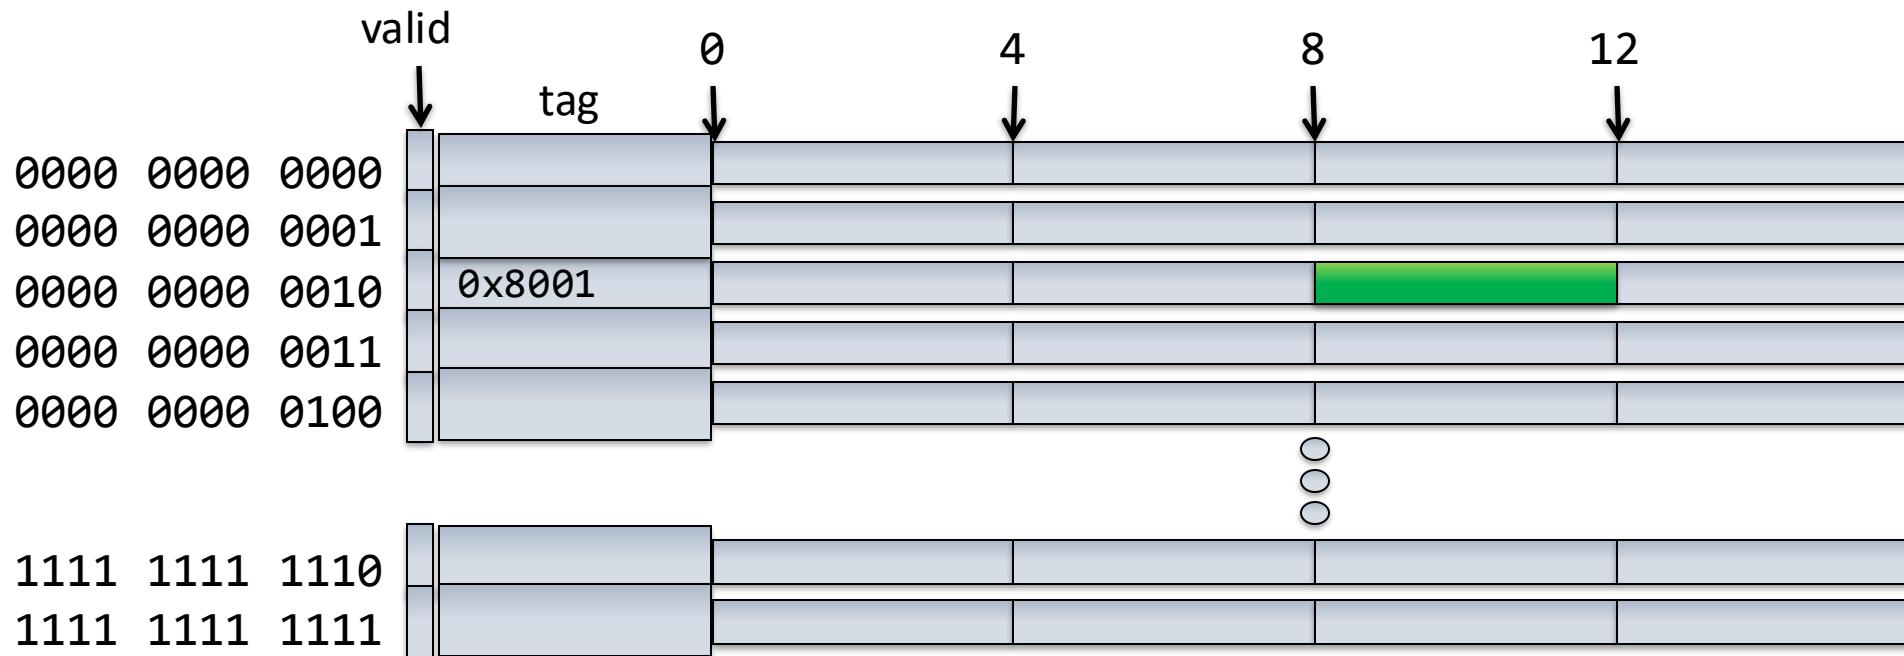
How do we know if the word in the cache is from
0xFF00 0028 or from 0x8001 0028?

When the processor accesses
address **0x8001 0028**



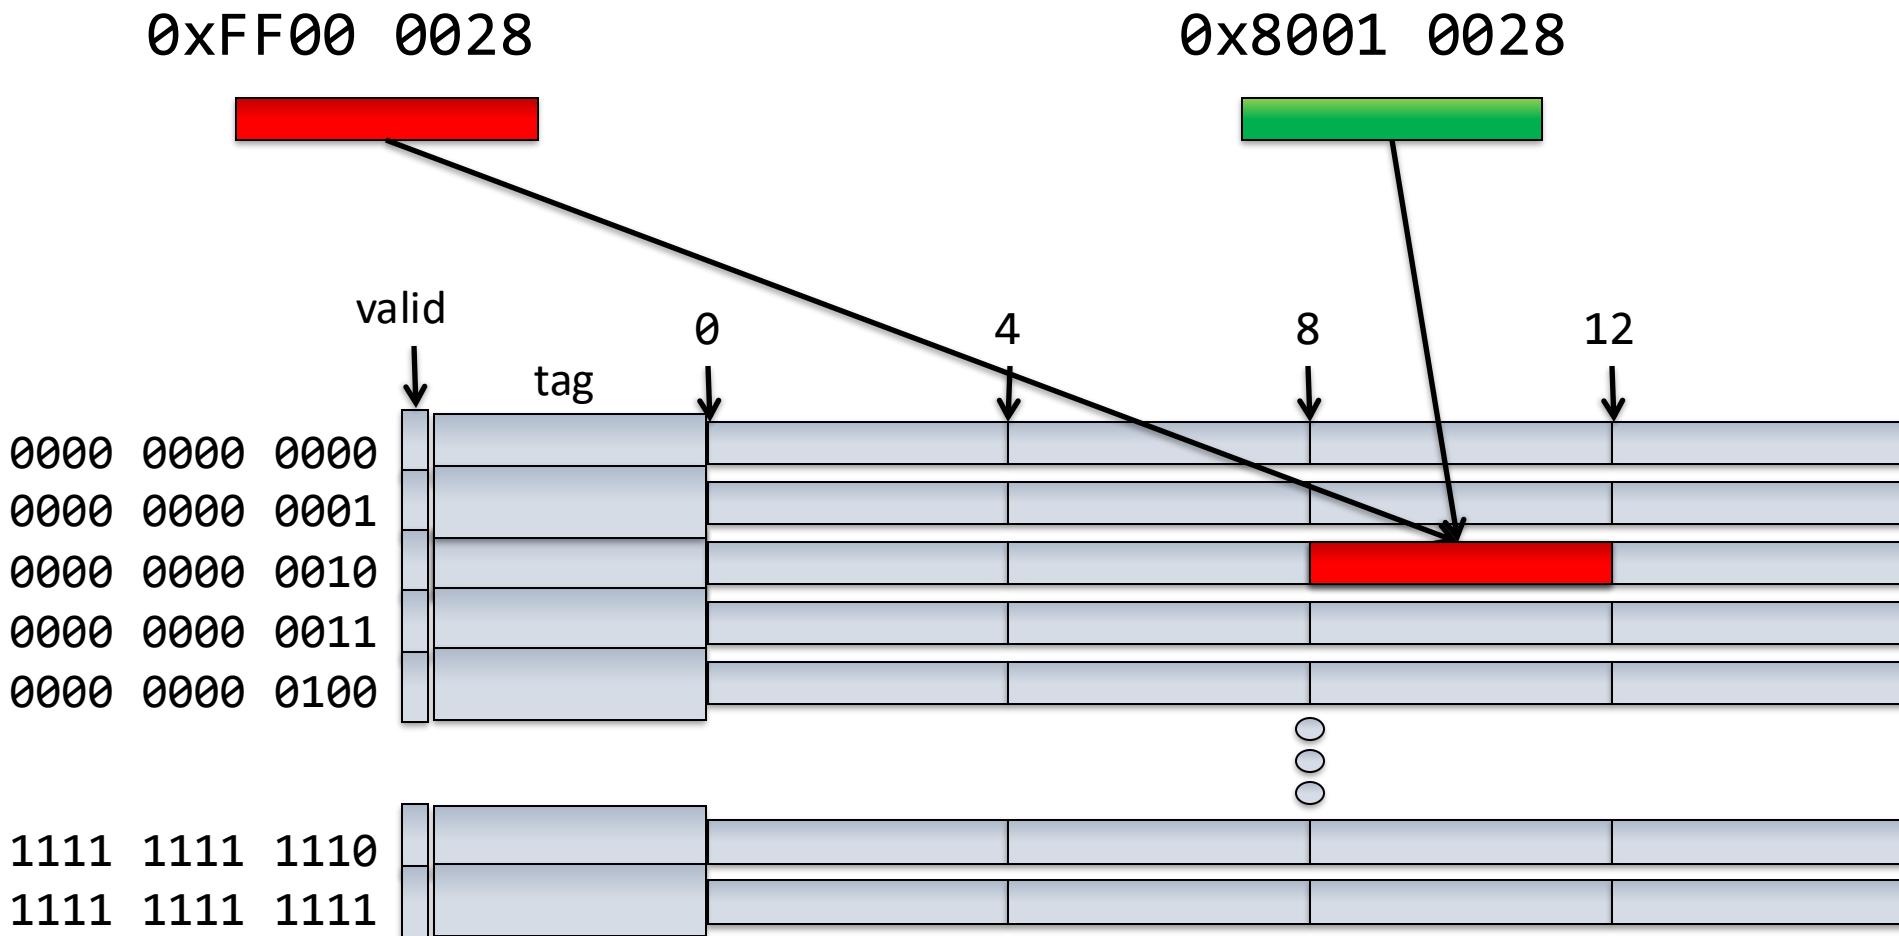
The *portion of the address stored along with each cache block to identify which address in the memory the block came from* is called **tag**

How do we know that the cache has a valid block for the requested address?



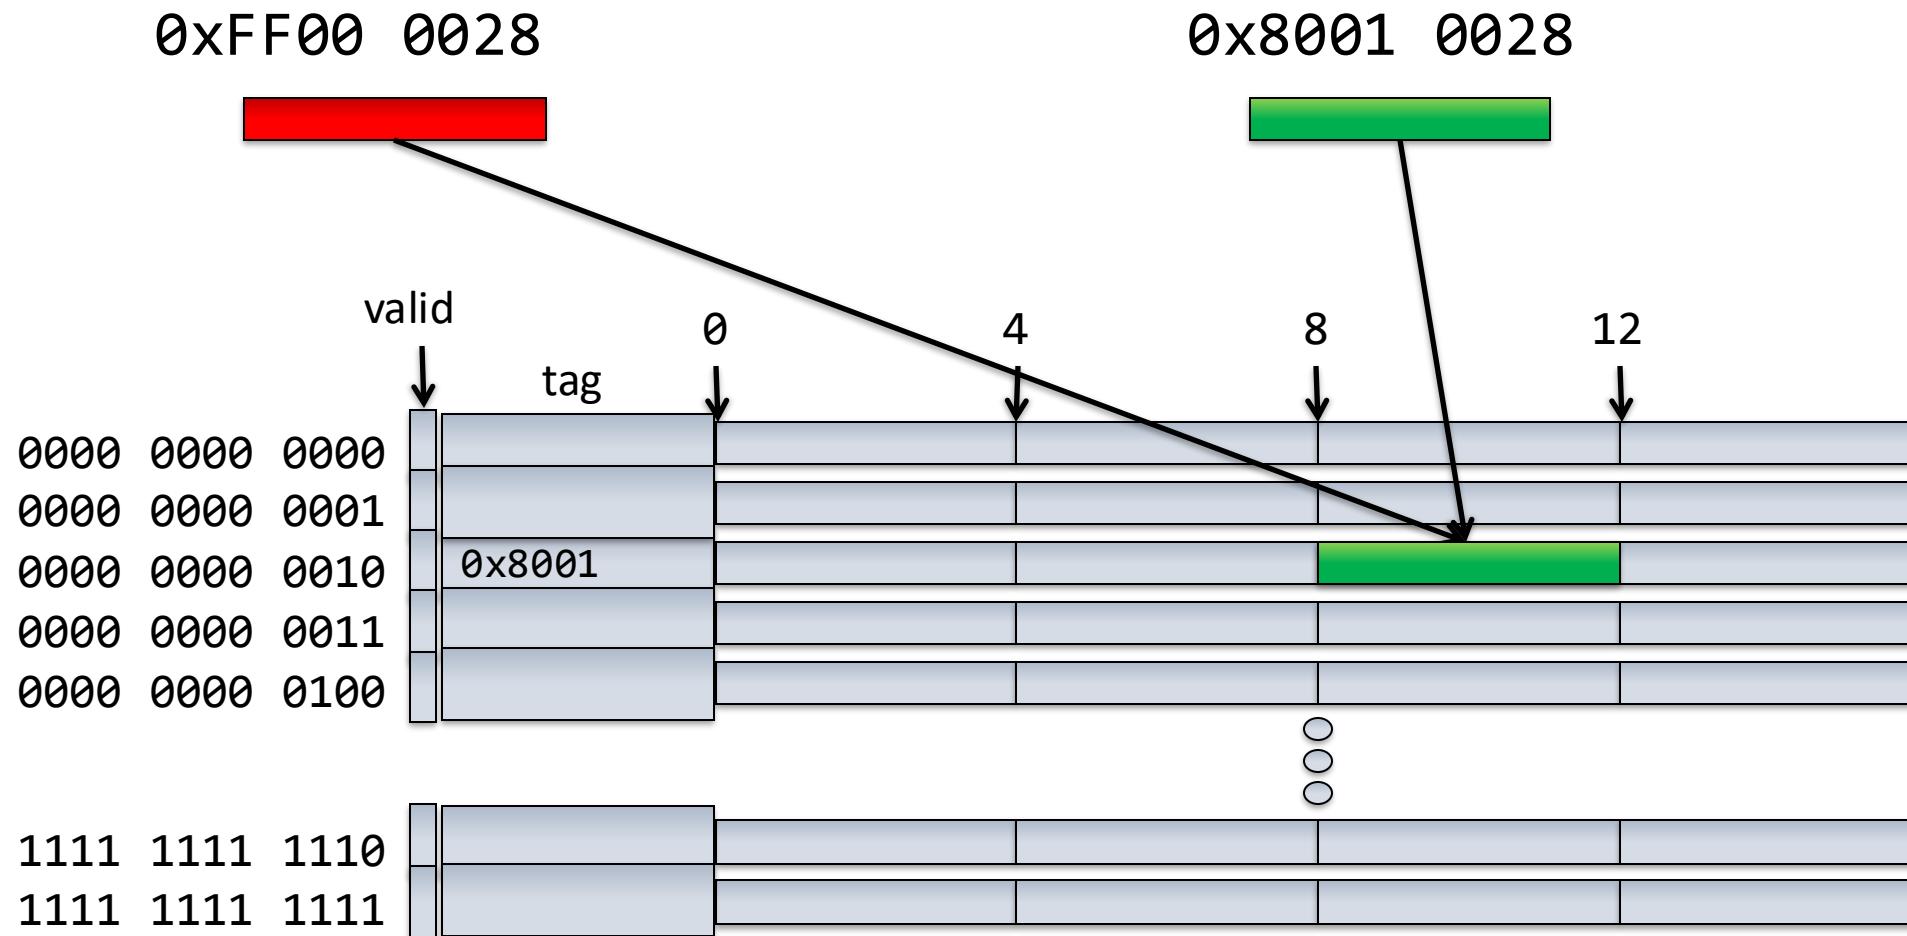
A **valid bit** associated with each block indicates if the content in the block is valid

In this design, there is only one place in the cache for each memory address



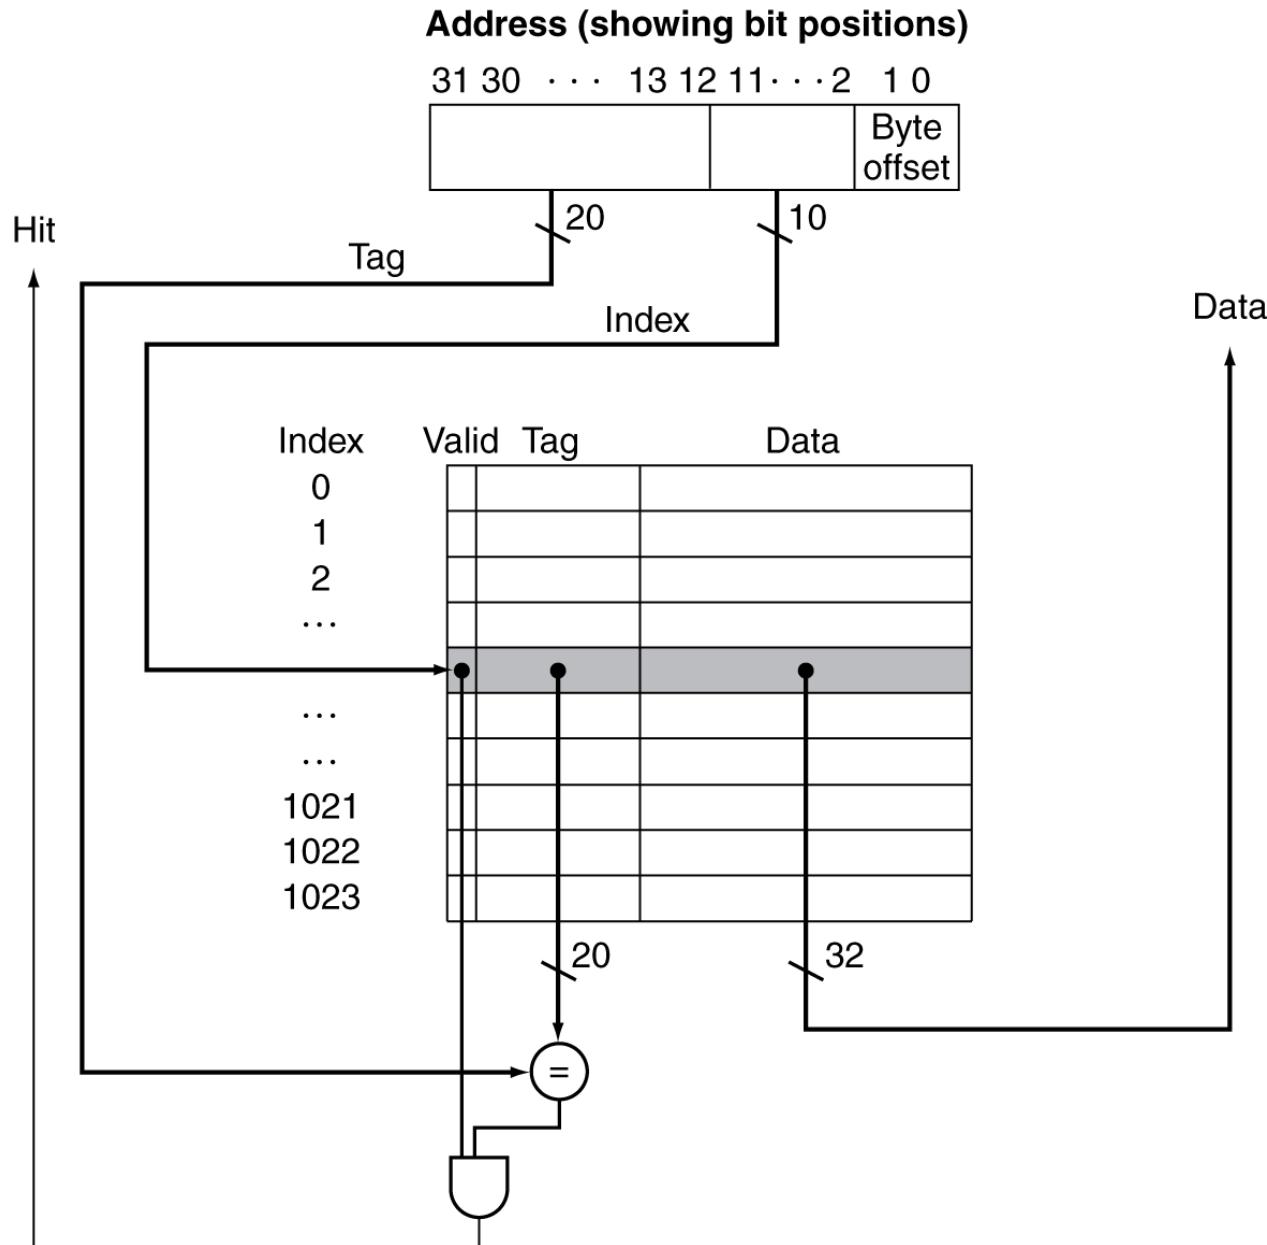
This mapping between memory addresses and cache blocks is called a
direct mapping

In this design, there is only one place in the cache for each memory address



The problem with this mapping is that there are often **conflicts** into the cache for competing memory addresses

Address Subdivision



We want to change the mapping to allow a given memory address to be mapped to more than one block in the cache

Let's group the cache blocks into sets

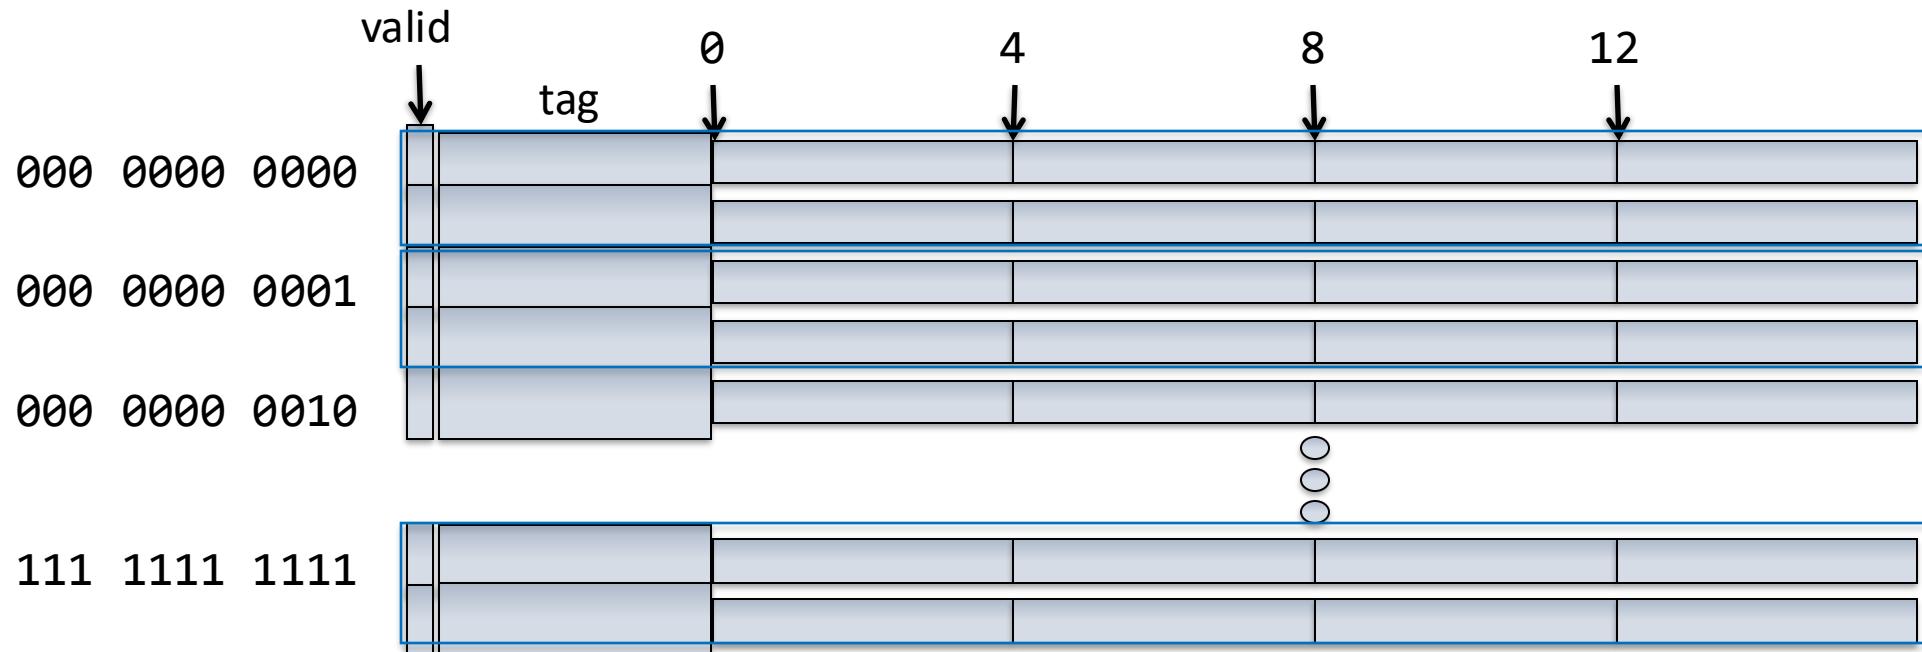
Example: lets put two blocks in each set

The cache stores 4K blocks

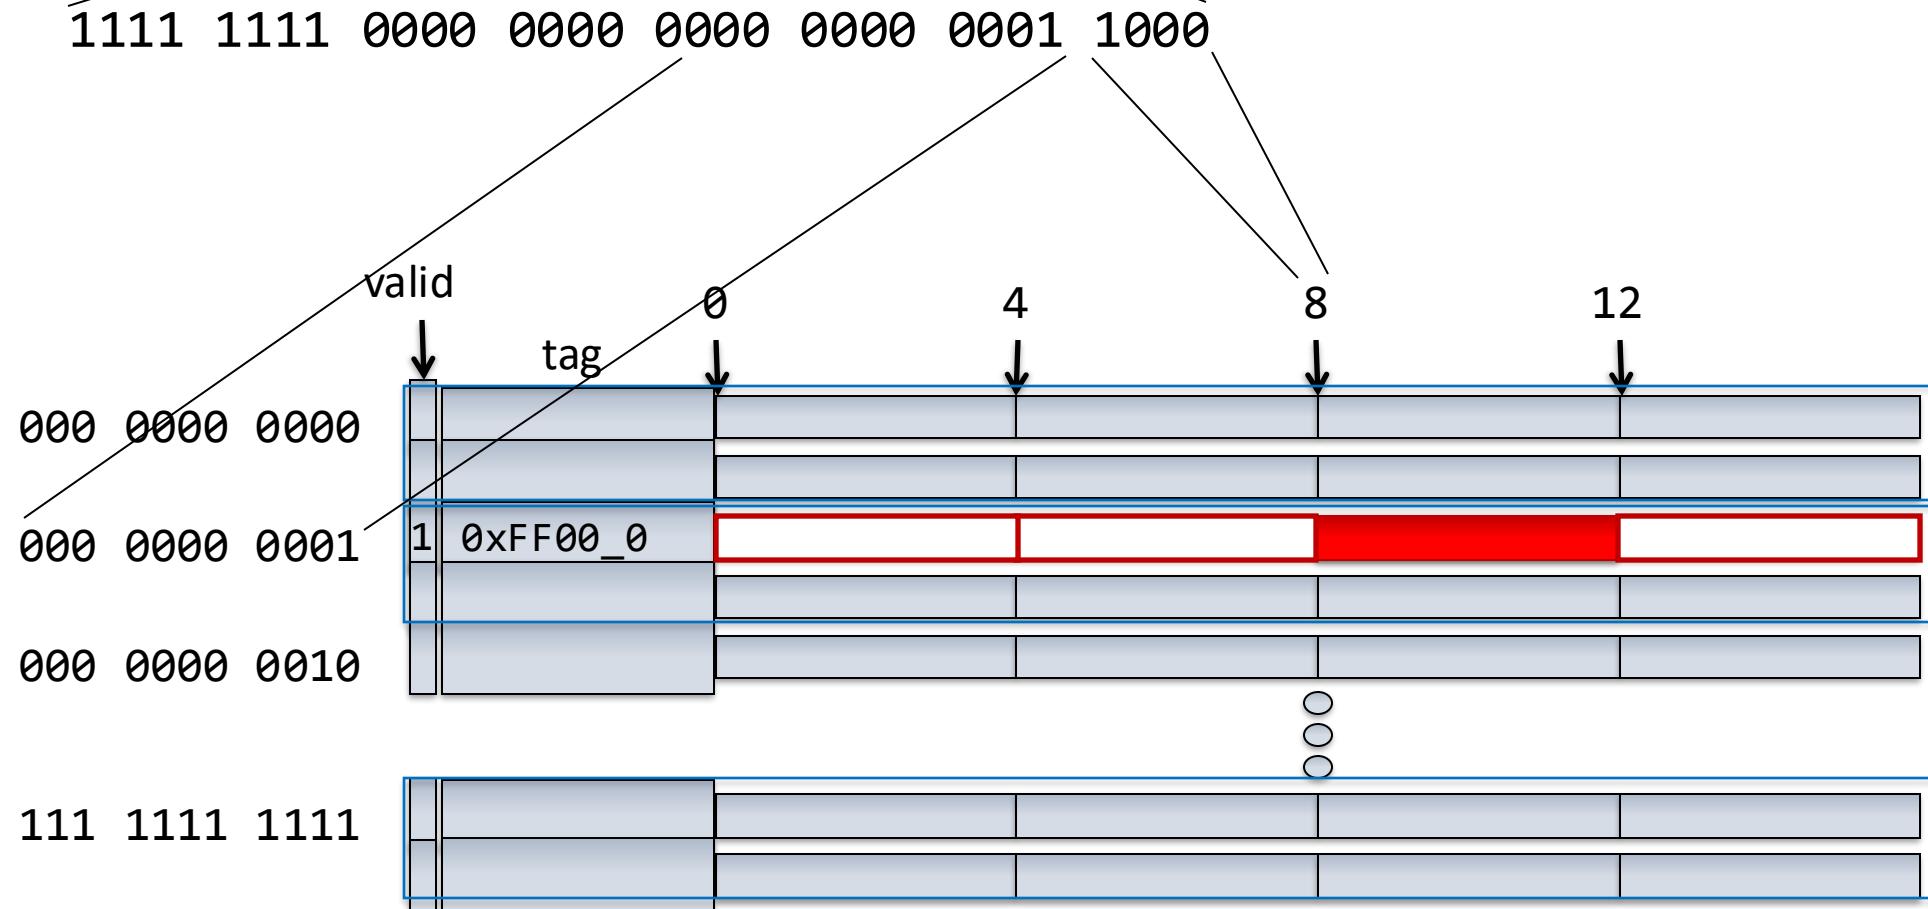
How many sets are in the cache? **2K sets**

We want to give a unique ID for each set

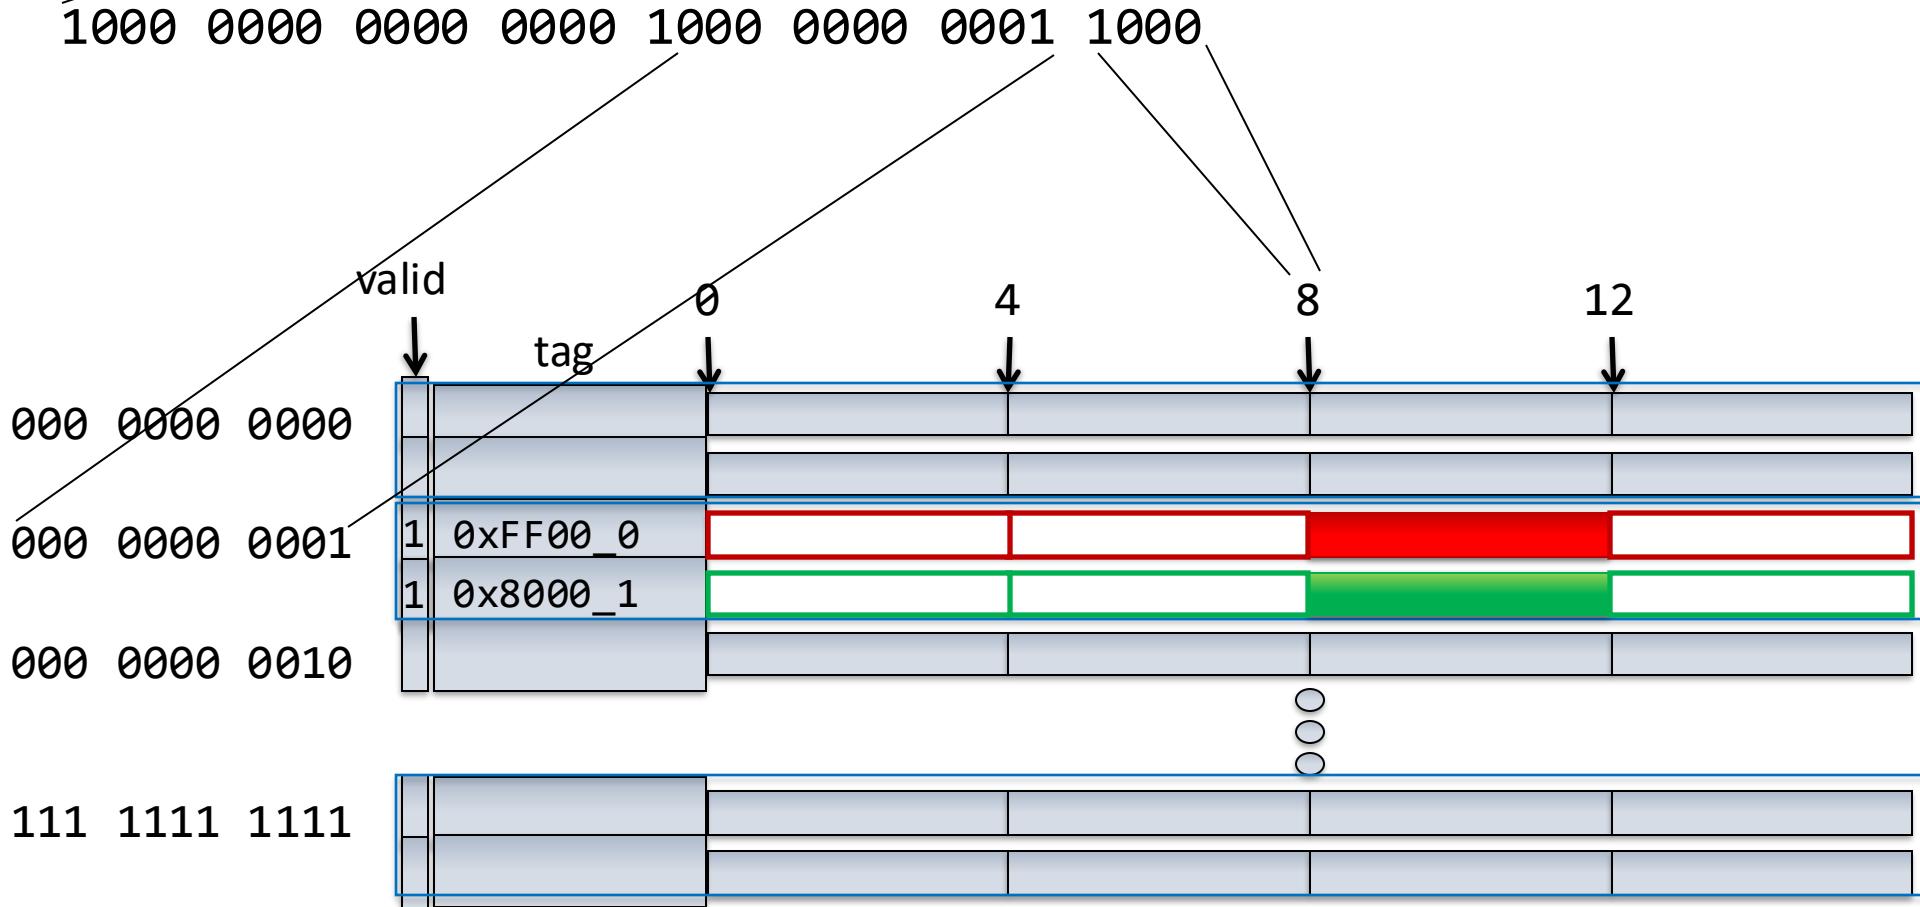
How many bits do we need for each set ID? **11 bits**



Address 0xFF00 0018

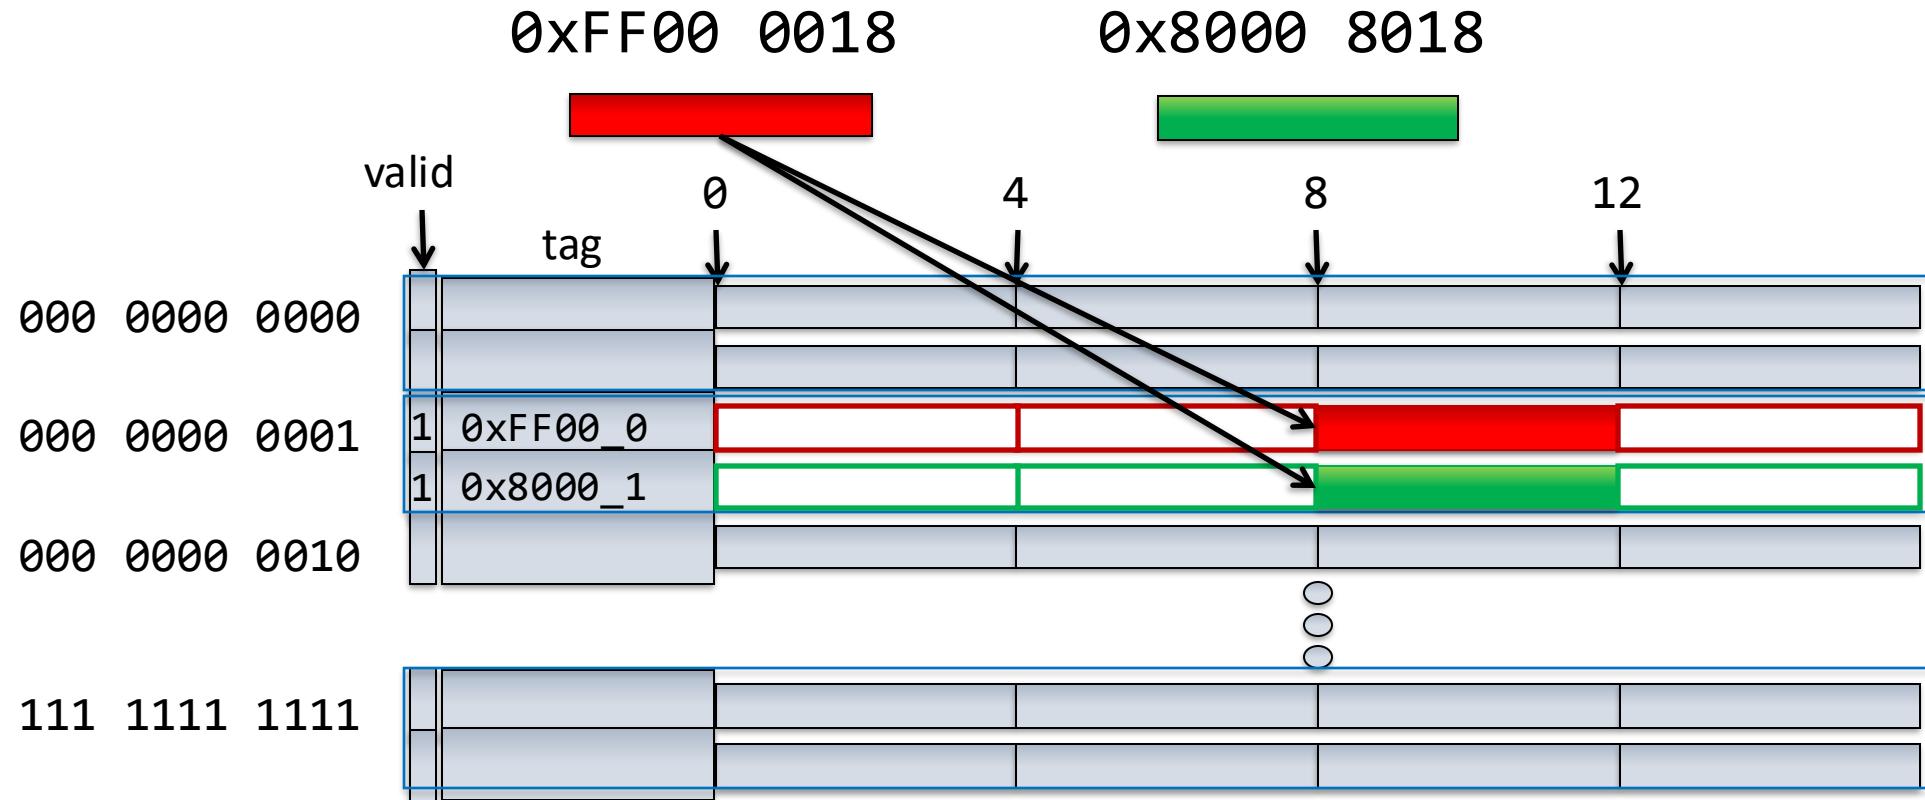


Address 0x8000 8018



Now both memory locations can be in the cache at the same time

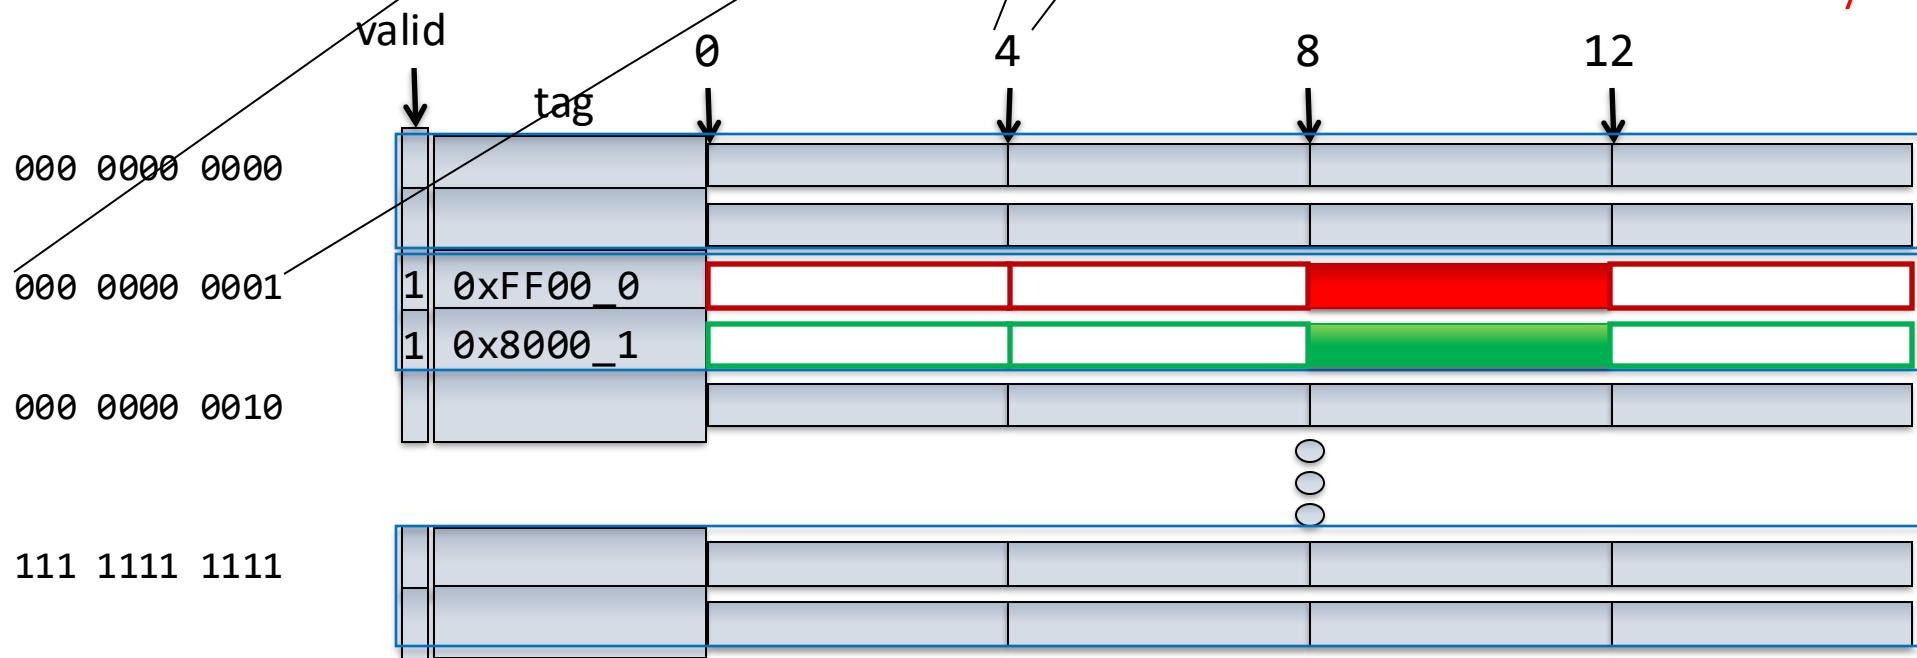
Each location can be in a set in two ways



This cache is two-way set associative

Address **0x7900 0014**

0111 1001 0000 0000 0000 0000 0001 0100

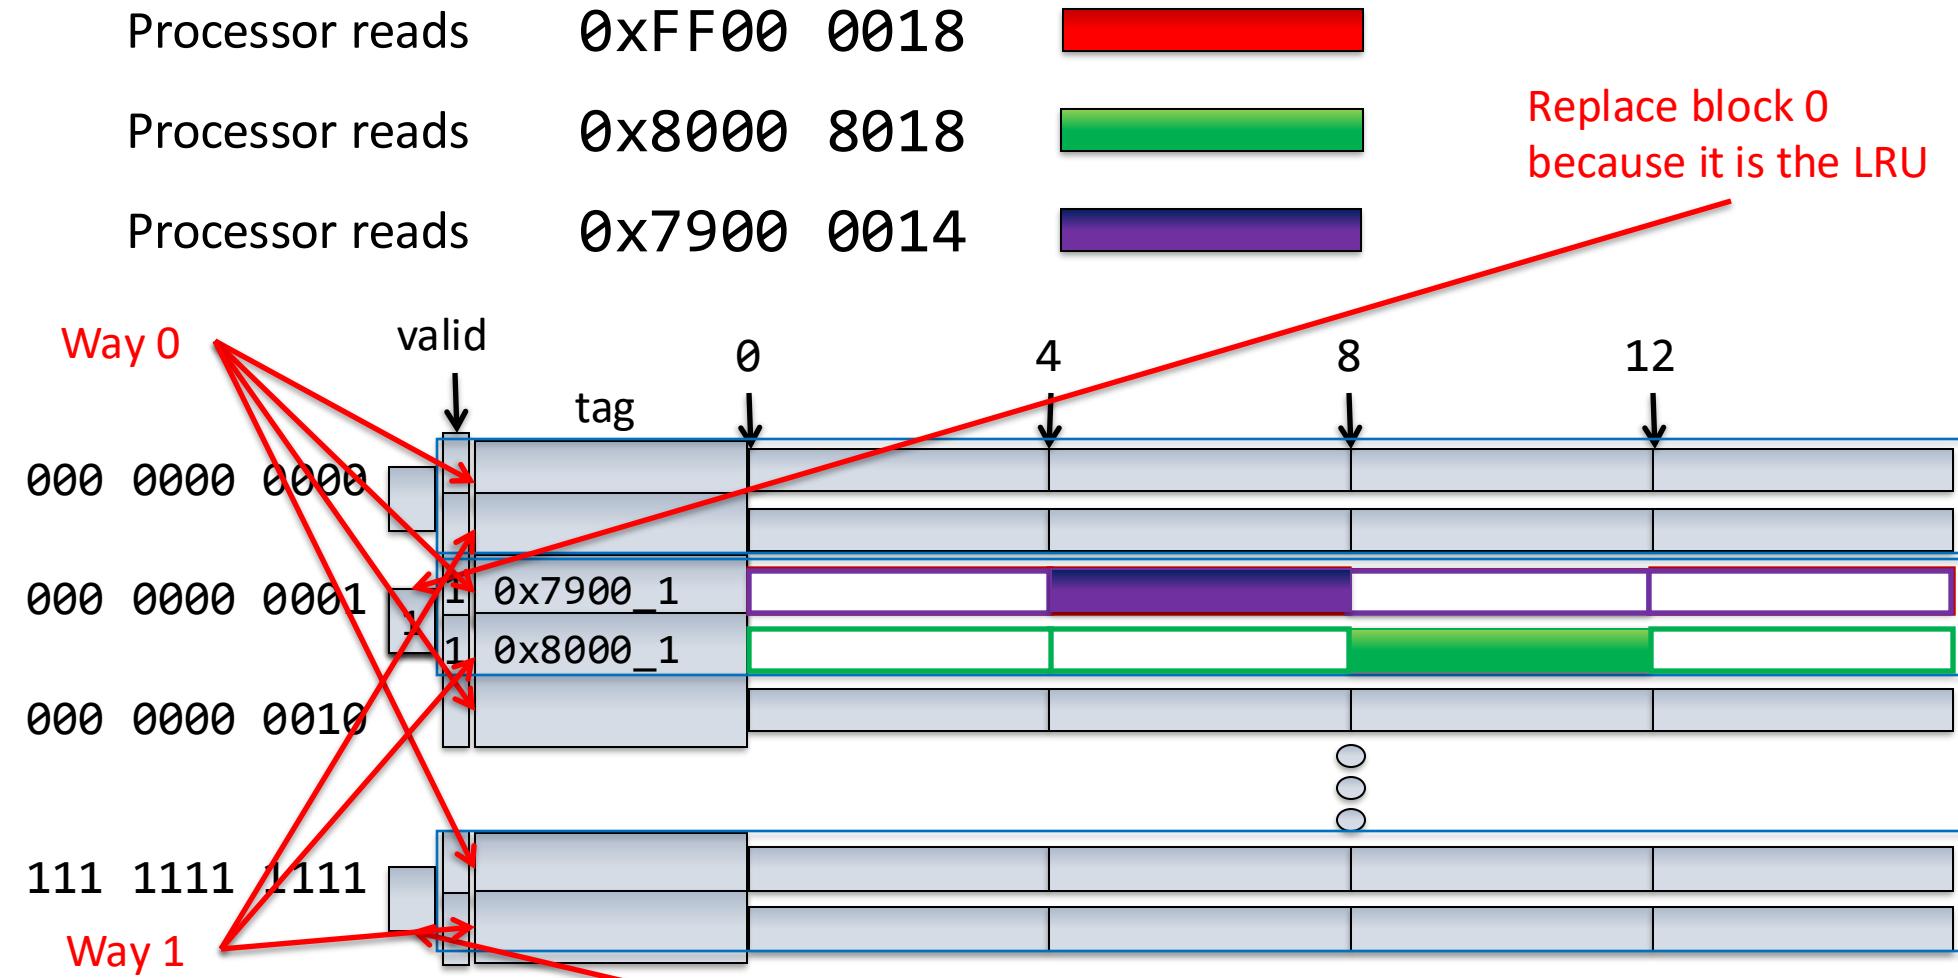


Which block should
be removed from the
cache to make
room for this block?

The one that has not
been used recently.

Tag in address does not match either tag in the set
⇒ This access is a cache miss

Least-Recently Used (LRU) bit



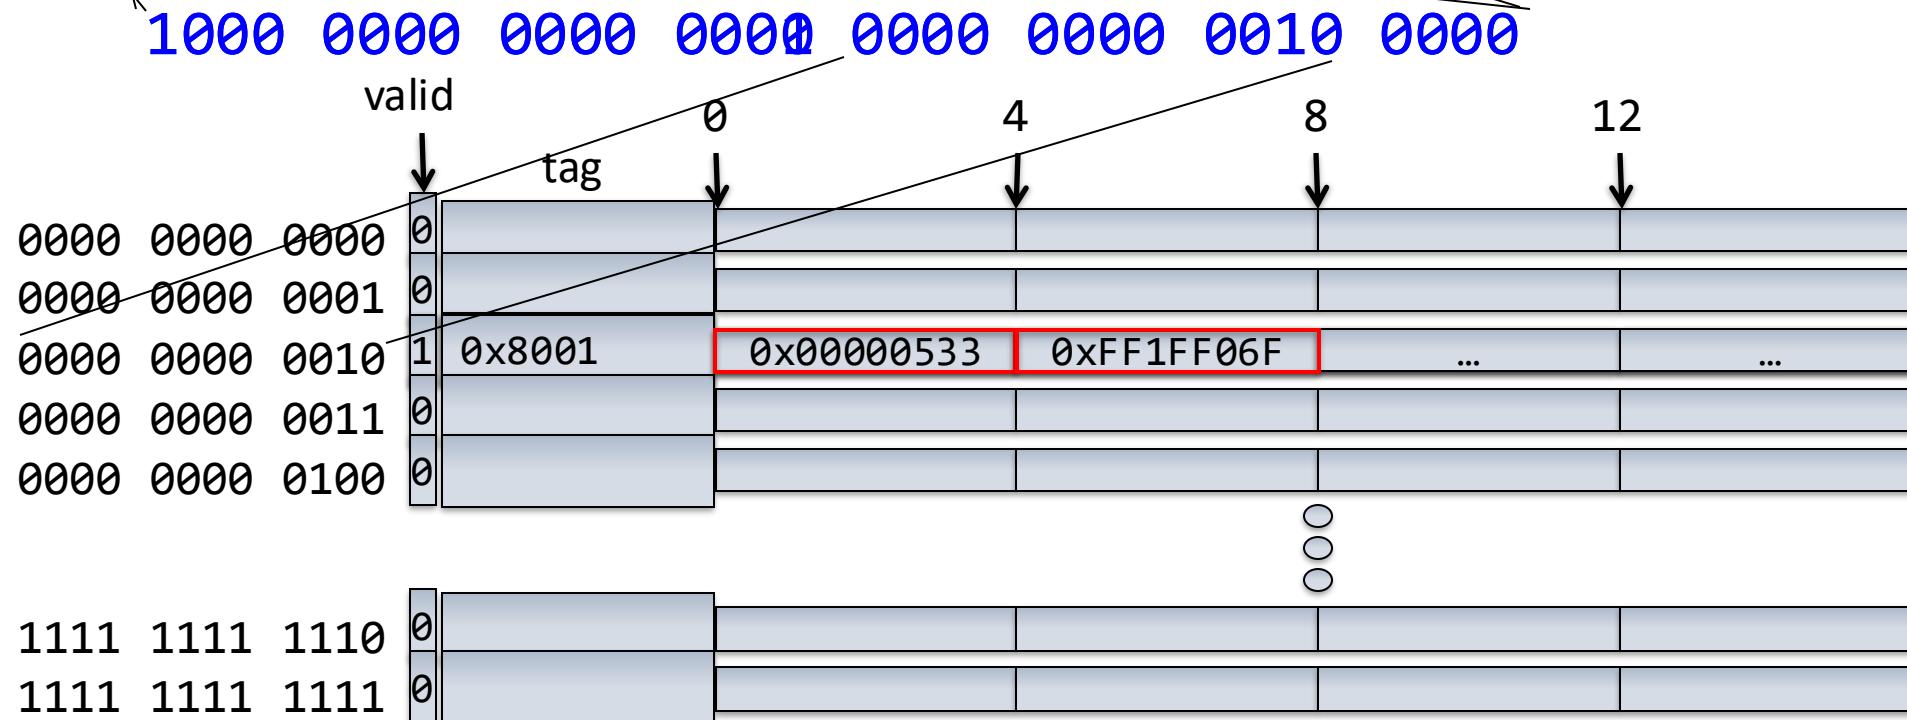
We associate an **LRU bit** with each set of blocks

The LRU bit indicates which block has not been used recently

Associativity and Performance: an example

0x8000 0020	loop:	add t1, t2, t3
0x8000 0024		beq t1, zero, false
0x8000 0028		jal zero, true
...		
0x8001 0020	false:	mv a0, zero
0x8001 0024	true:	jal zero, loop

Direct-mapped Cache



valid = 0 \Rightarrow cache miss

valid = 1 and tags match \Rightarrow cache hit

valid = 1 and tags do not match \Rightarrow cache miss

0x8000 0020	loop:	add	t1, t2, t3
0x8000 0024		beq	t1, zero, false
0x8000 0028		jal	zero, true
...			
0x8001 0020	false:	mv	a0, zero
0x8001 0024	true:	jal	zero, loop

Two-way
Set Associative
Cache

