

# Topic V31

Block Size and Associativity

Reading: (Section 5.3)

# Real Processor Caches

Intel Core duo (Yonah) has a 32 KB L1 cache that is 8-way set associative with 256-bit blocks

([http://www.xbitlabs.com/articles/mobile/display/core2duo\\_2.html](http://www.xbitlabs.com/articles/mobile/display/core2duo_2.html))



Assuming a 64-bit address, what is the address mapping (bits used for tag, index and offset) for this cache?


We need to answer these questions:

How many bytes per cache block (offset bits)

How many sets are there in the cache (index bits)

First find how many bytes per set

Intel Core duo (Yonah) has a 32 KB L1 cache that is 8-way set associative with 256-bit blocks and 64-bit address



$$\frac{32 \text{ KB}}{256 \frac{\text{bytes}}{\text{set}}} = \frac{32 \times 1024 \text{ bytes}}{256 \frac{\text{bytes}}{\text{set}}} = 32 \times 4 \text{ sets} = 128 \text{ sets} \Rightarrow 7 \text{ index bits}$$

$$64 \text{ address bits} - 7 \text{ index bits} - 5 \text{ offset bits} = 52 \text{ tag bits}$$

Intel Core duo (Yonah) has a 32 KB L1 cache that is 8-way set associative with 256-bit blocks and 64-bit address

---

offset bits: 5

index bits: 7

tag bits: 52




# LRU for Higher Associativity

How can LRU work for higher associativity?

True LRU is expensive, lets see an example

# LRU for a 4-way Set Associative Cache (example)



Assumption:

11: Least recently used


10: Second least recently used

01: Third least recently used

00: Most recently used.

For true LRU, if a reference with index 1 that matches on way 0 occurs, what must happen?

# Alternative to True LRU: Not Most Recently Utilized





Ensures that replacement is not in the MRU group

Way selection of the non-MRU group is random

# Tree-Based Algorithm (LRU approximation)

- Example: 4-way set associative cache.

References:



Requires  $n-1$  bits per set for a  $n$ -way set-associative cache.

Baer, p. 235

# How Much Associativity

Increased associativity decreases miss rate  
But with diminishing returns


Simulation of a system with 64KB D-cache, 16-word blocks,  
SPEC2000

1-way: 10.3%


2-way: 8.6%

4-way: 8.3%

# Set-associative Cache Organization



# Set-associative Cache Organization



# Replacement Policy

Direct mapped: no choice

Set associative

Prefer non-valid entry, if there is one

Otherwise, choose among entries in the set

Least-recently used (LRU)

Choose the one unused for the longest time

Simple for 2-way, manageable for 4-way, too hard beyond that

Random

Gives approximately the same performance

# Larger Block Size (example)

# Cache with 64 blocks, 16 bytes/block

Map address  $1060_{10}$  to its cache block

$$\text{Block address} = \lfloor 1060/16 \rfloor = 66$$

Block number = 66 modulo 64 = 2

# Block Size Considerations

Larger blocks should reduce miss rate

Due to spatial locality

But in a fixed-sized cache

Larger blocks  $\Rightarrow$  fewer of them

More competition  $\Rightarrow$  increased miss rate

Larger blocks  $\Rightarrow$  pollution

Larger miss penalty

Can override benefit of reduced miss rate