

Topic V36

Virtual Memory

Reading: (Section 5.7)

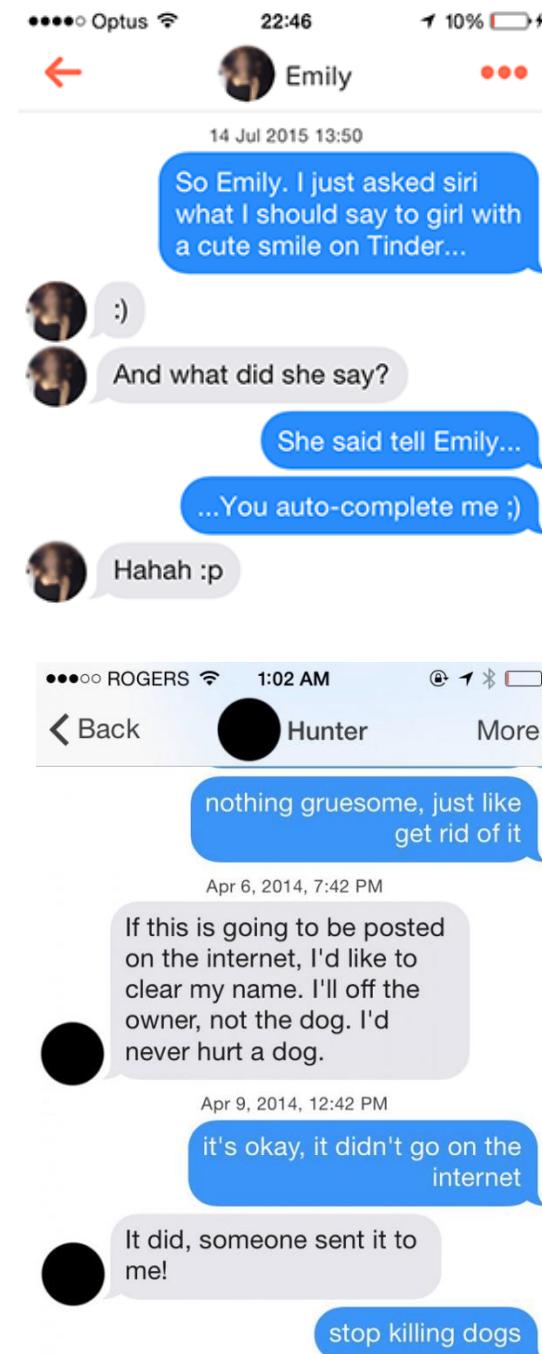
Motivation 1:
Protect Data

Emily's Chat Code

```
.text  
0x8000 0000    la      a0, MsgBuffer  
0x8000 0004    jal     ra, WriteMsg  
0x8000 0008    la      a0, MsgBuffer  
0x8000 000C    jal     ra, ReadReply  
  
...  
  
.data  
MsgBuffer: .space 256
```

Hunter's Chat Code

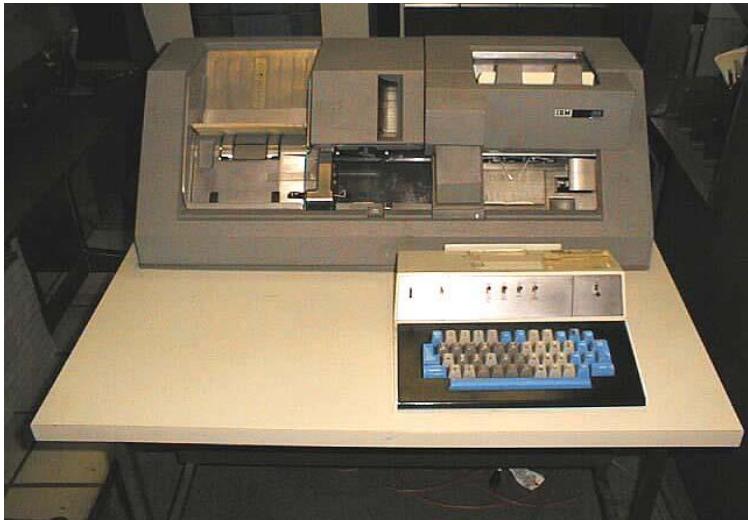
```
.text  
0x8000 0000    la      a0, MsgBuffer  
0x8000 0004    jal     ra, WriteMsg  
0x8000 0008    la      a0, MsgBuffer  
0x8000 000C    jal     ra, ReadReply  
  
...  
  
.data  
MsgBuffer: .space 256
```



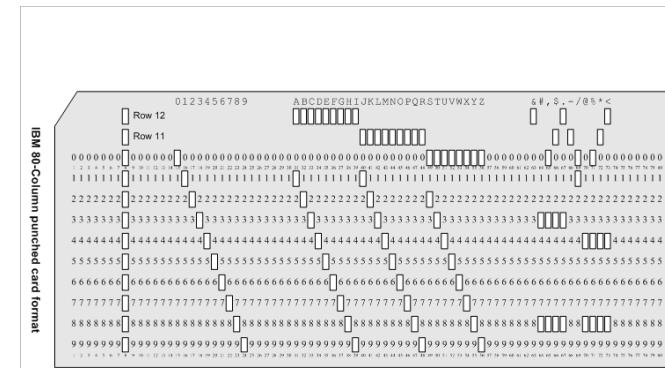
Motivation 2:
Provide Illusion of
Larger Memory



IBM 1130



Punch Card Perforator



Punch Card

Introduced in 1965

15-bit address space

Maximum of 64 Kbytes of memory

Later could have up to 4 terminals

Fortran compiler run on machine
with 8Kbytes of memory



48-bit address bus

Memory (RAM)

M2: up to 24 GB

M2 Pro: up to 32 GB

M2 Max: up to 96 GB

M2 Ultra: up to 192 GB



256GB to 8TB SSD



MacBook Air
Dec, 2023



48-bit address bus

Memory (RAM)

M2: up to 24 GB

M2 Pro: up to 32 GB

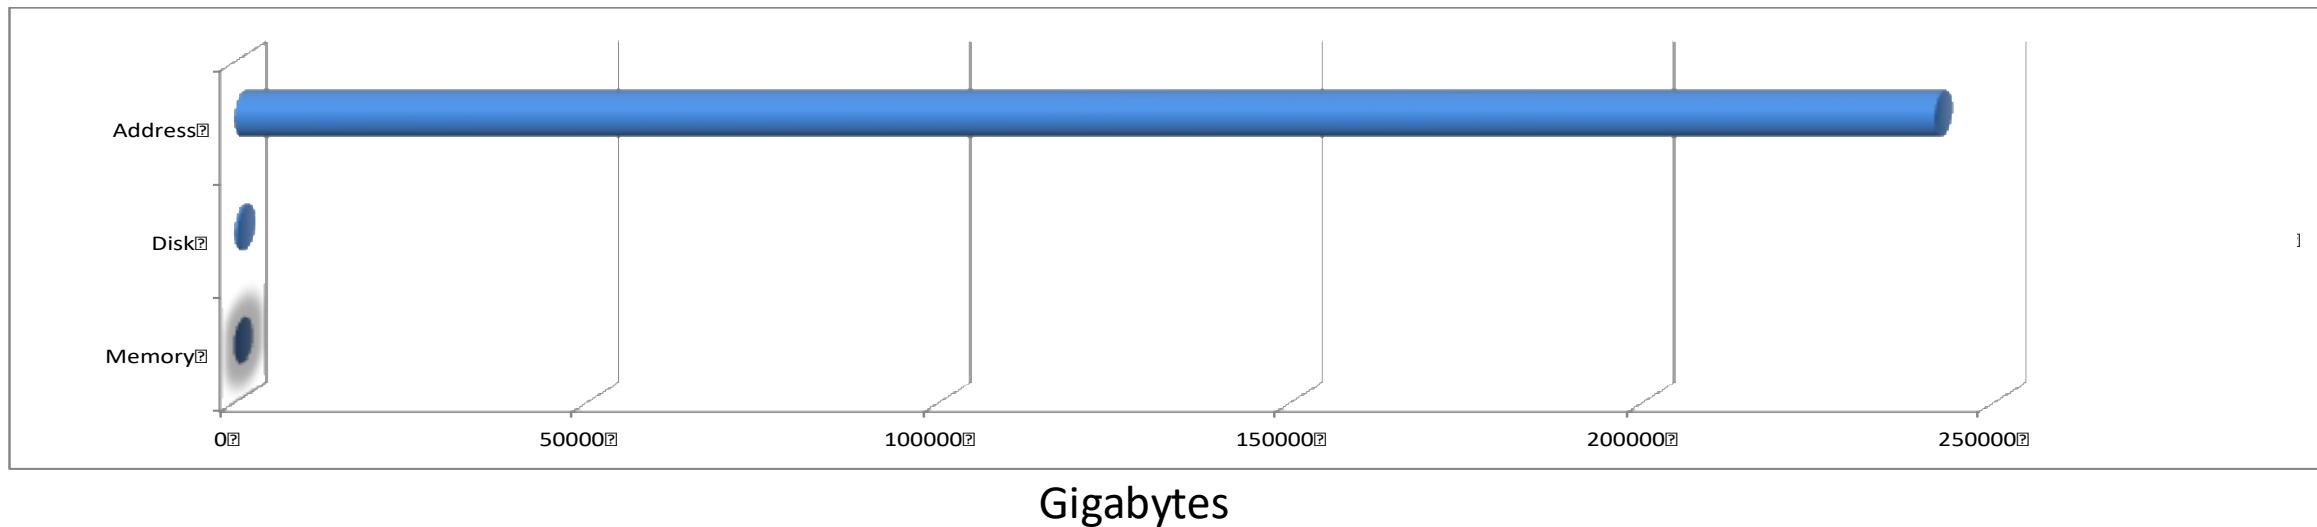
M2 Max: up to 96 GB

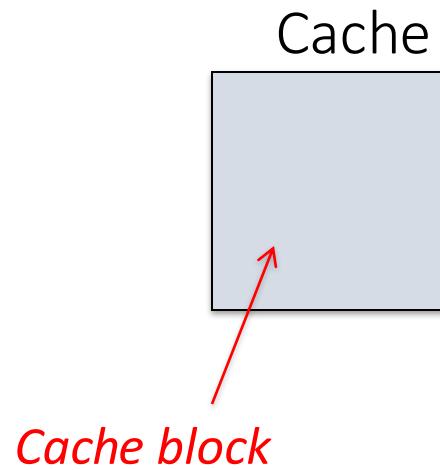
M2 Ultra: up to 192 GB



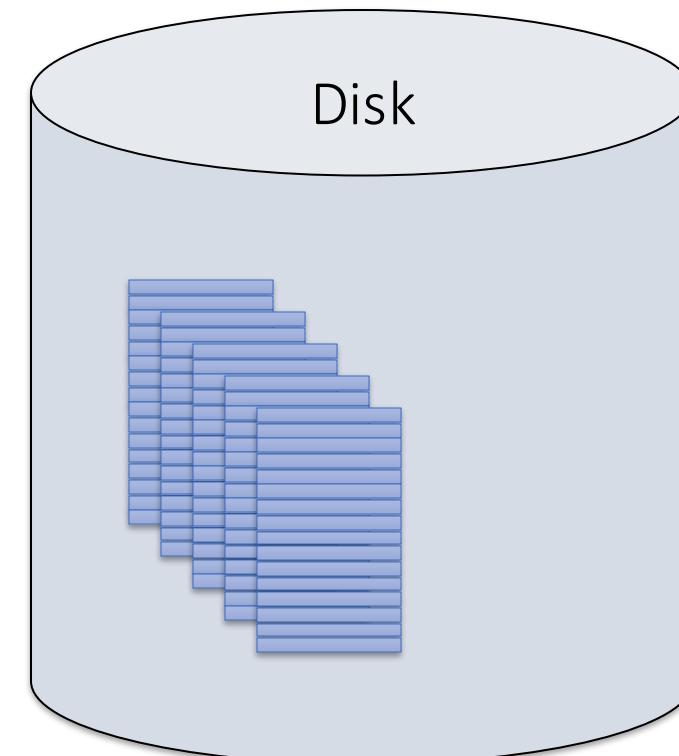
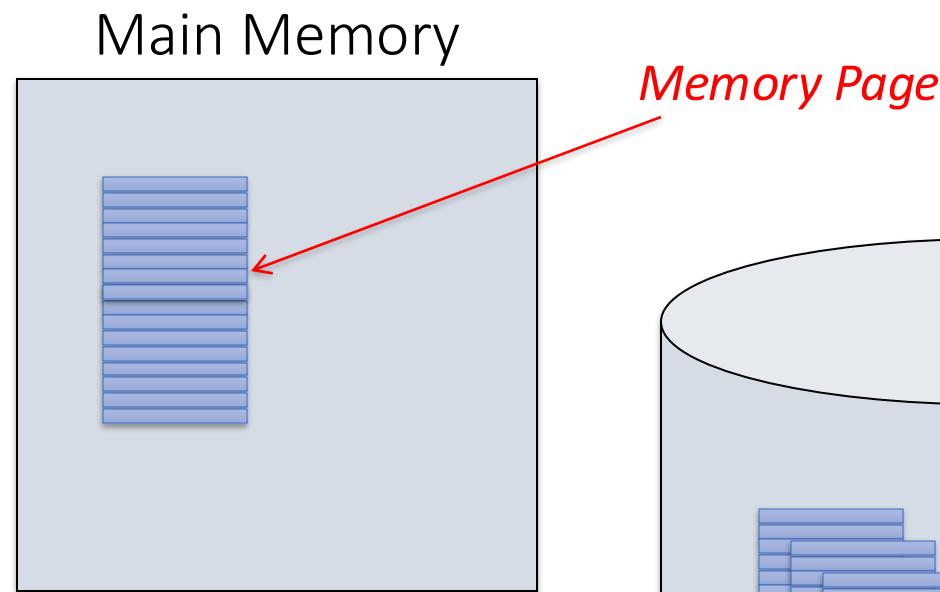
256GB to 8TB SSD

$$2^{48} = 2^8 \text{ TB} = 256 \text{ TB} = 241,664 \text{ GB}$$





Main Memory is used as a “cache” for the content of the disk

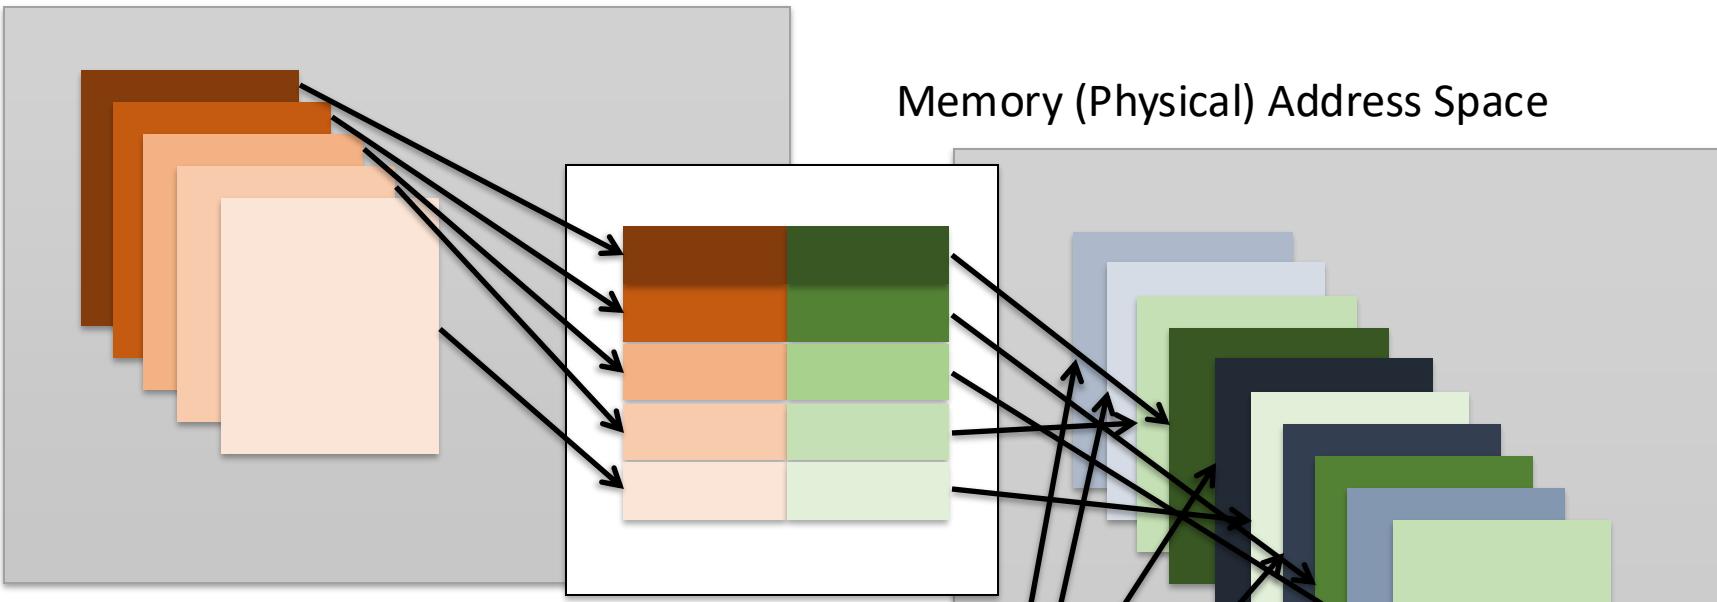


This “caching system” is called
Virtual Memory

Virtual Memory is managed
jointly by the hardware and OS

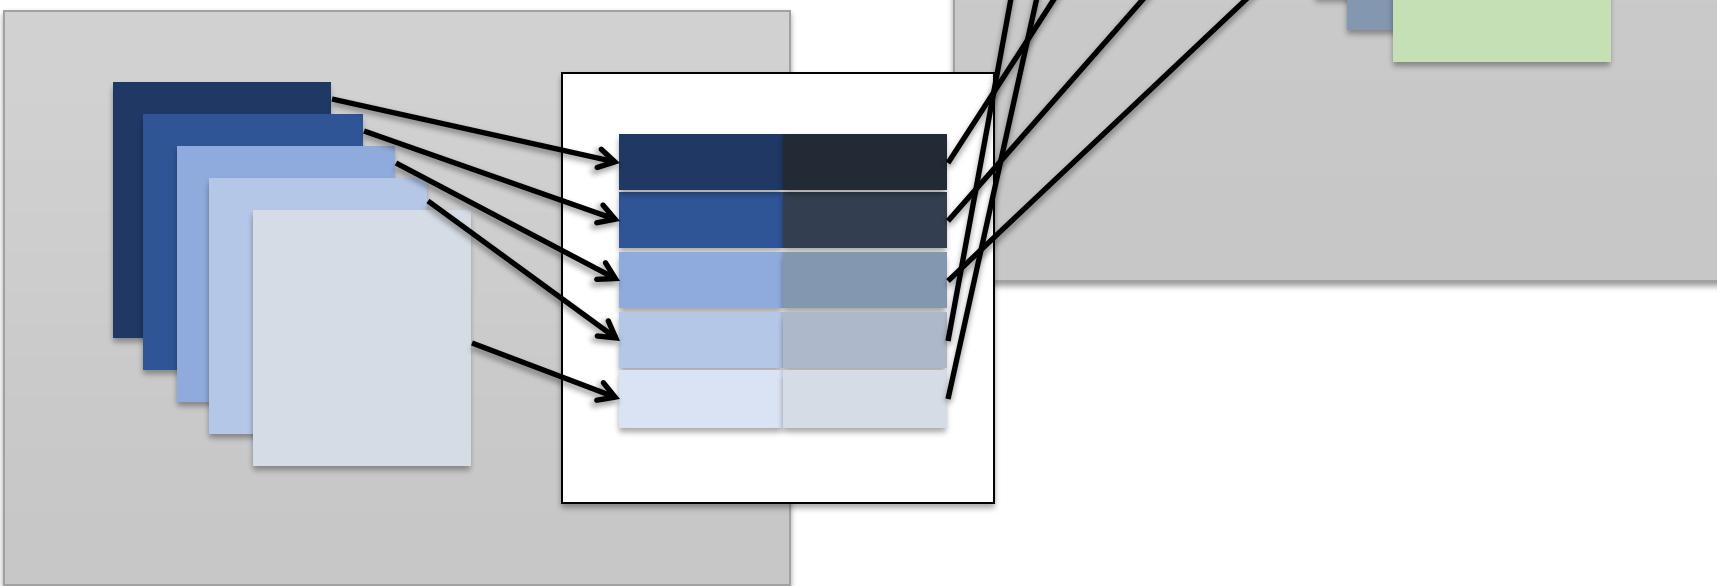
Virtual Memory

Program A Private Virtual Address Space



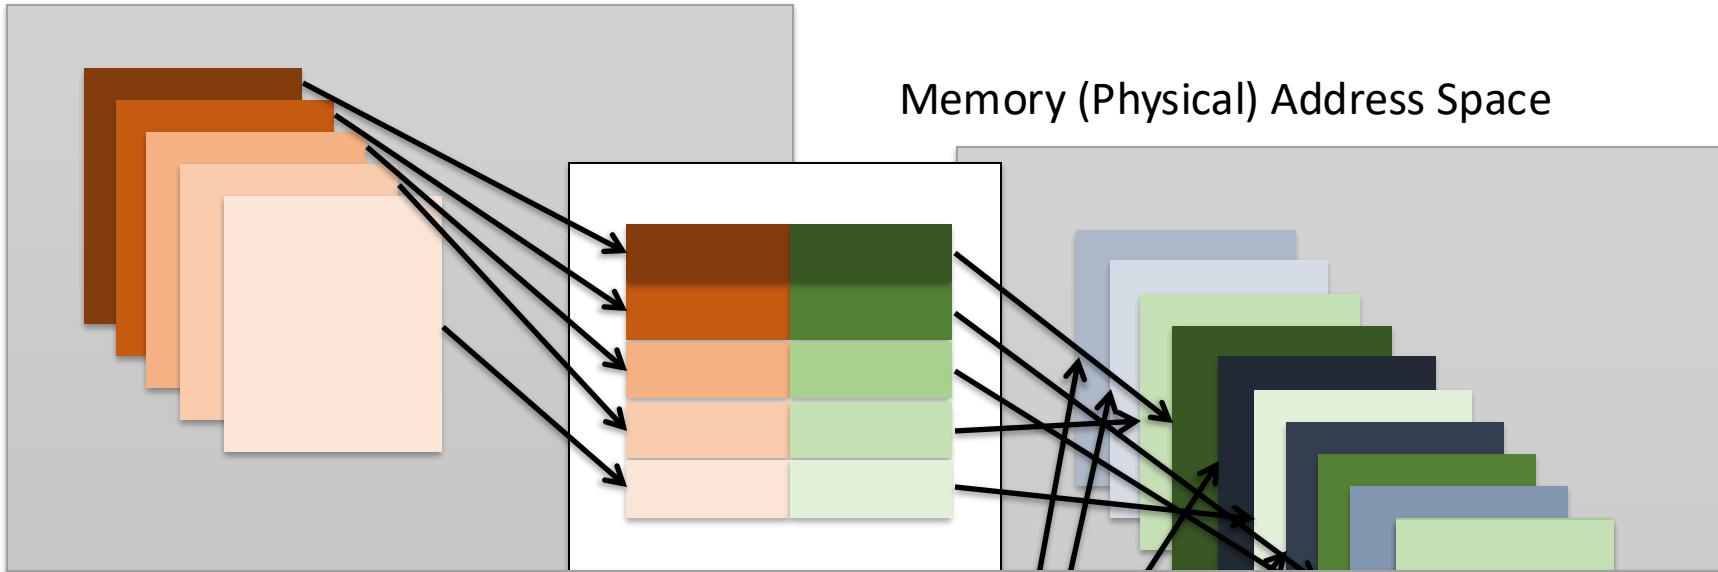
Memory (Physical) Address Space

Program B Private Virtual Address Space



Program A Private Virtual Address Space

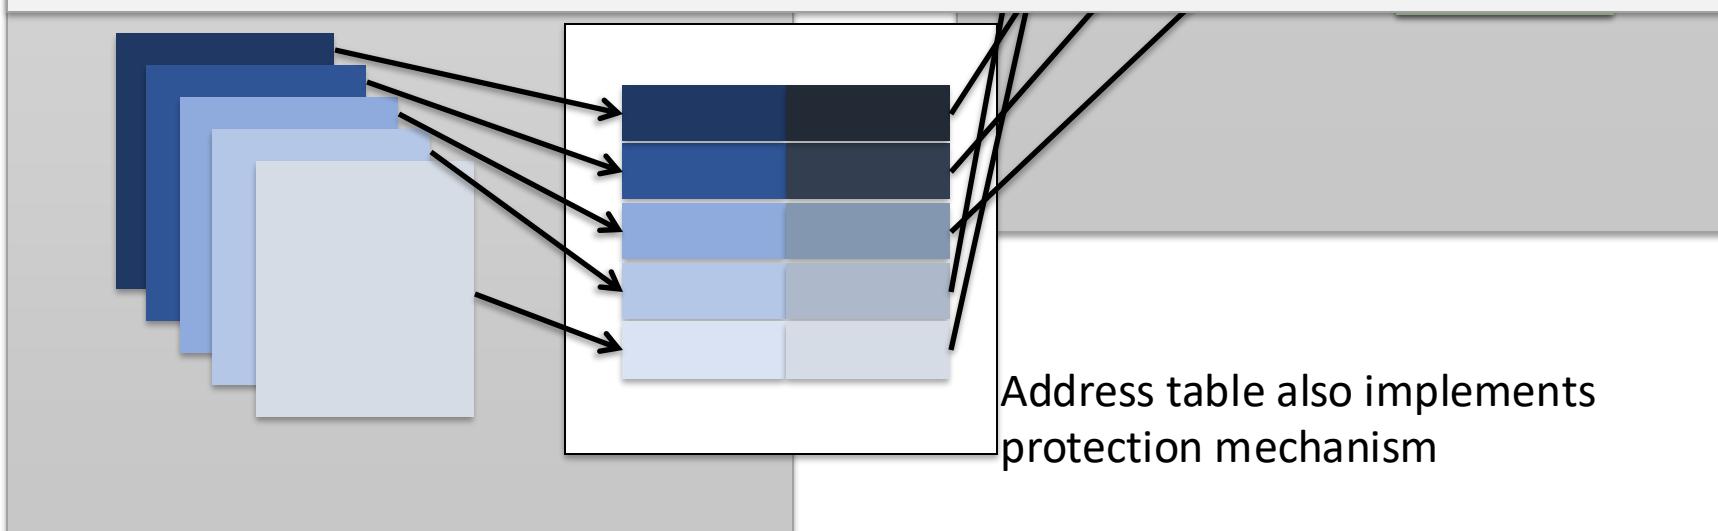
Hardware and OS implement address translation



Program B

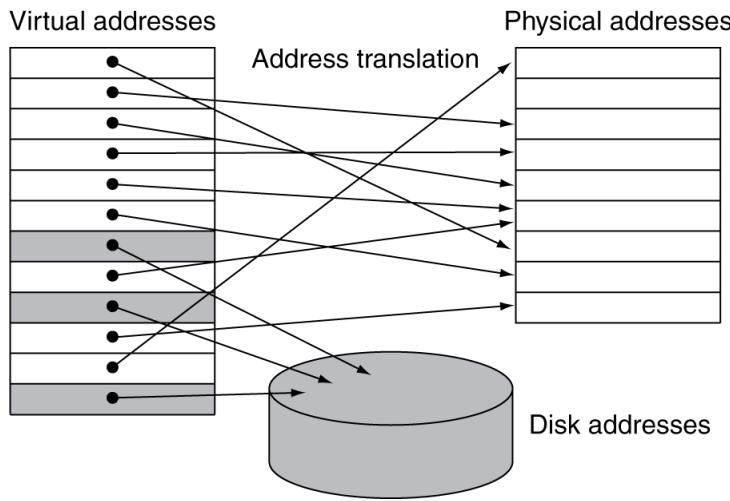
What happens if a page requested by the processor is not in memory?

Virtual Memory “miss” is called a *page fault*



Address Translation

Fixed-size pages (e.g., 4K)



Virtual address

31 30 29 28 27 15 14 13 12 11 10 9 8 3 2 1 0

Virtual page number

Page offset

Physical address

Try to **minimize page fault rate**:

- Fully associative page placement
- Smart replacement algorithms

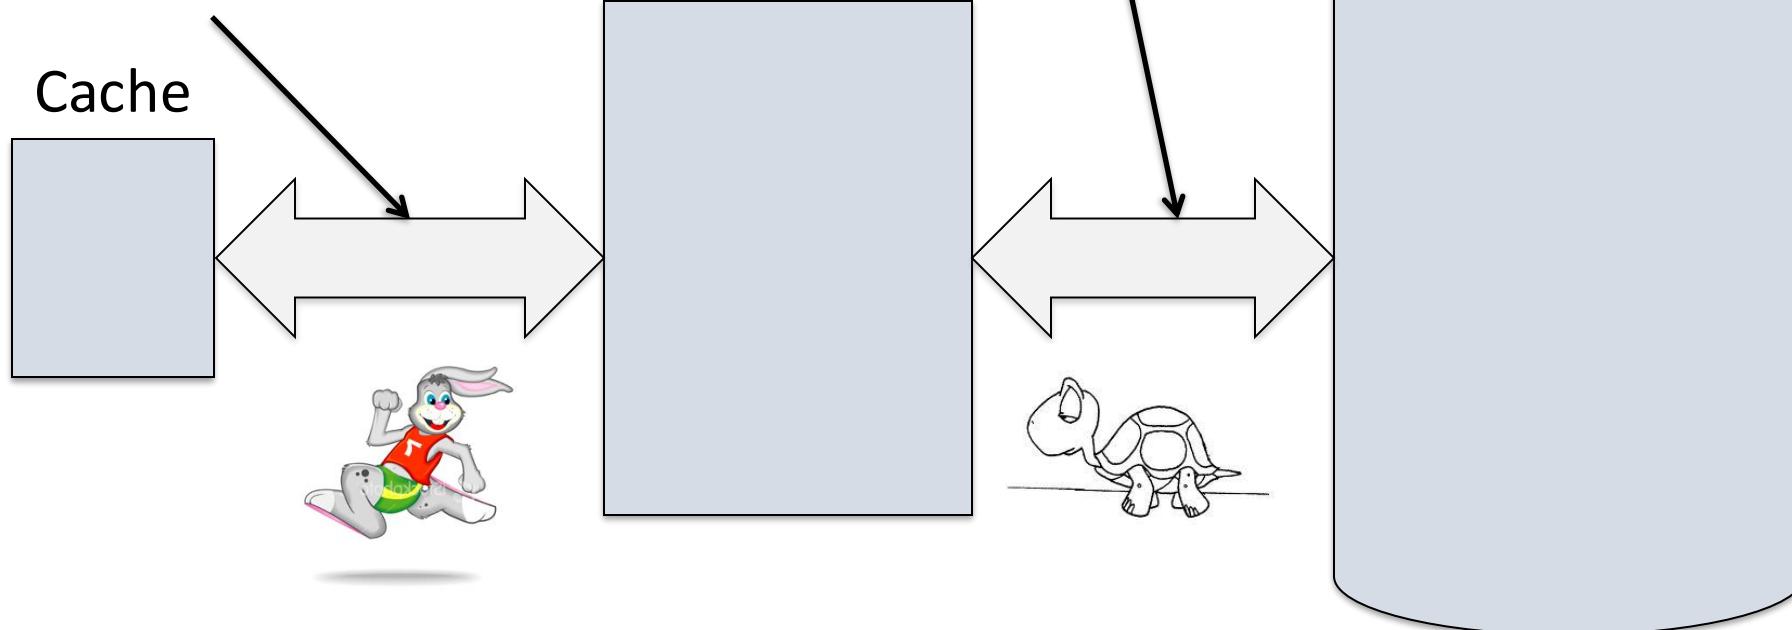
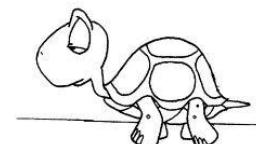
Fast: hundreds of cycles
handled by hardware,
processor waits for result

Cache

Main Memory

Disk

Slow: millions of cycles
exception to OS,
context switch

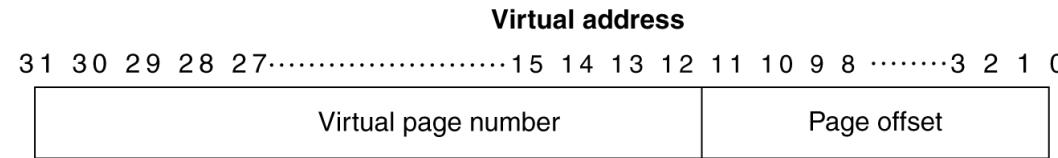


Unlike caches, virtual memory is not a place where data is stored

Virtual memory is:

- a mapping between two address spaces
- a mechanism that allows main memory to serve as a cache to disk storage

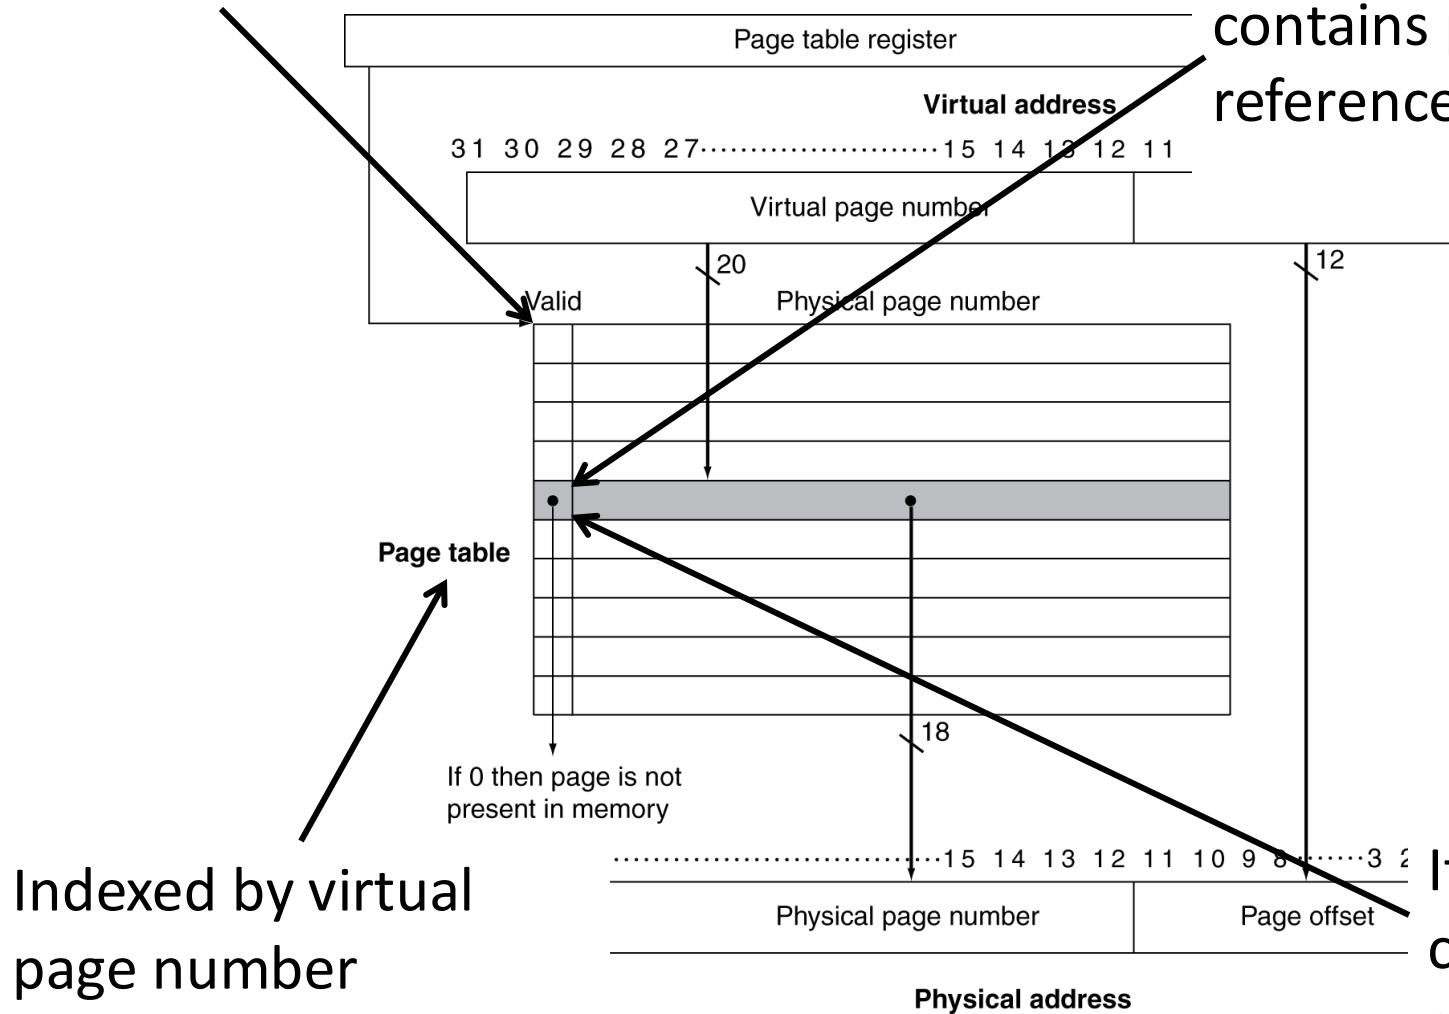
Translation Using a Page Table



Page Table Register in CPU
points to start address of
page table in memory

Using a Page Table

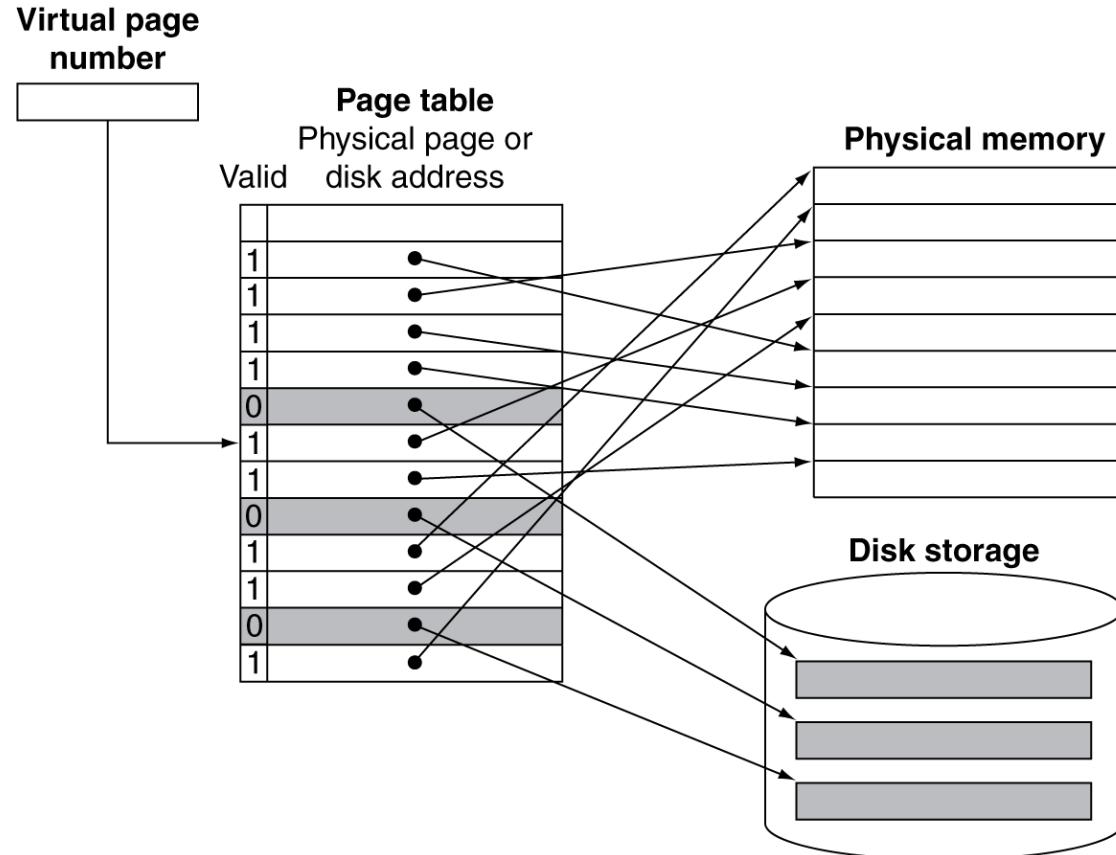
If page is in memory:
contains physical address,
referenced bit, dirty, protection info



Indexed by virtual
page number

If page is not in memory:
contains location in swap
space on disk

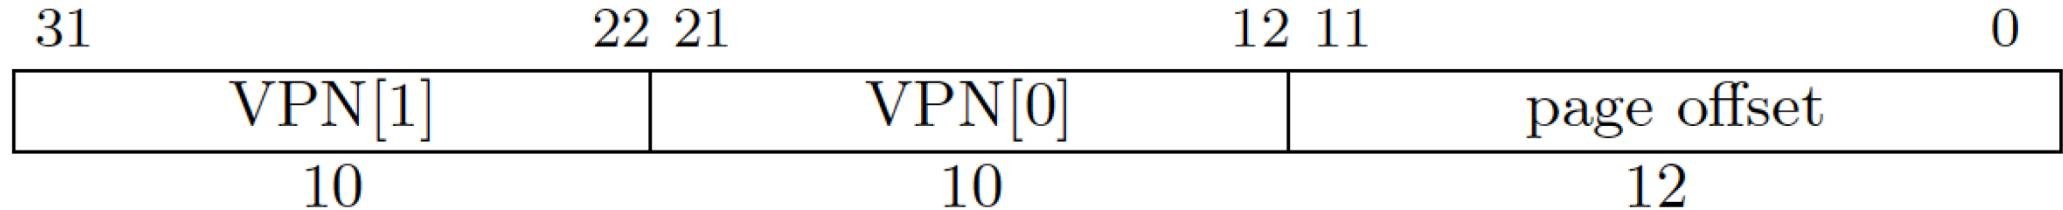
Mapping Pages to Storage



Address Translation in RISC-V

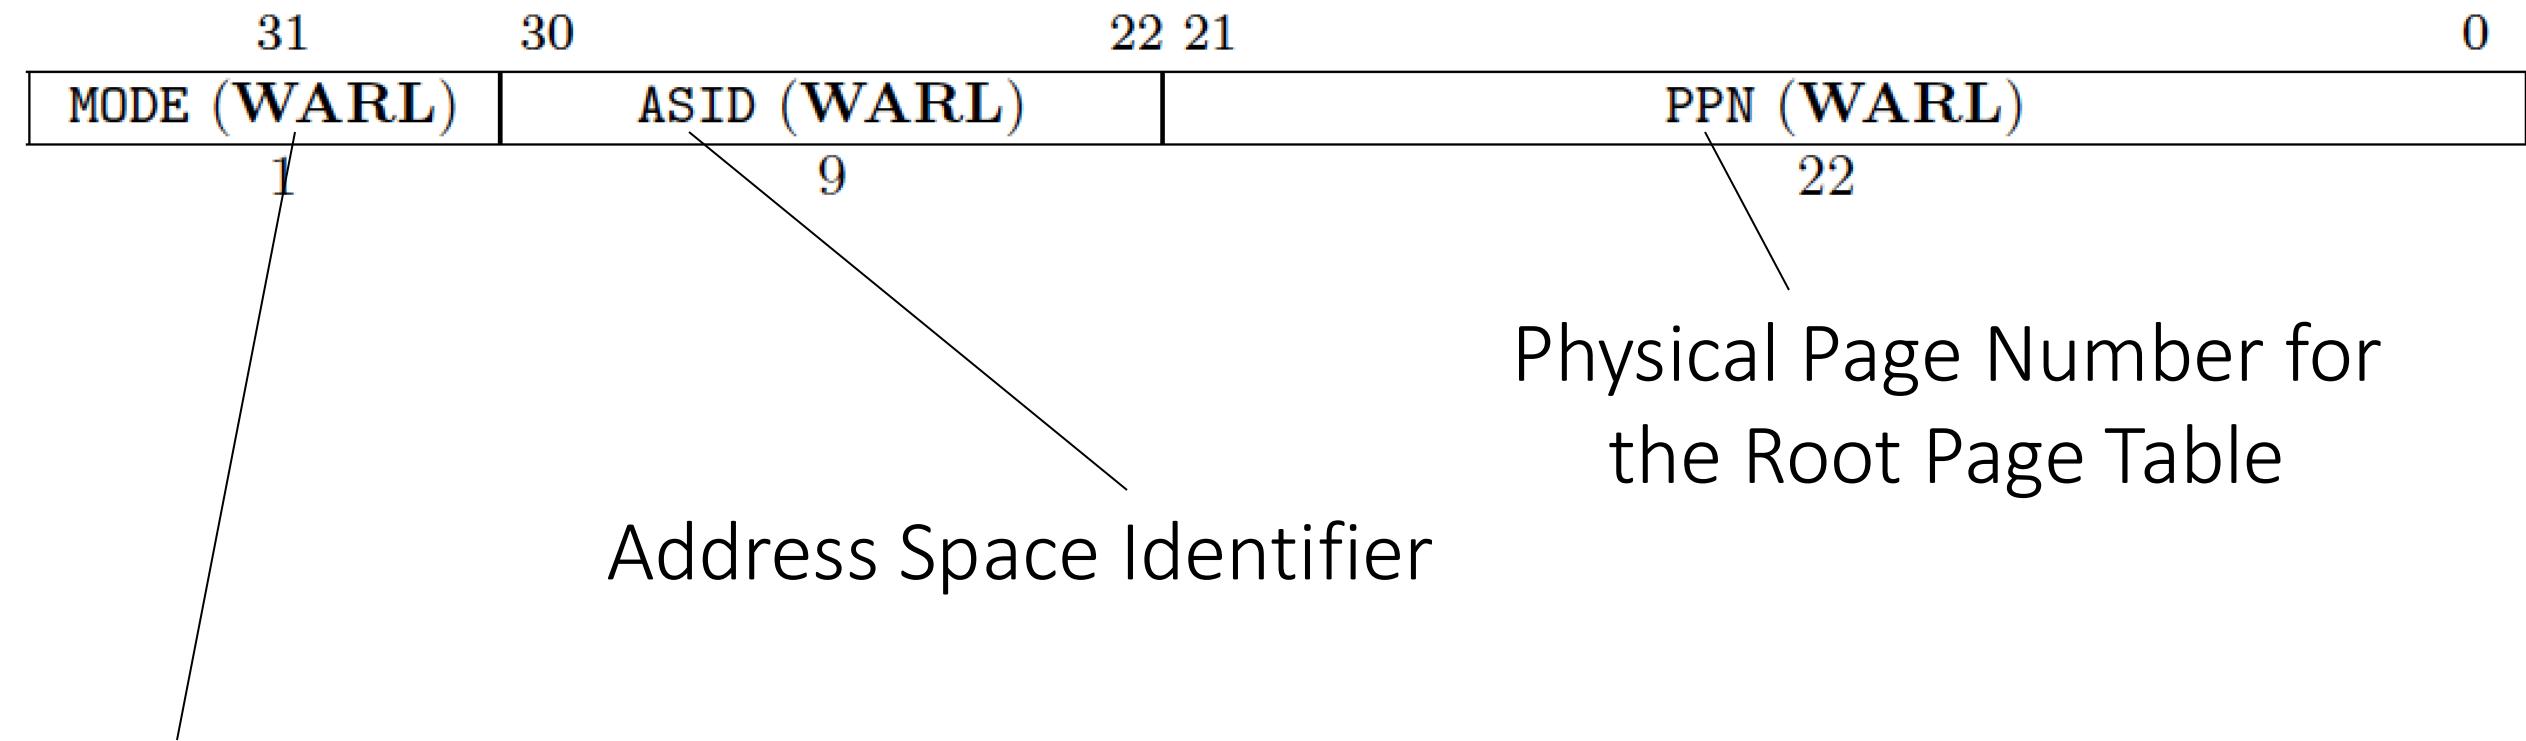
Radix-Based Page Table
And SATP register

Address Translation in RISC-V (32 bits)



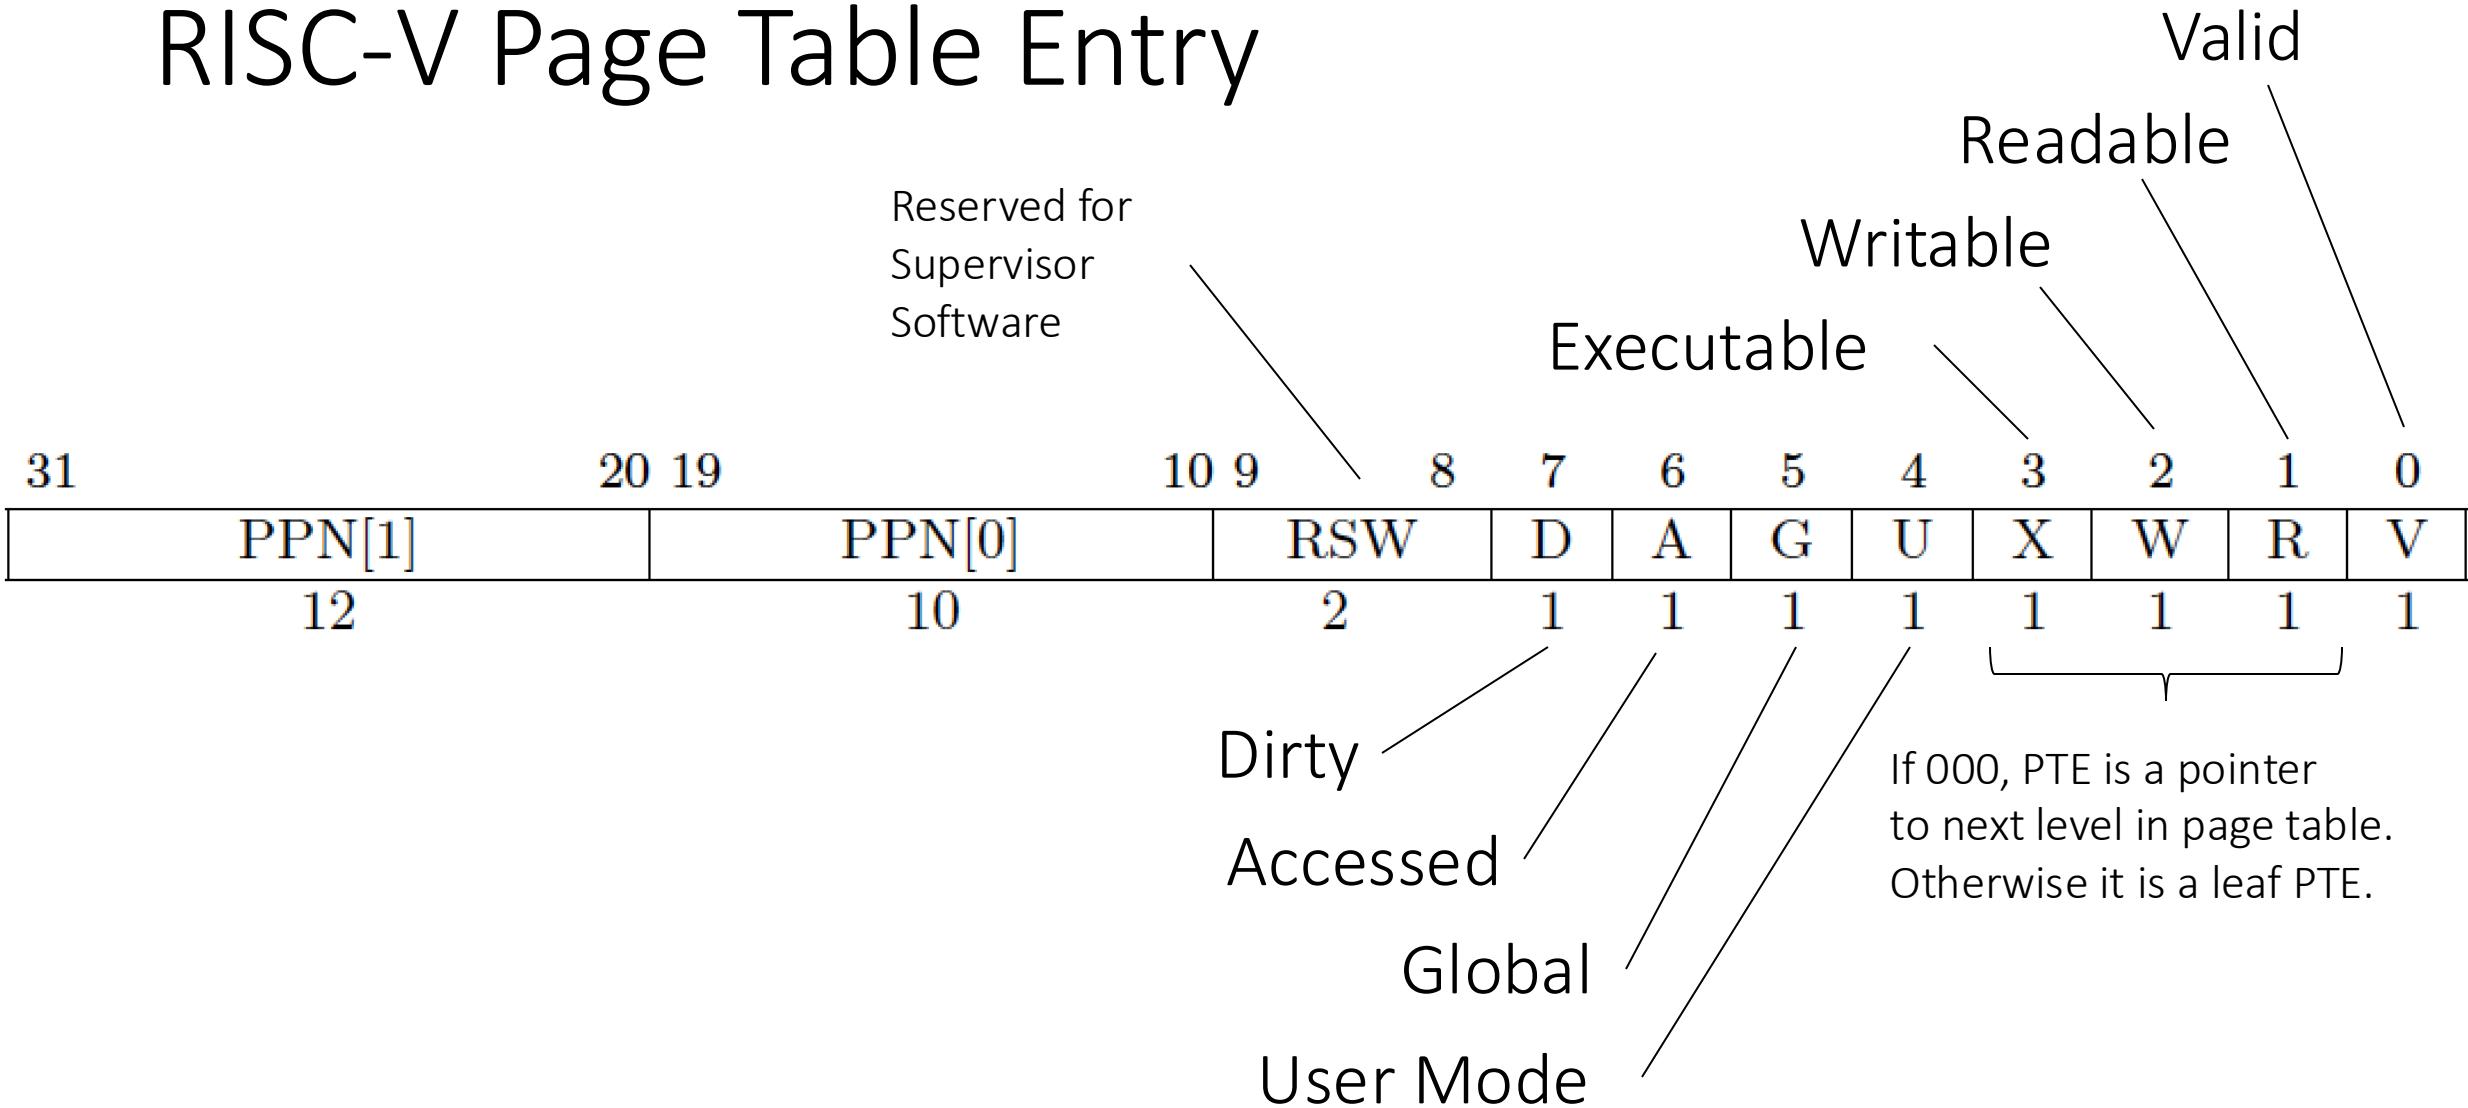
Sv32 virtual address.

RISC-V Supervisor Address Translation and Protection (SATP) register



WARL: Write Any Values, Reads Legal Values

RISC-V Page Table Entry



Radix Page Table

64-bit Virtual Address

	63	48	47	39	38	30	29	21	20	12	11	0
	sign ext.		idx 4		idx 3		idx 2		idx 1		offset	

Radix Page Table only makes new pages when accessing a new page

Each idx corresponds to a page table entry in its associated level

idx 4 -> which page table entry to access in the base page table

idx 3 -> which page table to access in that new page table from the previous page table entry

...

idx 1 -> which page table to access in memory from the last page table

offset -> the data to be accessed on that page

Hash, Don't Cache (the Page Table)

Idan Yaniv

Technion – Israel Institute of Technology

idanyani@cs.technion.ac.il

Dan Tsafrir

Technion – Israel Institute of Technology

dan@cs.technion.ac.il



Congress Center, Antibes Juan-les-Pins, June 14-18, 2016

Virtual Address -> Physical Address

64-bit Virtual
Address

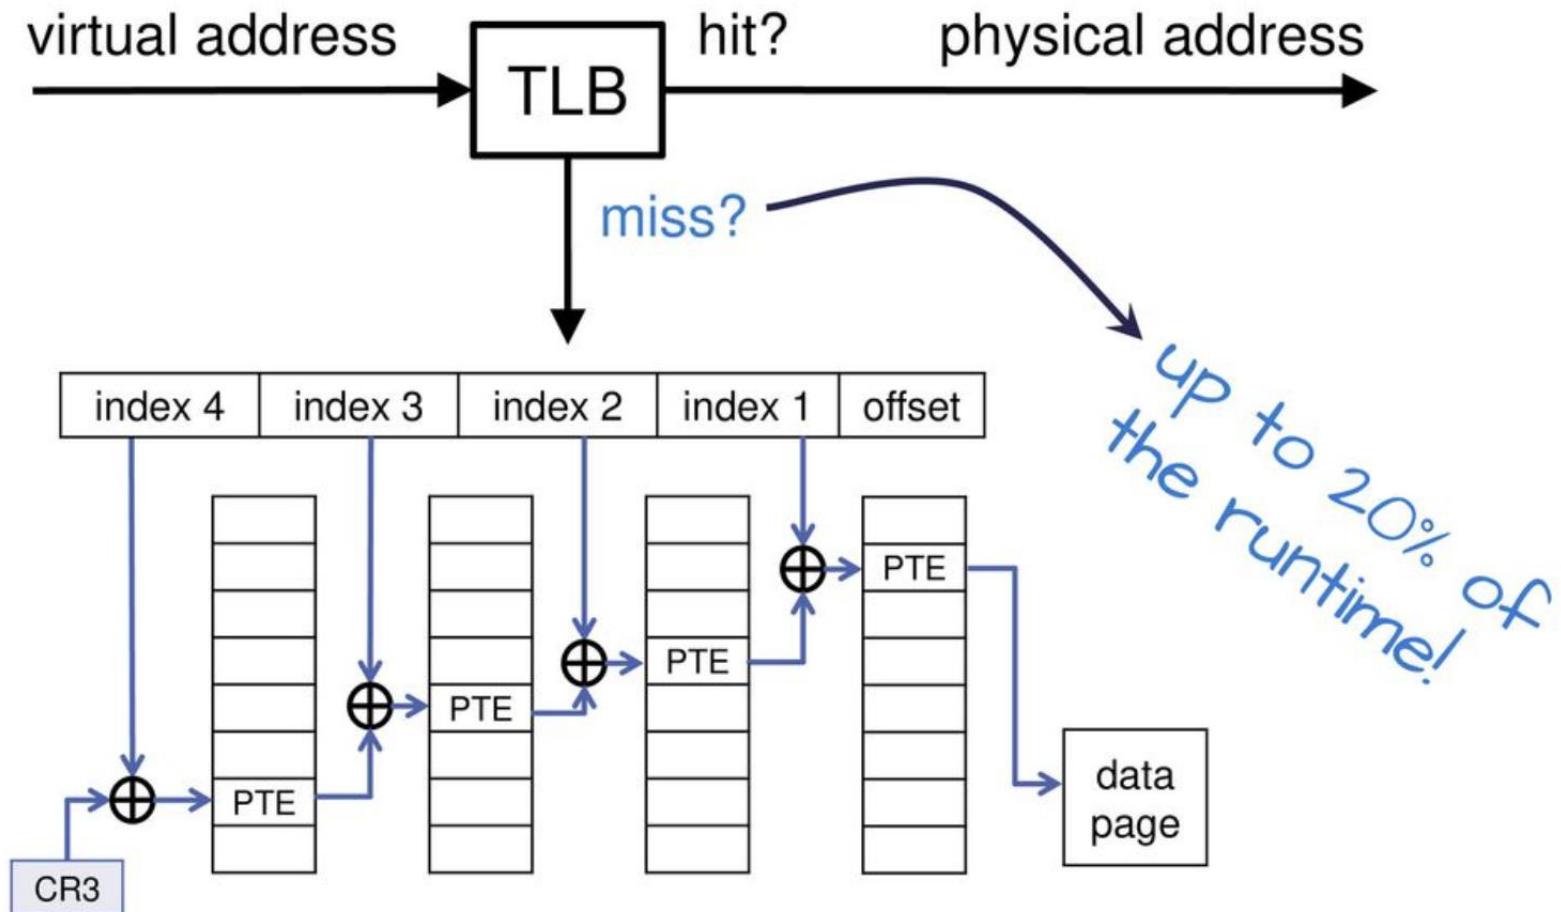
63	48	47	39	38	30	29	21	20	12	11	0
sign ext.	idx 4		idx 3		idx 2		idx 1		offset		

Control Register
contains base address
of root table

physical
address

SATP: Supervisor Address Translation and Protection
PPN: Physical Page Number

Radix page tables



Data TLB: 4-KB Pages, 4-way set associative, 64 entries

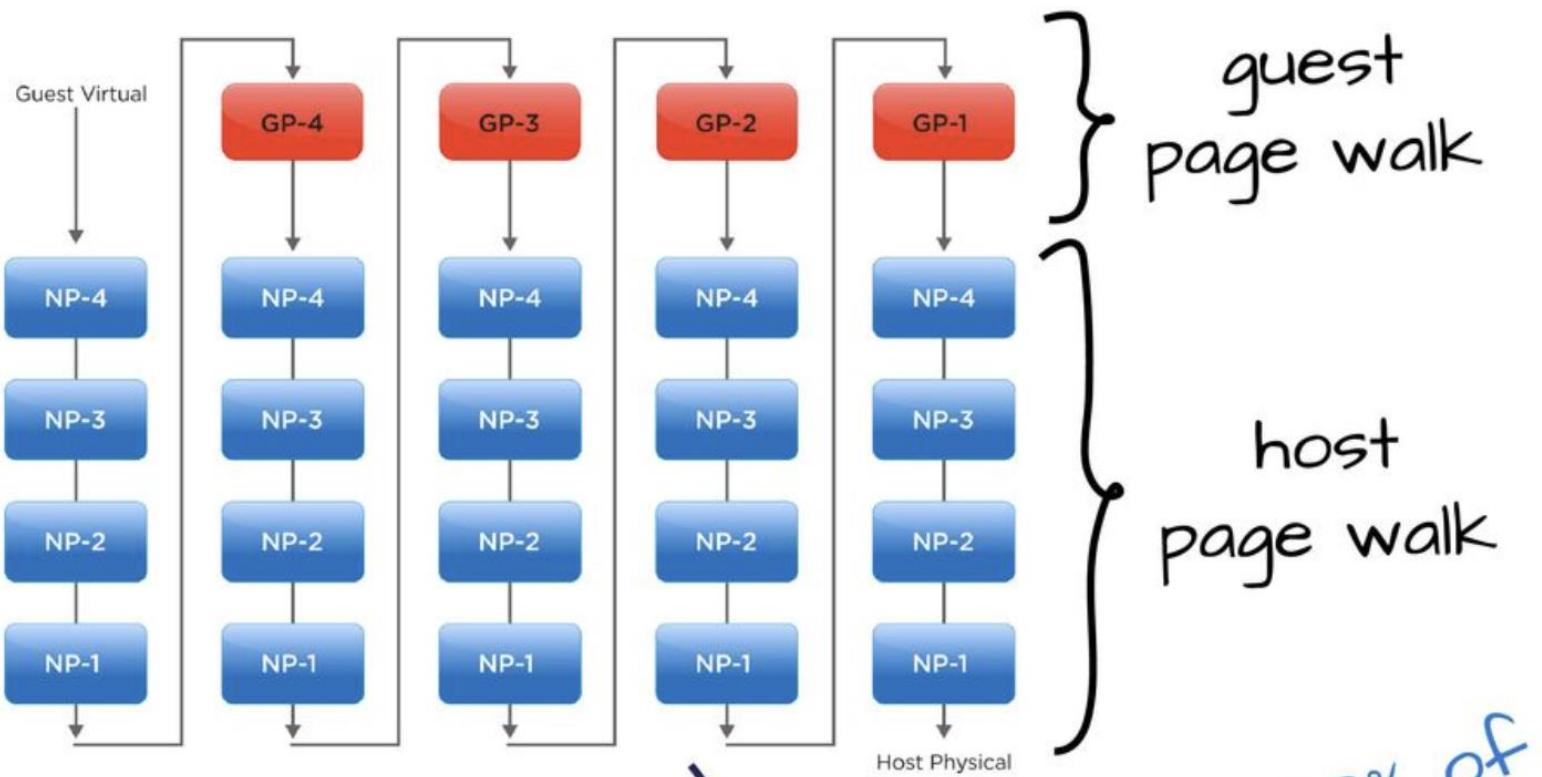
Instruction TLB: 4-KByte pages, 8-way set associative, 64 entries

Shared 2nd-Level TLB: 4-KB / 2-MB pages, 6-way associative, 1536 entries.

	[entries]	L1 cache	L2 cache
Ivy Bridge (2013)	512	64 KB	256 KB
Haswell (2013)	1024	64 KB	256 KB
Broadwell (2015)	1536	64 KB	256 KB



Virtualization requires 24 steps



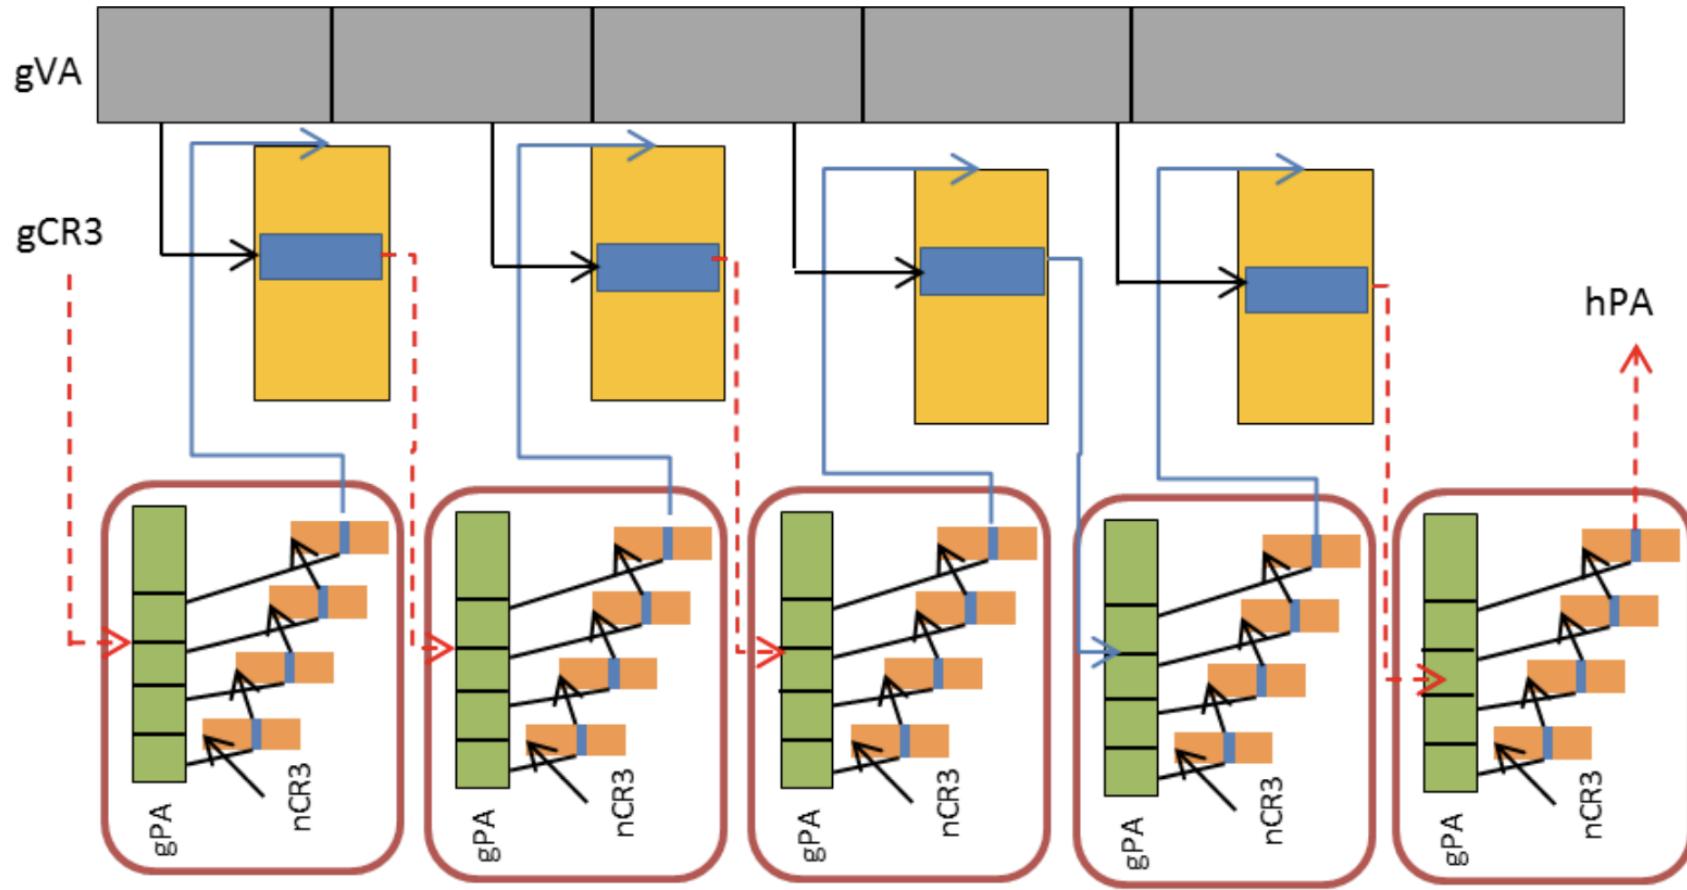
up to 90% of
the runtime!

TLB and Pagewalk Performance in Multicore Architectures with Large Die-Stacked DRAM Cache

Adarsh Patil

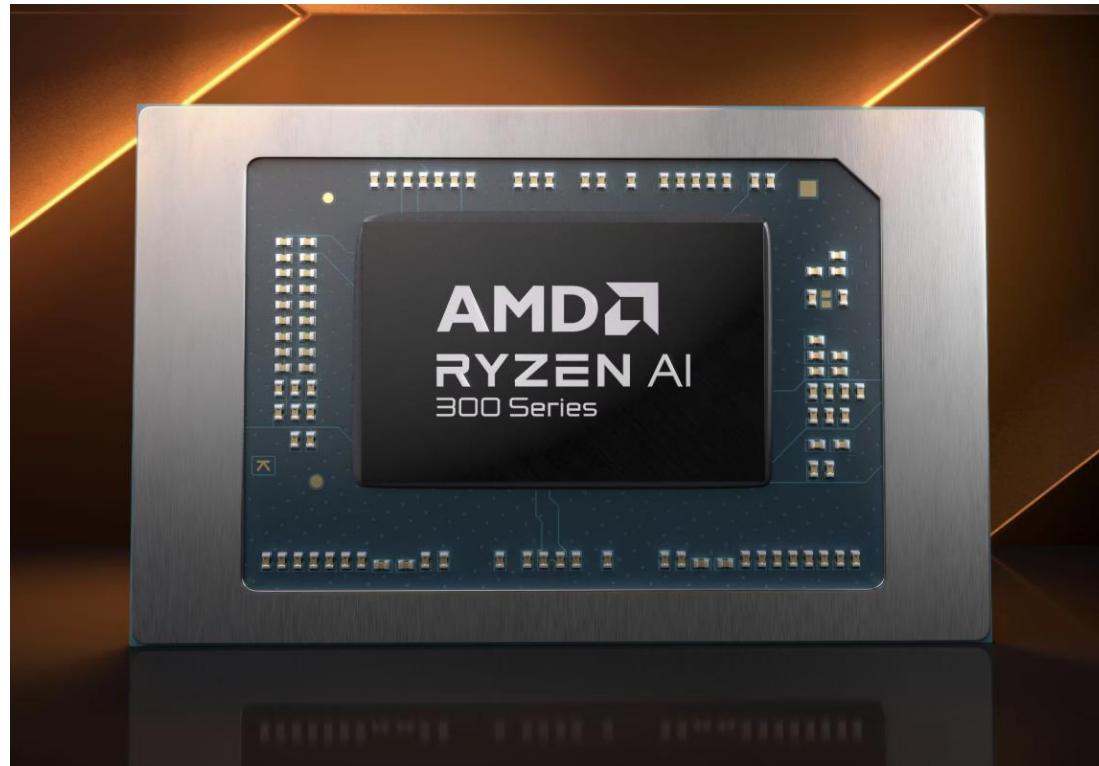
*Computer Science and Automation
Indian Institute of Sciences, Bangalore, IN
adarsh.patil@csa.iisc.ernet.in*



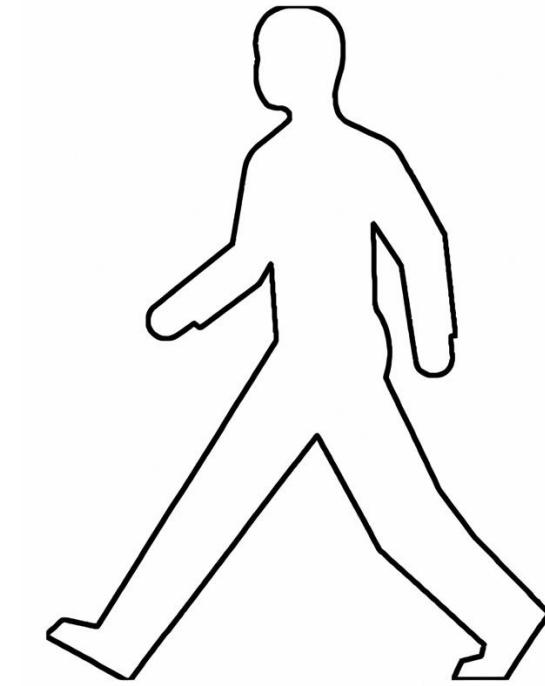


(b) 2D page walk in Virtualization

Page Walking Caches (PWC)



PWC is 16KB

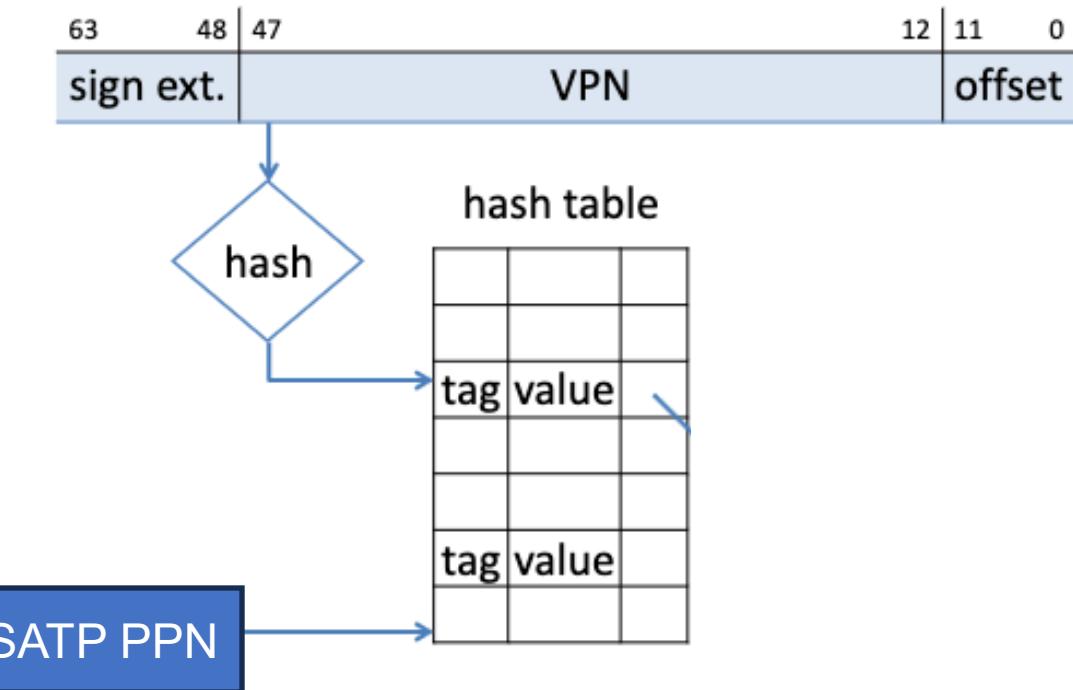


<https://esle.io/>

Alternatives to Radix Structure

Hash Page Table

- The Virtual Page Number (VPN) is hashed to find an entry in the hash table
- One large page table is shared among all processors, and the Hash is used as an index in that table for the translation.

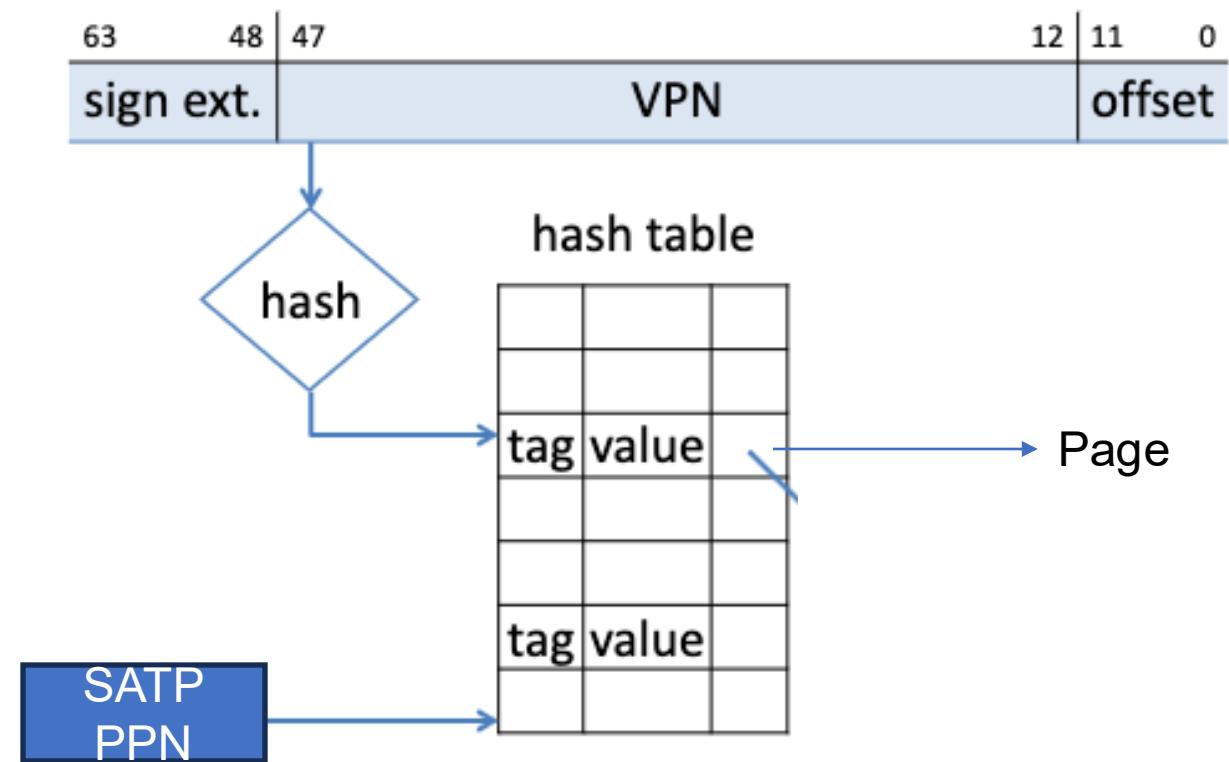


Traversing Hash Page Table

1) Hash VPN -> Position in Hash Table

2) Compare original VPN to tag

- a) VPN == tag
- b) VPN != tag



Radix vs Hash

	Radix	Hash

Overall, the literature mostly agrees that Radix is the winner!

Radix vs Hash Page Indexing

Perfect Hashing —

One memory access

Radix Structure —

Guaranteed at least
four memory accesses

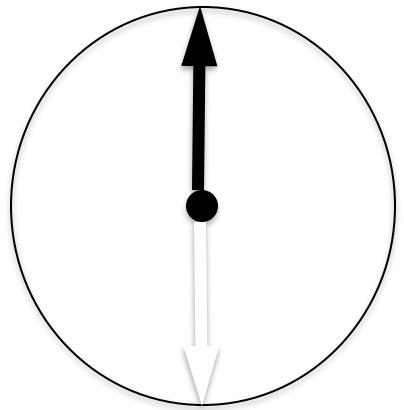
Page Replacement

To reduce page fault rate, prefer least-recently used (LRU) replacement

Reference bit (aka use bit) in PTE set to 1 on access to page

Periodically cleared to 0 by OS

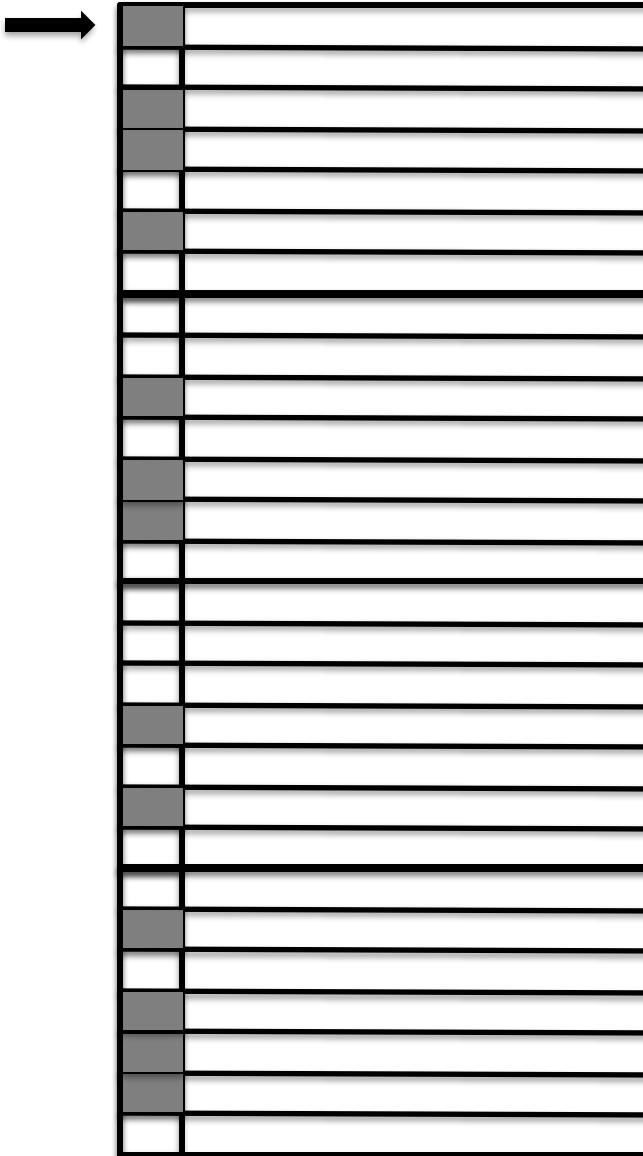
A page with reference bit = 0 has not been used recently



An LRU approximation used for page replacement

Search Pointer

Use Bit Page Table Entries



Writes

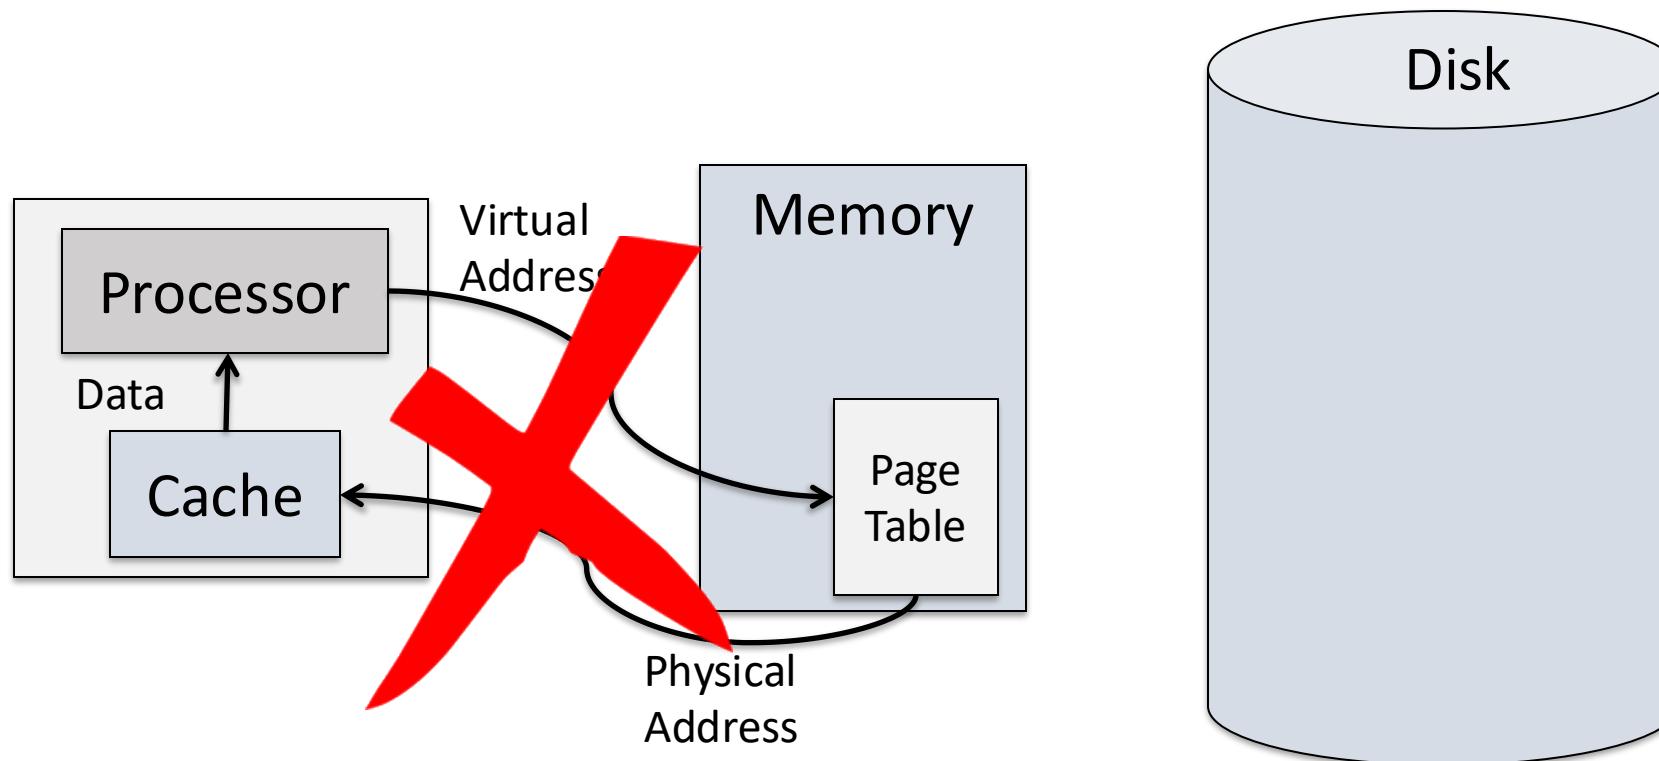
Disk writes take millions of cycles

Block at once, not on individual locations

Write through is impractical

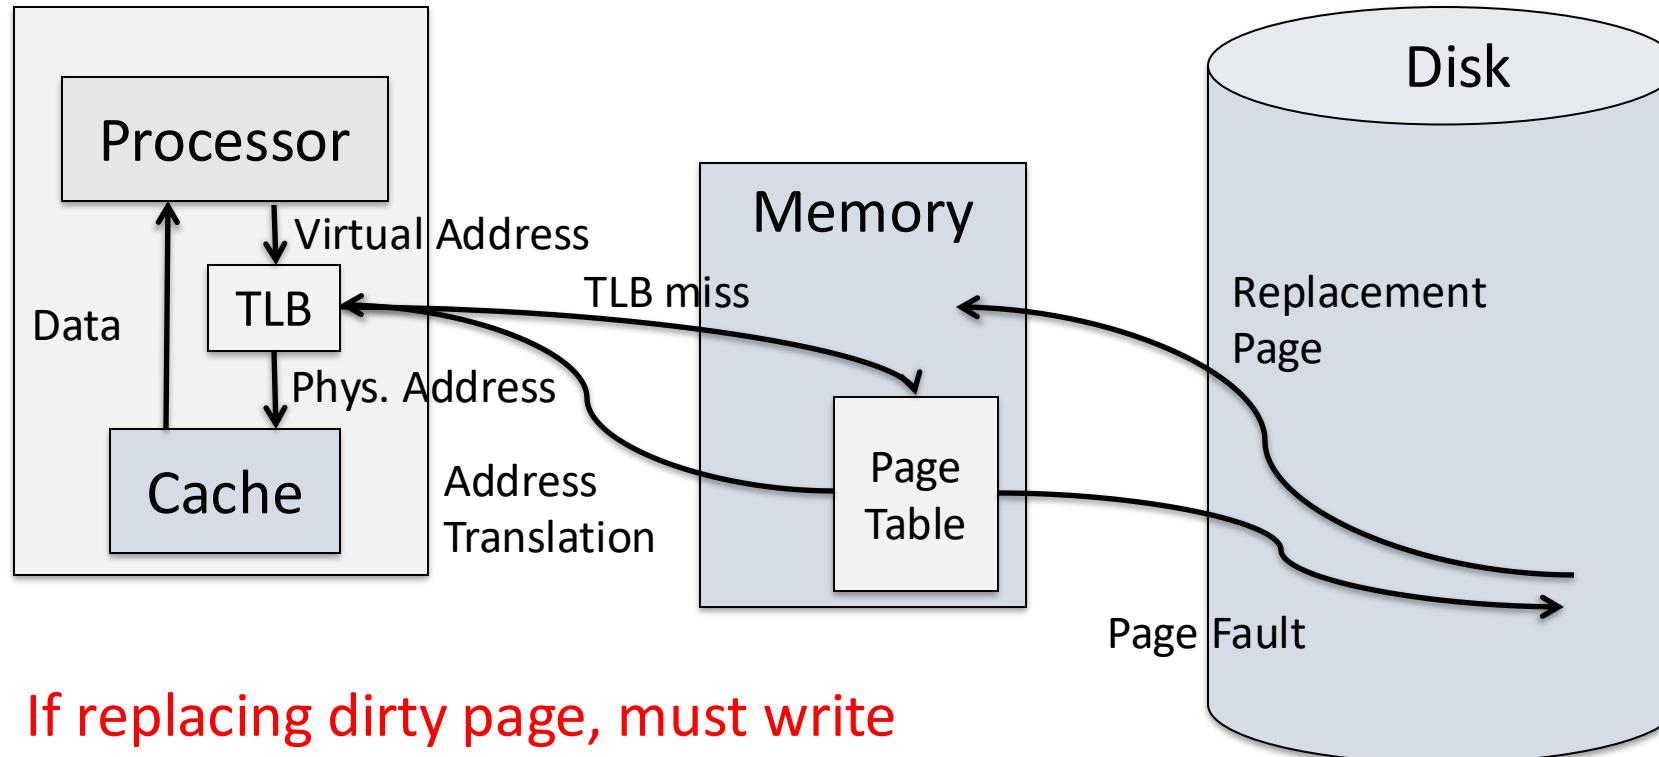
Use write-back

Do We Need an Extra Load for Each Memory Reference?



Translation-Lookaside Buffer (TLB)

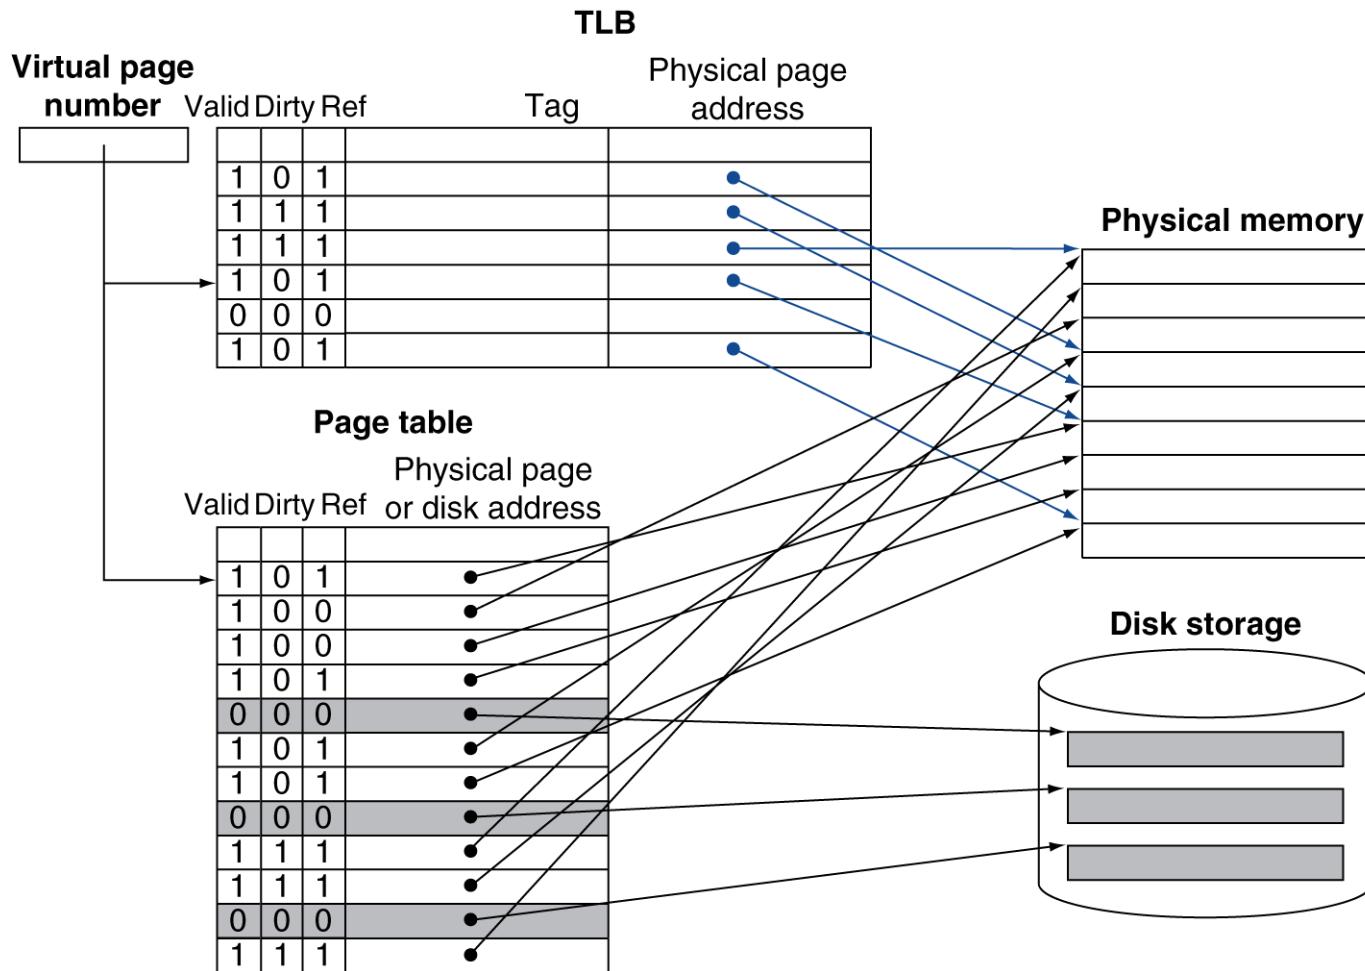
Page Fault is handled by OS



If replacing dirty page, must write it to disk first

Typical TLB: 16–512 Page Table Entries, 0.5–1 cycle for hit, 10–100 cycles for miss, 0.01%–1% miss rate

Fast Translation Using a TLB



TLB Misses

If page is in memory

- Load the Page Table Entry (PTE) from memory and retry

- Could be handled in hardware

- Can get complex for more complicated page table structures

- Or in software

- Raise a special exception, with optimized handler

If page is not in memory (page fault)

- OS handles fetching the page and updating the page table

TLB Miss Handler

TLB miss indicates either:

- Page present, but PTE not in TLB

- Page not present

Must recognize TLB miss before destination register overwritten

- Raise exception

Exception Handler copies PTE from memory to TLB

- Then restarts instruction

Page Fault Handler

Use faulting virtual address to find PTE

Locate page on disk

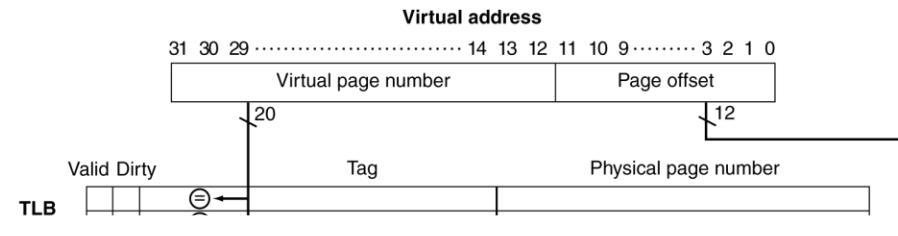
Choose page to replace

If dirty, write to disk first

Read page into memory and update page table

Make process runnable again

TLB and Cache Interaction



If cache tag uses physical address
Need to translate before cache lookup
Alternative: use virtual address tag
Complications due to aliasing

Memory Protection

Different tasks can share parts of their virtual address spaces

- But need to protect against errant access

- Requires OS assistance

Hardware support for OS protection

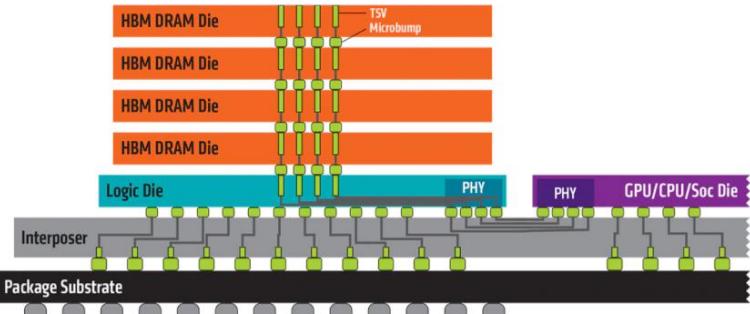
- Privileged supervisor mode (aka kernel mode)

- Privileged instructions

- Page tables and other state information only accessible in supervisor mode

- System call exception (e.g., ecall)

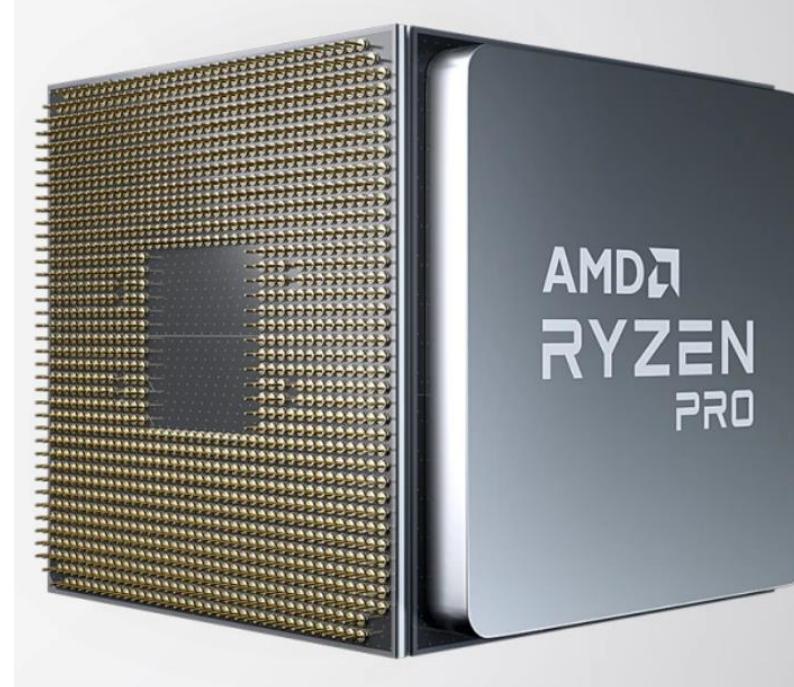
2021



Revolutionary HBM breaks the processing bottleneck

HBM is a new type of CPU/GPU memory ("RAM") that vertically stacks memory chips, like floors in a skyscraper. In doing so, it shortens your information commute. Those towers connect to the CPU or GPU through an ultra-fast interconnect called the "interposer." Several stacks of HBM are plugged into the interposer alongside a CPU or GPU, and that assembled module connects to a circuit board.

Though these HBM stacks are not physically integrated with the CPU or GPU, they are so closely and quickly connected via the interposer that HBM's characteristics are nearly indistinguishable from on-chip integrated RAM.



High Bandwidth
Memory