

# Topic V37

Memory Hierarchy Concepts

Readings: (Section 5.1, 5.3, 5.4, 5.6)

# The Memory Hierarchy

Common principles apply at all levels of the memory hierarchy

Based on notions of caching

At each level in the hierarchy

Block placement

Finding a block




Replacement on a miss

Write policy

# Block Placement

Determined by associativity

# Mapping



# Block Placement

Higher associativity reduces miss rate

Increases complexity, cost, and access time

# Finding a Block

Associativity	Location method	Tag comparisons
Direct mapped	Index	1

Hardware caches

Reduce comparisons to reduce cost

Virtual memory

Full table lookup makes full associativity feasible

Benefit in reduced miss rate

# Replacement

Choice of entry to replace on a miss

Least recently used (LRU)

Complex and costly hardware for high associativity

Not most-recently used with Random selection


Close to LRU, easier to implement

Virtual memory


LRU approximation with hardware support

# Write Policy

Write Back



Write Through



# Virtual Memory Write Policy

Only write-back policy is feasible, given disk write latency

# Source of Misses

Compulsory misses (aka cold start misses)

- First access to a block

Capacity misses

- Due to finite cache size

- A replaced block is later accessed again

Conflict misses (aka collision misses)

- In a non-fully associative cache

- Due to competition for entries in a set

- Would not occur in a fully associative cache of the same total size

# Cache Design Trade-offs

Design change	Effect on miss rate	Negative performance effect

# Virtual Machines

Host computer emulates guest operating system and machine resources

- Improved isolation of multiple guests

- Avoids security and reliability problems

- Aids sharing of resources

Virtualization has some performance impact

- Feasible with modern high-performance computers

Examples

- IBM VM/370 (1970s technology!)

- VMWare

- Microsoft Virtual PC

# Virtual Machine Monitor (VMM)

Maps virtual resources to physical resources

- Memory, I/O devices, CPUs

Guest code runs on native machine in user mode

- Traps to VMM on privileged instructions and access to protected resources

Guest OS may be different from host OS

VMM handles real I/O devices

- Emulates generic virtual I/O devices for guest

# Timer Virtualization (example)

In native machine, on timer interrupt

OS suspends current process, handles interrupt, selects and resumes next process

With Virtual Machine Monitor

VMM suspends current VM, handles interrupt, selects and resumes next VM


If a VM requires timer interrupts

VMM emulates a virtual timer

Emulates interrupt for VM when physical timer interrupt occurs

# Cache Coherence

# Example



1. P1 reads some data A
2. P2 reads same data A


// P1, P2 and memory have valid copies of A

3. P3 has a write miss on A
  - 3.1. A copy of A is brought into P3 Cache
  - 3.2. P3 modifies its copy of A

// Two possibilities: *write-update protocol* × *write-invalidate protocol*


# Example

1. P1 reads some data A
  2. P2 reads same data A  
*// P1, P2 and memory have valid copies of A*
  3. P3 has a write miss on A
    - 3.1. A copy of A is brought into P3 Cache
    - 3.2. P3 modifies its copy of A  
*// Two possibilities: *write-update protocol* × *write-invalidate protocol**
- // write-update protocol*  
*// (similar to write-through policy)*
- 3.3. P3 sends new A to all caches and to main memory  
*// A is valid in all caches and memory*



# Example

1. P1 reads some data A
2. P2 reads same data A  
*// P1, P2 and memory have valid copies of A*
3. P3 has a write miss on A
  - 3.1. A copy of A is brought into P3 Cache
  - 3.2. P3 modifies its copy of A  
*// Two possibilities: *write-update protocol* × *write-invalidate protocol**



*// write-update protocol*  
// (similar to write-through policy)

3.3. P3 sends new A to all caches and to main memory  
*// A is valid in all caches and memory*

*// write-invalidate protocol*  
// (similar to write-back policy)

3.3. P3 sends message invalidating A to bus  
// P3 has the only valid copy of A  
// memory has a stale copy of A

# Example

1. P1 reads some data A
2. P2 reads same data A

// P1, P2 and memory have valid copies of A

3. P3 has a write miss on A
  - 3.1. A copy of A is brought into P3 Cache
  - 3.2. P3 modifies its copy of A


// Two possibilities: *write-update protocol* × *write-invalidate protocol*

// write-update protocol

// (similar to write-through policy)

- 3.3. P3 sends new A to all caches and to main memory

// A is valid in all caches and memory



// write-invalidate protocol

// (similar to write-back policy)

- 3.3. P3 sends message invalidating A to bus

// P3 has the only valid copy of A

// memory still has a stale copy of A

# Cache Coherence Protocols

Operations performed by caches in multiprocessors to ensure coherence

- Migration of data to local caches

  - Reduces bandwidth for shared memory

- Replication of read-shared data

  - Reduces contention for access

## Snooping protocols

- Each cache monitors bus reads/writes

## Directory-based protocols

- Caches and memory record sharing status of blocks in a directory

# Memory Consistency

When are writes seen by other processors?

“Seen” means a read returns the written value

Can’t be instantaneously

Assumptions

A write completes only when all processors have seen it

A processor does not reorder writes with other accesses


Consequence

P writes X then writes Y  $\Rightarrow$  all processors that see new Y also see new X

Processors can reorder reads, but not writes

# Multilevel On-chip Caches

Intel Nehalem 4-core processor



Per core: 32KB L1 I-cache, 32KB L1 D-cache, 512KB L2 cache

# 2-level TLB Organization

	Intel Nehalem	AMD Opteron X4
--	---------------	----------------

# 3-level Cache Organization

	Intel Nehalem	AMD Opteron X4
L1 caches (per core)	L1 I-cache: 32KB, 64-byte blocks, 4-way, approx LRU replacement, hit time n/a	L1 I-cache: 32KB, 64-byte blocks, 2-way, LRU replacement, hit time 3 cycles

# Reducing Miss Penalty

Return requested word first

Then back-fill rest of block

Non-blocking miss processing

Hit under miss: allow hits to proceed

Miss under miss: allow multiple outstanding misses

Hardware prefetch: instructions and data

Opteron X4: bank-interleaved L1 D-cache

Two concurrent accesses per cycle

# Pitfalls

## Byte vs. word addressing

Example: 32-byte direct-mapped cache, 4-byte blocks

Byte 36 maps to block 1

Word 36 maps to block 4

## Ignoring memory system effects when writing or generating code

Example: iterating over rows vs. columns of arrays

Large strides result in poor locality

# Pitfalls

In multiprocessor with shared L2 or L3 cache

- Less associativity than cores results in conflict misses

- More cores ⇒ need to increase associativity

Using AMAT to evaluate performance of out-of-order processors

- Ignores effect of non-blocked accesses

- Instead, evaluate performance by simulation

# Concluding Remarks

Fast memories are small, large memories are slow

We really want fast, large memories ☺

Caching gives this illusion ☺

Principle of locality

Programs use a small part of their memory space frequently

Memory hierarchy

L1 cache < L2 cache < ... < DRAM memory < disk

Memory system design is critical for multiprocessors